

# Next Token Generation

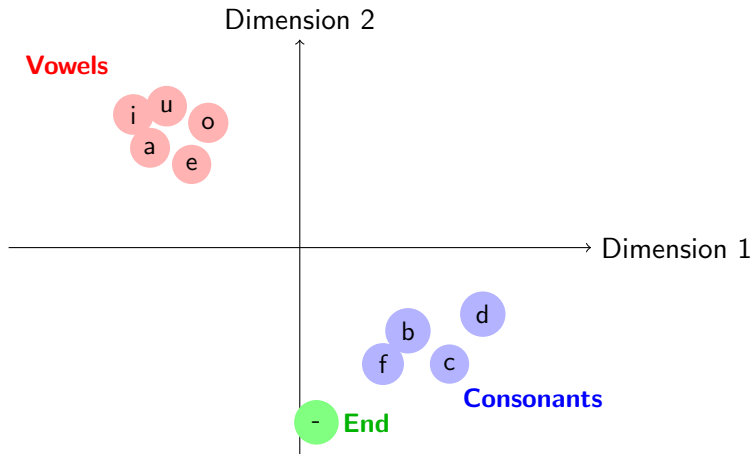
Nipun Batra

IIT Gandhinagar

July 29, 2025

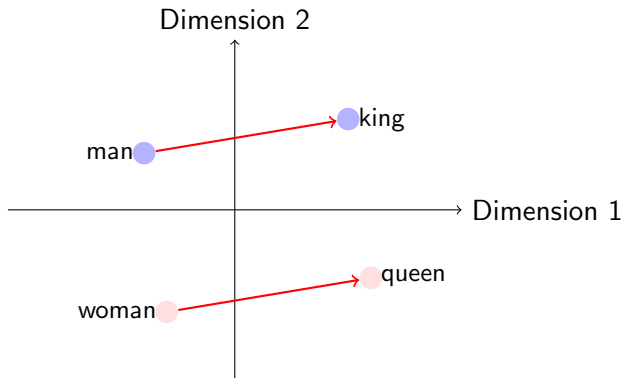
### **Vocabulary Size:**

26 letters + 1 hyphen = **27 characters**



# Word2Vec Analogy Example

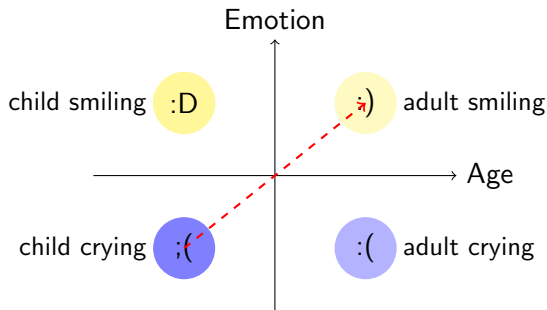
## Classic Word2Vec Relationship



**Relationship:**  $\text{queen} \approx \text{king} - \text{man} + \text{woman}$


# Analogy with Emotions

## Emotional Expression Analogy



**Relationship:** child crying = child smiling + adult crying - adult smiling

# Embedding Matrix/Table Concept



**Process:** Character  $\rightarrow$  Lookup in Embedding Table  $\rightarrow$  Dense Vector

# Embedding Table Structure

## $27 \times K$ Embedding Matrix

Char	D1	D2	...	DK
a	0.2	-0.1	...	0.8
b	-0.3	0.5	...	-0.2
c	0.1	0.3	...	0.4
$\vdots$	$\vdots$	$\vdots$	$\ddots$	$\vdots$
z	0.7	-0.4	...	0.1
-	0.0	0.9	...	-0.5

## Key Point

Each character maps to a K-dimensional vector.

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters



# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation
  - ▶ Learns meaningful character representations

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation
  - ▶ Learns meaningful character representations
- ▶ **Neural Network Weights:** MLP parameters

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation
  - ▶ Learns meaningful character representations
- ▶ **Neural Network Weights:** MLP parameters
  - ▶ Transform concatenated embeddings to output

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation
  - ▶ Learns meaningful character representations
- ▶ **Neural Network Weights:** MLP parameters
  - ▶ Transform concatenated embeddings to output
  - ▶ Learn classification patterns

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation
  - ▶ Learns meaningful character representations
- ▶ **Neural Network Weights:** MLP parameters
  - ▶ Transform concatenated embeddings to output
  - ▶ Learn classification patterns
- ▶ **Total Learnable Parameters:**



# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation
  - ▶ Learns meaningful character representations
- ▶ **Neural Network Weights:** MLP parameters
  - ▶ Transform concatenated embeddings to output
  - ▶ Learn classification patterns
- ▶ **Total Learnable Parameters:**
  - ▶ Embedding:  $27 \times K$

# Learnable Parameters

- ▶ **Embedding Matrix:**  $27 \times K$  parameters
  - ▶ Initially random
  - ▶ Updated during training via backpropagation
  - ▶ Learns meaningful character representations
- ▶ **Neural Network Weights:** MLP parameters
  - ▶ Transform concatenated embeddings to output
  - ▶ Learn classification patterns
- ▶ **Total Learnable Parameters:**
  - ▶ Embedding:  $27 \times K$
  - ▶ MLP:  $(\text{context\_size} \times K) \rightarrow \text{hidden} \rightarrow \dots \rightarrow 27$

## Example: 2D Embeddings for “abi”


Embedding Matrix ( $27 \times 2$ )

Input:  $X = [\text{“a”}, \text{“b”}, \text{“i”}]$

	D1	D2	
a	0.2	-0.1	$[0.2, -0.1]$
t	-0.3	0.5	$[-0.3, 0.5]$
...	...	...	
i	0.1	0.3	$[0.1, 0.3]$
...	...	...	
z	0.7	-0.4	
-	0.0	0.9	

# Concatenate the Embeddings

## Feature Vector Construction




**6D feature vector**

## Result

3 chars  $\times$  2D embeddings = 6D input to neural network

# Multi-Layer Perceptron Architecture



# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**



# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
2. **MLP Weights:** Neural network parameters for classification

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
2. **MLP Weights:** Neural network parameters for classification

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**
  1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
  2. **MLP Weights:** Neural network parameters for classification
- ▶ **Training Process:**

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**
  1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
  2. **MLP Weights:** Neural network parameters for classification
- ▶ **Training Process:**

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
2. **MLP Weights:** Neural network parameters for classification

- ▶ **Training Process:**

1. Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities



# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
2. **MLP Weights:** Neural network parameters for classification

- ▶ **Training Process:**

1. Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities
2. Compute cross-entropy loss

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
2. **MLP Weights:** Neural network parameters for classification

- ▶ **Training Process:**

1. Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities
2. Compute cross-entropy loss
3. Backward pass: Update both embeddings and MLP weights

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
2. **MLP Weights:** Neural network parameters for classification

- ▶ **Training Process:**

1. Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities
2. Compute cross-entropy loss
3. Backward pass: Update both embeddings and MLP weights

# Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ( $27 \times K$  parameters)
2. **MLP Weights:** Neural network parameters for classification

- ▶ **Training Process:**

1. Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities
2. Compute cross-entropy loss
3. Backward pass: Update both embeddings and MLP weights
4. Repeat for all training examples

# Sampling from the Learned Model

Test Input: “abi”

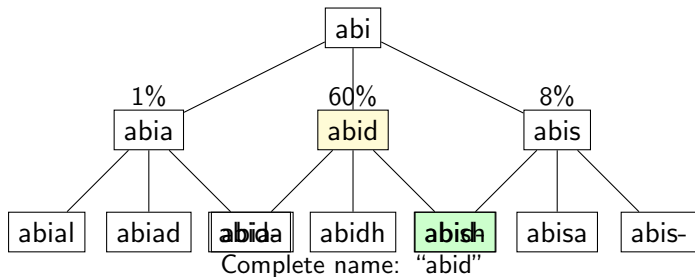
## Predicted Probability Distribution

Next Char	Probability	Next Char	Probability
a	0.01	n	0.05
b	0.01	o	0.02
c	0.03	p	0.01
d	<b>0.60</b>	q	0.00
e	0.02	r	0.03
f	0.01	s	0.08
...	...	...	...
-	0.05	z	0.01

Most Likely Continuation

“abi” → “abid” (60

# Generation Tree Structure



**Recursive Process:** Sample next character, append, repeat until end token

# Temperature in Softmax

## ► Standard Softmax:

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

# Temperature in Softmax

## ► Standard Softmax:

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$



# Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

# Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

# Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- ▶ **Temperature Effects:**

# Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- ▶ **Temperature Effects:**

- ▶  $T = 1$ : Standard probabilities

# Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- ▶ **Temperature Effects:**

- ▶  $T = 1$ : Standard probabilities
- ▶  $T \rightarrow 0$ : More peaked (deterministic)

# Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- ▶ **Temperature Effects:**

- ▶  $T = 1$ : Standard probabilities
- ▶  $T \rightarrow 0$ : More peaked (deterministic)

# Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- ▶ **Temperature Effects:**

- ▶  $T = 1$ : Standard probabilities
- ▶  $T \rightarrow 0$ : More peaked (deterministic)
- ▶  $T \rightarrow \infty$ : More uniform (random)

# Temperature Variations

**Context:** “abi” → Next character probabilities

Char	T=0.5 (Low)	T=1.0 (Default)	T=2.0 (High)
a	0.001	0.01	0.08
d	<b>0.95</b>	<b>0.60</b>	<b>0.25</b>
s	0.01	0.08	0.12
h	0.005	0.03	0.09
-	0.02	0.05	0.11
others	0.015	0.23	0.35

- **Low T:** Conservative, predictable



# Temperature Variations

**Context:** “abi” → Next character probabilities

Char	T=0.5 (Low)	T=1.0 (Default)	T=2.0 (High)
a	0.001	0.01	0.08
d	<b>0.95</b>	<b>0.60</b>	<b>0.25</b>
s	0.01	0.08	0.12
h	0.005	0.03	0.09
-	0.02	0.05	0.11
others	0.015	0.23	0.35

- **Low T:** Conservative, predictable

# Temperature Variations

**Context:** “abi” → Next character probabilities

Char	T=0.5 (Low)	T=1.0 (Default)	T=2.0 (High)
a	0.001	0.01	0.08
d	<b>0.95</b>	<b>0.60</b>	<b>0.25</b>
s	0.01	0.08	0.12
h	0.005	0.03	0.09
-	0.02	0.05	0.11
others	0.015	0.23	0.35

- ▶ **Low T:** Conservative, predictable
- ▶ **High T:** Creative, diverse

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling



# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models
  - ▶ GPT models use the same principle

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models
  - ▶ GPT models use the same principle
  - ▶ Scaled to words/subwords instead of characters

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models
  - ▶ GPT models use the same principle
  - ▶ Scaled to words/subwords instead of characters



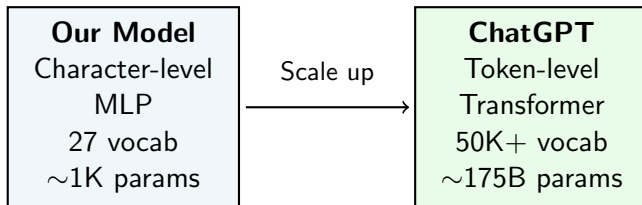
# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models
  - ▶ GPT models use the same principle
  - ▶ Scaled to words/subwords instead of characters
  - ▶ Transformer architecture instead of MLP

# Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models
  - ▶ GPT models use the same principle
  - ▶ Scaled to words/subwords instead of characters
  - ▶ Transformer architecture instead of MLP
  - ▶ Billions of parameters instead of thousands

# From Character-Level to ChatGPT



**Same fundamental principle: Predict the next token!**