

Next Token Generation

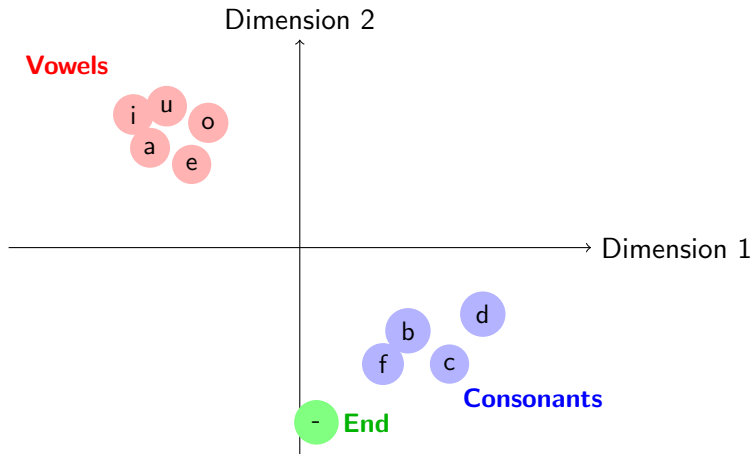
Nipun Batra

IIT Gandhinagar

July 29, 2025

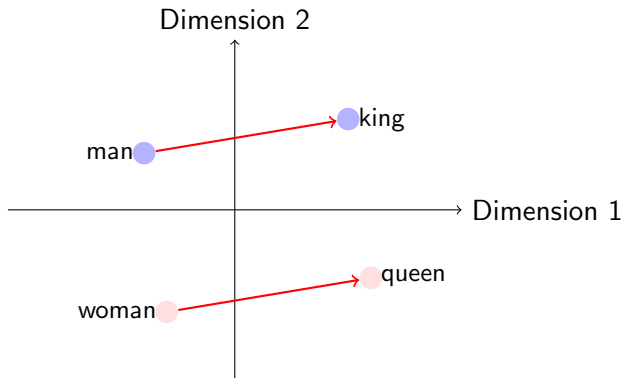
Vocabulary Size:

26 letters + 1 hyphen = **27 characters**



Word2Vec Analogy Example

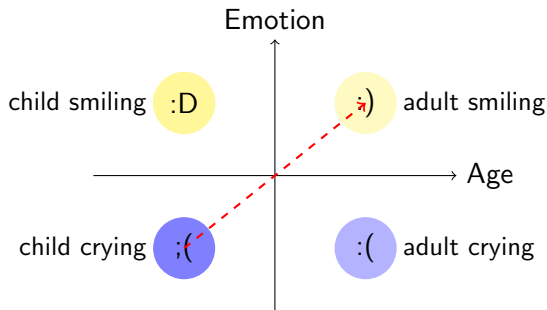
Classic Word2Vec Relationship



Relationship: $\text{queen} \approx \text{king} - \text{man} + \text{woman}$


Analogy with Emotions

Emotional Expression Analogy



Relationship: child crying = child smiling + adult crying - adult smiling

Embedding Matrix/Table Concept



Process: Character \rightarrow Lookup in Embedding Table \rightarrow Dense Vector

Embedding Table Structure

$27 \times K$ Embedding Matrix

Char	D1	D2	...	DK
a	0.2	-0.1	...	0.8
b	-0.3	0.5	...	-0.2
c	0.1	0.3	...	0.4
\vdots	\vdots	\vdots	\ddots	\vdots
z	0.7	-0.4	...	0.1
-	0.0	0.9	...	-0.5

Key Point

Each character maps to a K-dimensional vector.

Learnable Parameters

- ▶ **Embedding Matrix:** $27 \times K$ parameters
 - ▶ Initially random

Learnable Parameters

- ▶ **Embedding Matrix:** $27 \times K$ parameters
 - ▶ Initially random
 - ▶ Updated during training via backpropagation
 - ▶ Learns meaningful character representations
- ▶ **Neural Network Weights:** MLP parameters
 - ▶ Transform concatenated embeddings to output
 - ▶ Learn classification patterns
- ▶ **Total Learnable Parameters:**
 - ▶ Embedding: $27 \times K$
 - ▶ MLP: $(\text{context_size} \times K) \rightarrow \text{hidden} \rightarrow \dots \rightarrow 27$

Example: 2D Embeddings for “abi”

Embedding Matrix (27×2)

Input: $X = [\text{“a”}, \text{“b”}, \text{“i”}]$

	D1	D2	
a	0.2	-0.1	[0.2, -0.1]
t	-0.3	0.5	[-0.3, 0.5]
...	
i	0.1	0.3	[0.1, 0.3]
...	
z	0.7	-0.4	
-	0.0	0.9	

Concatenate the Embeddings

Feature Vector Construction

a: [0.2, -0.1]

b: [-0.3, 0.5]

i: [0.1, 0.3]

concat


[0.2, -0.1, -0.3, 0.5, 0.1, 0.3]

6D feature vector

Result

3 chars \times 2D embeddings = 6D input to neural network

Multi-Layer Perceptron Architecture



Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)

Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
2. **MLP Weights:** Neural network parameters for classification

Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**
 1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
 2. **MLP Weights:** Neural network parameters for classification
- ▶ **Training Process:**

Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
2. **MLP Weights:** Neural network parameters for classification

- ▶ **Training Process:**

1. Forward pass: Input \rightarrow Embeddings \rightarrow Concatenate \rightarrow MLP \rightarrow Probabilities
2. Compute cross-entropy loss
3. Backward pass: Update both embeddings and MLP weights

Training Objective

- ▶ **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- ▶ **What we learn:**

1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
2. **MLP Weights:** Neural network parameters for classification

- ▶ **Training Process:**

1. Forward pass: Input \rightarrow Embeddings \rightarrow Concatenate \rightarrow MLP \rightarrow Probabilities
2. Compute cross-entropy loss
3. Backward pass: Update both embeddings and MLP weights
4. Repeat for all training examples

Sampling from the Learned Model

Test Input: “abi”


Predicted Probability Distribution

Next Char	Probability	Next Char	Probability
a	0.01	n	0.05
b	0.01	o	0.02
c	0.03	p	0.01
d	0.60	q	0.00
e	0.02	r	0.03
f	0.01	s	0.08
...
-	0.05	z	0.01

Most Likely Continuation

“abi” → “abid” (60

Generation Tree Structure



Recursive Process: Sample next character, append, repeat until end token

Temperature in Softmax

► Standard Softmax:

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- ▶ **Temperature Effects:**

- ▶ $T = 1$: Standard probabilities
- ▶ $T \rightarrow 0$: More peaked (deterministic)

Temperature in Softmax

- ▶ **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- ▶ **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- ▶ **Temperature Effects:**

- ▶ $T = 1$: Standard probabilities
- ▶ $T \rightarrow 0$: More peaked (deterministic)
- ▶ $T \rightarrow \infty$: More uniform (random)

Temperature Variations

Context: “abi” → Next character probabilities

Char	T=0.5 (Low)	T=1.0 (Default)	T=2.0 (High)
a	0.001	0.01	0.08
d	0.95	0.60	0.25
s	0.01	0.08	0.12
h	0.005	0.03	0.09
-	0.02	0.05	0.11
others	0.015	0.23	0.35

- **Low T:** Conservative, predictable

Temperature Variations

Context: “abi” → Next character probabilities

Char	T=0.5 (Low)	T=1.0 (Default)	T=2.0 (High)
a	0.001	0.01	0.08
d	0.95	0.60	0.25
s	0.01	0.08	0.12
h	0.005	0.03	0.09
-	0.02	0.05	0.11
others	0.015	0.23	0.35

- ▶ **Low T:** Conservative, predictable
- ▶ **High T:** Creative, diverse

Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification

Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity

Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling

Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights

Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control

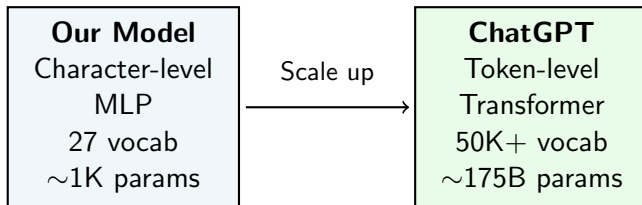
Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models
 - ▶ GPT models use the same principle
 - ▶ Scaled to words/subwords instead of characters

Key Takeaways

- ▶ **Core Idea:** Next token prediction as classification
- ▶ **Representation Learning:** Character embeddings capture similarity
- ▶ **Architecture:** Embeddings + MLP for sequence modeling
- ▶ **Training:** Joint learning of embeddings and classifier weights
- ▶ **Generation:** Autoregressive sampling with temperature control
- ▶ **Applications:** Foundation for modern language models
 - ▶ GPT models use the same principle
 - ▶ Scaled to words/subwords instead of characters
 - ▶ Transformer architecture instead of MLP
 - ▶ Billions of parameters instead of thousands

From Character-Level to ChatGPT



Same fundamental principle: Predict the next token!