

Next Token Generation

Nipun Batra

July 26, 2025

IIT Gandhinagar

Table of Contents

Introduction and Motivation

Problem Formulation

Case Study: Indian Names Generation

Training Data Generation

Representation Learning

Embedding Architecture

Neural Network Architecture

Training and Loss Function

Text Generation

Temperature and Sampling Strategies

Summary and Applications

Introduction and Motivation

Acknowledgment and Inspiration

- This lecture is inspired by the excellent work of **Andrej Karpathy**

Acknowledgment and Inspiration

- This lecture is inspired by the excellent work of **Andrej Karpathy**
- Search for *“Neural Networks: Zero to Hero”* to find his comprehensive tutorial series

Acknowledgment and Inspiration

- This lecture is inspired by the excellent work of **Andrej Karpathy**
- Search for “*Neural Networks: Zero to Hero*” to find his comprehensive tutorial series
- These concepts are fundamental to understanding modern language models

Acknowledgment and Inspiration

- This lecture is inspired by the excellent work of **Andrej Karpathy**
- Search for *“Neural Networks: Zero to Hero”* to find his comprehensive tutorial series
- These concepts are fundamental to understanding modern language models
- Direct relevance to **ChatGPT** and other large language models:

Acknowledgment and Inspiration

- This lecture is inspired by the excellent work of **Andrej Karpathy**
- Search for *“Neural Networks: Zero to Hero”* to find his comprehensive tutorial series
- These concepts are fundamental to understanding modern language models
- Direct relevance to **ChatGPT** and other large language models:
 - ChatGPT generates text by predicting the next token

Acknowledgment and Inspiration

- This lecture is inspired by the excellent work of **Andrej Karpathy**
- Search for *“Neural Networks: Zero to Hero”* to find his comprehensive tutorial series
- These concepts are fundamental to understanding modern language models
- Direct relevance to **ChatGPT** and other large language models:
 - ChatGPT generates text by predicting the next token
 - Same underlying principle scaled to billions of parameters

Acknowledgment and Inspiration

- This lecture is inspired by the excellent work of **Andrej Karpathy**
- Search for “*Neural Networks: Zero to Hero*” to find his comprehensive tutorial series
- These concepts are fundamental to understanding modern language models
- Direct relevance to **ChatGPT** and other large language models:
 - ChatGPT generates text by predicting the next token
 - Same underlying principle scaled to billions of parameters
 - Understanding next token prediction is key to understanding LLMs

Given the sequence “app”, what is the next character?

Problem Formulation

Next Character Prediction as Classification

Next Character Prediction as Classification

a

p

p

Input: “app”

Next Character Prediction as Classification

Output: Character Probabilities

a

p

p

Input: “app”

Character	Probability
a	0.05
b	0.02
c	0.03
...	...
l	0.35
m	0.01
...	...
y	0.08
z	0.01
- (end)	0.15

Classification Task: Predict probability distribution over all

Case Study: Indian Names Generation

Dataset: Indian Names

Training Dataset

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran

Goal: Learn to generate new, realistic Indian names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh
- Sunita, Vikash, Zara

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Dataset Properties:

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh
- Sunita, Vikash, Zara

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Dataset Properties:

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh
- Sunita, Vikash, Zara

- 1000+ unique names

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh
- Sunita, Vikash, Zara

Dataset Properties:

- 1000+ unique names
- Diverse regional origins

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh
- Sunita, Vikash, Zara

Dataset Properties:

- 1000+ unique names
- Diverse regional origins
- Various lengths (4-10 chars)

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh
- Sunita, Vikash, Zara

Dataset Properties:

- 1000+ unique names
- Diverse regional origins
- Various lengths (4-10 chars)
- Both male & female names

Goal: Learn to generate new, realistic Indian names

Dataset: Indian Names

Training Dataset

Sample Names:

- Abid, Abhidha, Adesh
- Aditya, Arjun, Kiran
- Krishna, Lakshmi, Meera
- Nisha, Priya, Rajesh
- Sunita, Vikash, Zara

Dataset Properties:

- 1000+ unique names
- Diverse regional origins
- Various lengths (4-10 chars)
- Both male & female names
- Rich phonetic patterns

Goal: Learn to generate new, realistic Indian names

Assumptions and Constraints

- **Character Set:** Only 26 lowercase letters (a-z)

Assumptions and Constraints

- **Character Set:** Only 26 lowercase letters (a-z)
- **End Marker:** Hyphen (-) indicates end of name

Assumptions and Constraints

- **Character Set:** Only 26 lowercase letters (a-z)
- **End Marker:** Hyphen (-) indicates end of name
- **Length Constraint:** Names are between 4 and 10 characters

Assumptions and Constraints

- **Character Set:** Only 26 lowercase letters (a-z)
- **End Marker:** Hyphen (-) indicates end of name
- **Length Constraint:** Names are between 4 and 10 characters
- **Simplification:** No spaces, special characters, or uppercase letters

Assumptions and Constraints

- **Character Set:** Only 26 lowercase letters (a-z)
- **End Marker:** Hyphen (-) indicates end of name
- **Length Constraint:** Names are between 4 and 10 characters
- **Simplification:** No spaces, special characters, or uppercase letters

Assumptions and Constraints

- **Character Set:** Only 26 lowercase letters (a-z)
- **End Marker:** Hyphen (-) indicates end of name
- **Length Constraint:** Names are between 4 and 10 characters
- **Simplification:** No spaces, special characters, or uppercase letters

Vocabulary Size:

26 letters + 1 hyphen = **27 characters**

Training Data Generation

Generate Training Dataset

Example: “abid” → Training Examples

Context Length: 3 characters

Input (X)				Target (Y)
Char 1	Char 2	Char 3	Context	Next Char
-	-	-	“__”	a
-	-	a	“-a”	b
-	a	b	“-ab”	i
a	b	i	“abi”	d
b	i	d	“bid”	-

Training Examples

From one name “abid”, we create **5 training examples** using a 7 / 23

Representation Learning

The Idea: Character Embeddings

- **Goal:** Learn vector representations for each character

The Idea: Character Embeddings

- **Goal:** Learn vector representations for each character
- **Hypothesis:** Similar characters should have similar embeddings

The Idea: Character Embeddings


- **Goal:** Learn vector representations for each character
- **Hypothesis:** Similar characters should have similar embeddings
- **Benefit:** Capture semantic relationships between characters

The Idea: Character Embeddings

- **Goal:** Learn vector representations for each character
- **Hypothesis:** Similar characters should have similar embeddings
- **Benefit:** Capture semantic relationships between characters

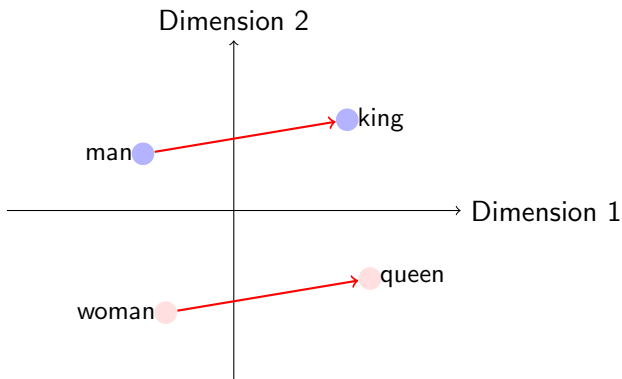
The Idea: Character Embeddings

- **Goal:** Learn vector representations for each character
- **Hypothesis:** Similar characters should have similar embeddings
- **Benefit:** Capture semantic relationships between characters



Word2Vec Analogy Example


Classic Word2Vec Relationship



Relationship: $\text{queen} \approx \text{king} - \text{man} + \text{woman}$

Analogy with Emotions


Emotional Expression Analogy



Relationship: child crying = child smiling + adult crying - adult smiling

Embedding Architecture

Embedding Matrix/Table Concept



Process: Character → Lookup in Embedding Table → Dense Vector

Embedding Table Structure

$27 \times K$ Embedding Matrix

Char	D1	D2	...	DK
a	0.2	-0.1	...	0.8
b	-0.3	0.5	...	-0.2
c	0.1	0.3	...	0.4
⋮	⋮	⋮	⋮	⋮
z	0.7	-0.4	...	0.1
-	0.0	0.9	...	-0.5

Key Point

Each character maps to a K-dimensional vector.

- **Embedding Matrix:** $27 \times K$ parameters

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation
 - Learns meaningful character representations

Learnable Parameters

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation
 - Learns meaningful character representations
- **Neural Network Weights:** MLP parameters

Learnable Parameters

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation
 - Learns meaningful character representations
- **Neural Network Weights:** MLP parameters
 - Transform concatenated embeddings to output

Learnable Parameters

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation
 - Learns meaningful character representations
- **Neural Network Weights:** MLP parameters
 - Transform concatenated embeddings to output
 - Learn classification patterns

Learnable Parameters

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation
 - Learns meaningful character representations
- **Neural Network Weights:** MLP parameters
 - Transform concatenated embeddings to output
 - Learn classification patterns
- **Total Learnable Parameters:**

Learnable Parameters

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation
 - Learns meaningful character representations
- **Neural Network Weights:** MLP parameters
 - Transform concatenated embeddings to output
 - Learn classification patterns
- **Total Learnable Parameters:**
 - Embedding: $27 \times K$

Learnable Parameters

- **Embedding Matrix:** $27 \times K$ parameters
 - Initially random
 - Updated during training via backpropagation
 - Learns meaningful character representations
- **Neural Network Weights:** MLP parameters
 - Transform concatenated embeddings to output
 - Learn classification patterns
- **Total Learnable Parameters:**
 - Embedding: $27 \times K$
 - MLP: $(\text{context_size} \times K) \rightarrow \text{hidden} \rightarrow \dots \rightarrow 27$

Neural Network Architecture

Example: 2D Embeddings for “abi”


Embedding Matrix (27×2)

Input: $X = [\text{“a”}, \text{“b”}, \text{“i”}]$

	D1	D2	
a	0.2	-0.1	[0.2, -0.1]
t	-0.3	0.5	[-0.3, 0.5]
...	
i	0.1	0.3	[0.1, 0.3]
...	
z	0.7	-0.4	
-	0.0	0.9	

Concatenate the Embeddings


Feature Vector Construction



Result

3 chars \times 2D embeddings = 6D input to neural network

Multi-Layer Perceptron Architecture



Training and Loss Function

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**

1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**
 1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
 2. **MLP Weights:** Neural network parameters for classification

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**
 1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
 2. **MLP Weights:** Neural network parameters for classification
- **Training Process:**

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**
 1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
 2. **MLP Weights:** Neural network parameters for classification
- **Training Process:**
 1. Forward pass: Input \rightarrow Embeddings \rightarrow Concatenate \rightarrow MLP \rightarrow Probabilities

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**
 1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
 2. **MLP Weights:** Neural network parameters for classification
- **Training Process:**
 1. Forward pass: Input \rightarrow Embeddings \rightarrow Concatenate \rightarrow MLP \rightarrow Probabilities
 2. Compute cross-entropy loss

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**

1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
2. **MLP Weights:** Neural network parameters for classification

- **Training Process:**

1. Forward pass: Input \rightarrow Embeddings \rightarrow Concatenate \rightarrow MLP \rightarrow Probabilities
2. Compute cross-entropy loss
3. Backward pass: Update both embeddings and MLP weights

Training Objective

- **Loss Function:** Cross-entropy loss for multi-class classification

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^{27} y_{i,c} \log(\hat{y}_{i,c}) \quad (1)$$

- **What we learn:**

1. **Embedding Matrix:** Character representations ($27 \times K$ parameters)
2. **MLP Weights:** Neural network parameters for classification

- **Training Process:**

1. Forward pass: Input \rightarrow Embeddings \rightarrow Concatenate \rightarrow MLP \rightarrow Probabilities
2. Compute cross-entropy loss
3. Backward pass: Update both embeddings and MLP weights
4. Repeat for all training examples

Text Generation

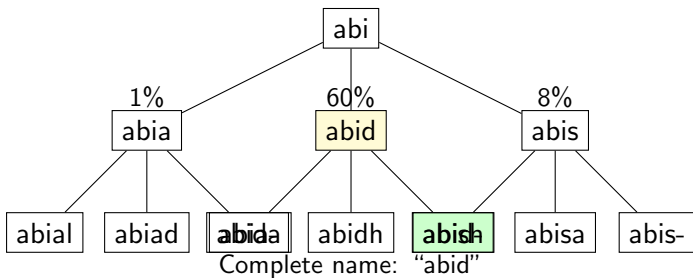
Sampling from the Learned Model

Test Input: "abi"

Predicted Probability Distribution

Next Char	Probability	Next Char	Probability
a	0.01	n	0.05
b	0.01	o	0.02
c	0.03	p	0.01
d	0.60	q	0.00
e	0.02	r	0.03
f	0.01	s	0.08
...
-	0.05	z	0.01

Generation Tree Structure



Recursive Process: Sample next character, append, repeat until end token

Temperature and Sampling Strategies

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- **Temperature Effects:**

Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- **Temperature Effects:**

- $T = 1$: Standard probabilities

Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- **Temperature Effects:**

- $T = 1$: Standard probabilities
- $T \rightarrow 0$: More peaked (deterministic)

Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (2)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (3)$$

- **Temperature Effects:**

- $T = 1$: Standard probabilities
- $T \rightarrow 0$: More peaked (deterministic)
- $T \rightarrow \infty$: More uniform (random)

Temperature Variations

Context: “abi” → Next character probabilities

Char	T=0.5 (Low)	T=1.0 (Default)	T=2.0 (High)
a	0.001	0.01	0.08
d	0.95	0.60	0.25
s	0.01	0.08	0.12
h	0.005	0.03	0.09
-	0.02	0.05	0.11
others	0.015	0.23	0.35

- **Low T:** Conservative, predictable

Temperature Variations

Context: “abi” → Next character probabilities

Char	T=0.5 (Low)	T=1.0 (Default)	T=2.0 (High)
a	0.001	0.01	0.08
d	0.95	0.60	0.25
s	0.01	0.08	0.12
h	0.005	0.03	0.09
-	0.02	0.05	0.11
others	0.015	0.23	0.35

- **Low T:** Conservative, predictable
- **High T:** Creative, diverse

Summary and Applications

Key Takeaways

- **Core Idea:** Next token prediction as classification

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling
- **Training:** Joint learning of embeddings and classifier weights

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling
- **Training:** Joint learning of embeddings and classifier weights
- **Generation:** Autoregressive sampling with temperature control

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling
- **Training:** Joint learning of embeddings and classifier weights
- **Generation:** Autoregressive sampling with temperature control
- **Applications:** Foundation for modern language models

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling
- **Training:** Joint learning of embeddings and classifier weights
- **Generation:** Autoregressive sampling with temperature control
- **Applications:** Foundation for modern language models
 - GPT models use the same principle

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling
- **Training:** Joint learning of embeddings and classifier weights
- **Generation:** Autoregressive sampling with temperature control
- **Applications:** Foundation for modern language models
 - GPT models use the same principle
 - Scaled to words/subwords instead of characters


Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling
- **Training:** Joint learning of embeddings and classifier weights
- **Generation:** Autoregressive sampling with temperature control
- **Applications:** Foundation for modern language models
 - GPT models use the same principle
 - Scaled to words/subwords instead of characters
 - Transformer architecture instead of MLP

Key Takeaways

- **Core Idea:** Next token prediction as classification
- **Representation Learning:** Character embeddings capture similarity
- **Architecture:** Embeddings + MLP for sequence modeling
- **Training:** Joint learning of embeddings and classifier weights
- **Generation:** Autoregressive sampling with temperature control
- **Applications:** Foundation for modern language models
 - GPT models use the same principle
 - Scaled to words/subwords instead of characters
 - Transformer architecture instead of MLP
 - Billions of parameters instead of thousands

From Character-Level to ChatGPT



Same fundamental principle: Predict the next token!