

# Next Token Generation

---

Nipun Batra

IIT Gandhinagar

August 3, 2025

# Inspiration and Relevance

**Inspired by the great lecture from Andrej Karpathy**

- Search for "Neural Networks: Zero to Hero" to find the original lecture

# Inspiration and Relevance

**Inspired by the great lecture from Andrej Karpathy**

- Search for "Neural Networks: Zero to Hero" to find the original lecture

# Inspiration and Relevance

## **Inspired by the great lecture from Andrej Karpathy**

- Search for "Neural Networks: Zero to Hero" to find the original lecture
- This approach is fundamental to modern language models

# Inspiration and Relevance

## **Inspired by the great lecture from Andrej Karpathy**

- Search for "Neural Networks: Zero to Hero" to find the original lecture
- This approach is fundamental to modern language models

# Inspiration and Relevance

**Inspired by the great lecture from Andrej Karpathy**

- Search for "Neural Networks: Zero to Hero" to find the original lecture
- This approach is fundamental to modern language models
- **Direct connection to ChatGPT:**

**Understanding this simple version helps grasp  
ChatGPT's foundation!**

# Inspiration and Relevance

## Inspired by the great lecture from Andrej Karpathy

- Search for "Neural Networks: Zero to Hero" to find the original lecture
- This approach is fundamental to modern language models
- **Direct connection to ChatGPT:**
  - Same core principle: predict the next token

**Understanding this simple version helps grasp  
ChatGPT's foundation!**

# Inspiration and Relevance

## Inspired by the great lecture from Andrej Karpathy

- Search for "Neural Networks: Zero to Hero" to find the original lecture
- This approach is fundamental to modern language models
- **Direct connection to ChatGPT:**
  - Same core principle: predict the next token
  - Scaled up from characters to words/subwords

**Understanding this simple version helps grasp  
ChatGPT's foundation!**



# Inspiration and Relevance

## **Inspired by the great lecture from Andrej Karpathy**

- Search for "Neural Networks: Zero to Hero" to find the original lecture
- This approach is fundamental to modern language models
- **Direct connection to ChatGPT:**
  - Same core principle: predict the next token
  - Scaled up from characters to words/subwords
  - Uses transformer architecture instead of MLP

**Understanding this simple version helps grasp  
ChatGPT's foundation!**

What is the Next Character?

**app**

**?**

**What is the next character?**

# Classification Task

**app**

**?**

**We can pose this as a classification task**

# Classification Task

**app**

**?**

**We can pose this as a classification task**

**Input:**

app

# Classification Task

**app**

**?**

**We can pose this as a classification task**

**Input:**  
app

**Output: Probability Distribution**

Char	Prob	Char	Prob
a	0.01	n	0.01
b	0.01	o	0.01
c	0.01	p	0.01
...	...	...	...
l	<b>0.45</b>	z	0.01
m	0.01	-	0.05

# Generate Indian Names

**Specific Problem: Generate Indian Names**

**Dataset:**

# Generate Indian Names

**Specific Problem: Generate Indian Names**

**Dataset:**

Abid  
Abhidha  
Adesh  
Aditya  
Agam  
...  
:  
Yash  
Yogesh  
Zara

# Generate Indian Names

**Specific Problem: Generate Indian Names**

**Dataset:**

Abid  
Abhidha  
Adesh  
Aditya  
Agam  
...  
:  
Yash  
Yogesh  
Zara



# Generate Indian Names

## Specific Problem: Generate Indian Names

### Dataset:

Abid  
Abhidha  
Adesh  
Aditya  
Agam  
...  
:  
Yash  
Yogesh  
Zara

- Collection of Indian names

# Generate Indian Names

## Specific Problem: Generate Indian Names

### Dataset:

Abid  
Abhidha  
Adesh  
Aditya  
Agam  
...  
:  
Yash  
Yogesh  
Zara

- Collection of Indian names
- Each name represents a sequence

# Generate Indian Names

## Specific Problem: Generate Indian Names

### Dataset:

Abid  
Abhidha  
Adesh  
Aditya  
Agam  
...  
:  
Yash  
Yogesh  
Zara

- Collection of Indian names
- Each name represents a sequence
- Goal: Learn to generate similar names

# Assumptions

**We'll make a few assumptions:**

1. **Character set:** Only use 26 lowercase characters (a-z)

# Assumptions

**We'll make a few assumptions:**

1. **Character set:** Only use 26 lowercase characters (a-z)

# Assumptions

**We'll make a few assumptions:**

1. **Character set:** Only use 26 lowercase characters (a-z)
2. **End marker:** A hyphen (-) indicates the end character

# Assumptions

**We'll make a few assumptions:**

1. **Character set:** Only use 26 lowercase characters (a-z)
2. **End marker:** A hyphen (-) indicates the end character

# Assumptions

**We'll make a few assumptions:**

1. **Character set:** Only use 26 lowercase characters (a-z)
2. **End marker:** A hyphen (-) indicates the end character
3. **Length constraint:** Names are between 4 and 10 characters

**Total vocabulary size:  $26 + 1 = 27$  characters**



# Generate Training Dataset

## Creating Training Data from "abid"

Using history/context of 3 characters:


X (Input)		Y (Target)
[-, -, -]	→	a a>-, a
	→	b a, b>a, b
	→	i a, b, i>a, b, i
	→	d b, i, d>b, i, d
	→	-

**Result: 5 training examples from one name "abid"**

# Representation Learning

## Important Idea: Representation Learning


- Learn a vector representation for each character



# Representation Learning


## Important Idea: Representation Learning

- Learn a vector representation for each character
- Hope that similar characters will be closer in vector space



# Word2Vec Reference


## Classic Word2Vec Relationship



**Relationship:**  $\text{queen} \approx \text{king} - \text{man} + \text{woman}$

# Analogy with Smileys


## Emotional Expression Analogy



**Relationship:** child crying = child smiling + adult crying - adult smiling

# Embedding Matrix/Table

## Main Idea: Embedding Matrix/Table



**Process:** Character → Lookup in Embedding Table → Dense Vector

# 27 × K Embedding Matrix

## Embedding Table Structure

Char	D1	D2	...	DK
a	0.2	-0.1	...	0.8
b	-0.3	0.5	...	-0.2
c	0.1	0.3	...	0.4
⋮	⋮	⋮	⋮	⋮
z	0.7	-0.4	...	0.1
-	0.0	0.9	...	-0.5

**This overall becomes a 27 × K dimensional matrix**

# Learnable Matrix

**This matrix is learnable!**

- **Initially:** Random values



# Learnable Matrix

**This matrix is learnable!**

- **Initially:** Random values

# Learnable Matrix

**This matrix is learnable!**

- **Initially:** Random values
- **During training:** Updated via backpropagation

# Learnable Matrix

**This matrix is learnable!**

- **Initially:** Random values
- **During training:** Updated via backpropagation

# Learnable Matrix

**This matrix is learnable!**

- **Initially:** Random values
- **During training:** Updated via backpropagation
- **After training:** Contains meaningful character representations

# Learnable Matrix

**This matrix is learnable!**

- **Initially:** Random values
- **During training:** Updated via backpropagation
- **After training:** Contains meaningful character representations

# Learnable Matrix

**This matrix is learnable!**

- **Initially:** Random values
- **During training:** Updated via backpropagation
- **After training:** Contains meaningful character representations
- **Similar characters:** Will have similar embedding vectors


**The network learns both the embeddings AND the classification weights!**

# Overall Architecture (2D Example)

Example with  $X = \text{"abi"}$  and 2D embeddings


Input:  $X = [\text{"a"}, \text{"b"}, \text{"i"}]$

Embedding Matrix ( $27 \times 2$ )



# Concatenate the Embeddings

## Feature Vector Creation for X = "abi"




**The feature vector pulls up embeddings and concatenates them**



# Multi-Layer Perceptron

## Neural Architecture



**Eventually shows 27-class output vector**

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**
  1. The embedding matrix ( $27 \times K$  parameters)

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**
  1. The embedding matrix ( $27 \times K$  parameters)
  2. The MLP weights (neural network parameters)

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**
  1. The embedding matrix ( $27 \times K$  parameters)
  2. The MLP weights (neural network parameters)

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**
  1. The embedding matrix ( $27 \times K$  parameters)
  2. The MLP weights (neural network parameters)
- **Training Process:**



# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**
  1. The embedding matrix ( $27 \times K$  parameters)
  2. The MLP weights (neural network parameters)
- **Training Process:**
  - Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**
  1. The embedding matrix ( $27 \times K$  parameters)
  2. The MLP weights (neural network parameters)
- **Training Process:**
  - Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities
  - Compute cross-entropy loss against true next character

# Cross-Entropy Loss

## Learning Process

- **Loss Function:** Use cross-entropy loss to learn
- **We are learning two things:**
  1. The embedding matrix ( $27 \times K$  parameters)
  2. The MLP weights (neural network parameters)
- **Training Process:**
  - Forward pass: Input  $\rightarrow$  Embeddings  $\rightarrow$  Concatenate  $\rightarrow$  MLP  $\rightarrow$  Probabilities
  - Compute cross-entropy loss against true next character
  - Backward pass: Update both embeddings and MLP weights

# Generate/Sample from Learned Model

Test Input: "abi"

Probability vector for next character:

Next Char	Probability	Next Char	Probability
a	0.01	n	0.05
b	0.01	o	0.02
c	0.03	p	0.01
d	<b>0.60</b>	...	...
...	...	z	0.01

- ABIA would be 1%

# Generate/Sample from Learned Model

Test Input: "abi"

Probability vector for next character:

Next Char	Probability	Next Char	Probability
a	0.01	n	0.05
b	0.01	o	0.02
c	0.03	p	0.01
d	<b>0.60</b>	...	...
...	...	z	0.01

- ABIA would be 1%
- ABIB would be 1%

# Generate/Sample from Learned Model

Test Input: "abi"


Probability vector for next character:

Next Char	Probability	Next Char	Probability
a	0.01	n	0.05
b	0.01	o	0.02
c	0.03	p	0.01
d	<b>0.60</b>	...	...
...	...	z	0.01

- ABIA would be 1%
- ABIB would be 1%
- **ABID would be 60%**

# Tree Structure

## Generation as Tree Structure



Had we chosen A, it starts a new branch. Had we chosen D, it starts a new branch, etc.

# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$



# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$

# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (2)$$

# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (2)$$

# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (2)$$

- **Temperature Effects:**

# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (2)$$

- **Temperature Effects:**

- $T = 1$ : Default/standard probabilities

# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (2)$$

- **Temperature Effects:**

- $T = 1$ : Default/standard probabilities
- $T \rightarrow 0$ : Very low temperature  $\rightarrow$  more peaked (deterministic)

# Temperature Term

## Temperature in Softmax

- **Standard Softmax:**

$$P(y_i) = \frac{e^{z_i}}{\sum_{j=1}^{27} e^{z_j}} \quad (1)$$

- **Temperature-scaled Softmax:**

$$P(y_i) = \frac{e^{z_i/T}}{\sum_{j=1}^{27} e^{z_j/T}} \quad (2)$$

- **Temperature Effects:**

- $T = 1$ : Default/standard probabilities
- $T \rightarrow 0$ : Very low temperature  $\rightarrow$  more peaked (deterministic)
- $T \rightarrow \infty$ : Very high temperature  $\rightarrow$  more uniform (random)

# Temperature Variations

## How sampling differs across temperatures

Next Char	Default T=1.0	Low T T=0.5	High T T=2.0
a	0.01	0.001	0.08
d	<b>0.60</b>	<b>0.95</b>	<b>0.25</b>
s	0.08	0.01	0.12
h	0.03	0.005	0.09
-	0.05	0.02	0.11
others	0.23	0.015	0.35

- **Low Temperature:** Conservative, predictable generation



# Temperature Variations

## How sampling differs across temperatures

Next Char	Default T=1.0	Low T T=0.5	High T T=2.0
a	0.01	0.001	0.08
d	<b>0.60</b>	<b>0.95</b>	<b>0.25</b>
s	0.08	0.01	0.12
h	0.03	0.005	0.09
-	0.05	0.02	0.11
others	0.23	0.015	0.35

- **Low Temperature:** Conservative, predictable generation
- **High Temperature:** Creative, diverse generation