

# Gradient Descent

---

Nipun Batra

July 25, 2025

IIT Gandhinagar


# Revision

---


# Contour Plot And Gradients

$$z = f(x, y) = x^2 + y^2$$

Surface Plot




Contour Plot




# Contour Plot And Gradients

$$z = f(x, y) = x^2 + y^2$$

Surface Plot



Contour Plot




Gradient denotes the direction of steepest ascent or the direction in which there is a maximum increase in  $f(x,y)$


## Contour Plot And Gradients

$$z = f(x, y) = x^2 + y^2$$

Surface Plot



Contour Plot



Gradient denotes the direction of steepest ascent or the direction in which there is a maximum increase in  $f(x,y)$

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x,y)}{\partial x} \\ \frac{\partial f(x,y)}{\partial y} \end{bmatrix} = \begin{bmatrix} 2x \\ 2y \end{bmatrix}$$

# Introduction

---

# Optimization algorithms

- We often want to minimize/maximize a function

- We often want to minimize/maximize a function
- We wanted to minimize the cost function:

$$f(\theta) = (\mathbf{y} - \mathbf{X}\theta)^T (\mathbf{y} - \mathbf{X}\theta) \quad (1)$$



# Optimization algorithms

- We often want to minimize/maximize a function
- We wanted to minimize the cost function:

$$f(\theta) = (\mathbf{y} - \mathbf{X}\theta)^T(\mathbf{y} - \mathbf{X}\theta) \quad (1)$$

- Note, here  $\theta$  is the parameter vector

# Optimization algorithms

- In general, we have following components:

# Optimization algorithms

- In general, we have following components:
- Maximize or Minimize a function subject to some constraints

# Optimization algorithms

- In general, we have following components:
- Maximize or Minimize a function subject to some constraints
- Today, we will focus on unconstrained optimization (no constraints)

# Optimization algorithms

- In general, we have following components:
- Maximize or Minimize a function subject to some constraints
- Today, we will focus on unconstrained optimization (no constraints)
- We will focus on minimization

# Optimization algorithms

- In general, we have following components:
- Maximize or Minimize a function subject to some constraints
- Today, we will focus on unconstrained optimization (no constraints)
- We will focus on minimization
- Goal:

$$\theta^* = \arg \min_{\theta} f(\theta) \quad (2)$$

- Gradient descent is an optimization algorithm

- Gradient descent is an optimization algorithm
- It is used to find the minimum of a function in unconstrained settings



- Gradient descent is an optimization algorithm
- It is used to find the minimum of a function in unconstrained settings
- It is an iterative algorithm

- Gradient descent is an optimization algorithm
- It is used to find the minimum of a function in unconstrained settings
- It is an iterative algorithm
- It is a first order optimization algorithm

# Introduction

- Gradient descent is an optimization algorithm
- It is used to find the minimum of a function in unconstrained settings
- It is an iterative algorithm
- It is a first order optimization algorithm
- It is a local search algorithm/greedy

# Gradient Descent Algorithm

1. Initialize  $\theta$  to some random value

# Gradient Descent Algorithm

1. Initialize  $\theta$  to some random value
2. Compute the gradient of the cost function at  $\theta$ ,  $\nabla f(\theta)$

# Gradient Descent Algorithm

1. Initialize  $\theta$  to some random value
2. Compute the gradient of the cost function at  $\theta$ ,  $\nabla f(\theta)$
3. For Iteration  $i$  ( $i = 1, 2, \dots$ ) or until convergence:

# Gradient Descent Algorithm

1. Initialize  $\theta$  to some random value
2. Compute the gradient of the cost function at  $\theta$ ,  $\nabla f(\theta)$
3. For Iteration  $i$  ( $i = 1, 2, \dots$ ) or until convergence:
  - $\theta_i \leftarrow \theta_{i-1} - \alpha \nabla f(\theta_{i-1})$

# Taylor's Series

---



## Taylor's Series

- Taylor's series is a way to approximate a function  $f(x)$  around a point  $x_0$  using a polynomial

## Taylor's Series

- Taylor's series is a way to approximate a function  $f(x)$  around a point  $x_0$  using a polynomial
- The polynomial is given by

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (3)$$

## Taylor's Series

- Taylor's series is a way to approximate a function  $f(x)$  around a point  $x_0$  using a polynomial
- The polynomial is given by

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (3)$$

- The vector form of the above equation is given by:

$$f(\mathbf{x}) = f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T (\mathbf{x} - \mathbf{x}_0) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_0)^T \nabla^2 f(\mathbf{x}_0) (\mathbf{x} - \mathbf{x}_0) + \dots \quad (4)$$

## Taylor's Series

- Taylor's series is a way to approximate a function  $f(x)$  around a point  $x_0$  using a polynomial
- The polynomial is given by

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (3)$$

- The vector form of the above equation is given by:

$$f(\mathbf{x}) = f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T (\mathbf{x} - \mathbf{x}_0) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_0)^T \nabla^2 f(\mathbf{x}_0) (\mathbf{x} - \mathbf{x}_0) + \dots \quad (4)$$

- where  $\nabla^2 f(\mathbf{x}_0)$  is the Hessian matrix and  $\nabla f(\mathbf{x}_0)$  is the gradient vector

## Taylor's Series

- Let us consider  $f(x) = \cos(x)$  and  $x_0 = 0$

## Taylor's Series

- Let us consider  $f(x) = \cos(x)$  and  $x_0 = 0$
- Then, we have:

## Taylor's Series

- Let us consider  $f(x) = \cos(x)$  and  $x_0 = 0$
- Then, we have:
- $f(x_0) = \cos(0) = 1$

## Taylor's Series

- Let us consider  $f(x) = \cos(x)$  and  $x_0 = 0$
- Then, we have:
- $f(x_0) = \cos(0) = 1$
- $f'(x_0) = -\sin(0) = 0$



## Taylor's Series

- Let us consider  $f(x) = \cos(x)$  and  $x_0 = 0$
- Then, we have:
- $f(x_0) = \cos(0) = 1$
- $f'(x_0) = -\sin(0) = 0$
- $f''(x_0) = -\cos(0) = -1$

## Taylor's Series

- Let us consider  $f(x) = \cos(x)$  and  $x_0 = 0$
- Then, we have:
- $f(x_0) = \cos(0) = 1$
- $f'(x_0) = -\sin(0) = 0$
- $f''(x_0) = -\cos(0) = -1$
- We can write the second order Taylor's series as:

## Taylor's Series

- Let us consider  $f(x) = \cos(x)$  and  $x_0 = 0$
- Then, we have:
- $f(x_0) = \cos(0) = 1$
- $f'(x_0) = -\sin(0) = 0$
- $f''(x_0) = -\cos(0) = -1$
- We can write the second order Taylor's series as:
- $f(x) = 1 + 0(x - 0) + \frac{-1}{2!}(x - 0)^2 = 1 - \frac{x^2}{2}$

## Taylor's series

- Let us consider another example:  $f(x) = x^2 + 2$  and  $x_0 = 2$

## Taylor's series

- Let us consider another example:  $f(x) = x^2 + 2$  and  $x_0 = 2$
- Question: How does the first order Taylor's series approximation look like?

## Taylor's series

- Let us consider another example:  $f(x) = x^2 + 2$  and  $x_0 = 2$
- Question: How does the first order Taylor's series approximation look like?
- First order Taylor's series approximation is given by:

## Taylor's series

- Let us consider another example:  $f(x) = x^2 + 2$  and  $x_0 = 2$
- Question: How does the first order Taylor's series approximation look like?
- First order Taylor's series approximation is given by:
- $f(x) = f(x_0) + f'(x_0)(x - x_0) = 6 + 4(x - 2) = 4x - 2$

## Taylor's Series (Alternative form)

- We have:

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (5)$$



## Taylor's Series (Alternative form)

- We have:

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (5)$$

- Let us consider  $x = x_0 + \Delta x$  where  $\Delta x$  is a small quantity

## Taylor's Series (Alternative form)

- We have:

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (5)$$

- Let us consider  $x = x_0 + \Delta x$  where  $\Delta x$  is a small quantity
- Then, we have:

$$f(x_0 + \Delta x) = f(x_0) + \frac{f'(x_0)}{1!}\Delta x + \frac{f''(x_0)}{2!}\Delta x^2 + \dots \quad (6)$$

## Taylor's Series (Alternative form)

- We have:

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (5)$$

- Let us consider  $x = x_0 + \Delta x$  where  $\Delta x$  is a small quantity
- Then, we have:

$$f(x_0 + \Delta x) = f(x_0) + \frac{f'(x_0)}{1!}\Delta x + \frac{f''(x_0)}{2!}\Delta x^2 + \dots \quad (6)$$

- Let us assume  $\Delta x$  is small enough such that  $\Delta x^2$  and higher order terms can be ignored

## Taylor's Series (Alternative form)

- We have:

$$f(x) = f(x_0) + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots \quad (5)$$

- Let us consider  $x = x_0 + \Delta x$  where  $\Delta x$  is a small quantity
- Then, we have:

$$f(x_0 + \Delta x) = f(x_0) + \frac{f'(x_0)}{1!}\Delta x + \frac{f''(x_0)}{2!}\Delta x^2 + \dots \quad (6)$$

- Let us assume  $\Delta x$  is small enough such that  $\Delta x^2$  and higher order terms can be ignored
- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!}\Delta x$

## Taylor's Series to Gradient Descent

- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!} \Delta x$

## Taylor's Series to Gradient Descent

- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!} \Delta x$
- Or, in vector form:  $f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$

## Taylor's Series to Gradient Descent

- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!} \Delta x$
- Or, in vector form:  $f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- Goal: Find  $\Delta \mathbf{x}$  such that  $f(\mathbf{x}_0 + \Delta \mathbf{x})$  is minimized

## Taylor's Series to Gradient Descent

- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!} \Delta x$
- Or, in vector form:  $f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- Goal: Find  $\Delta \mathbf{x}$  such that  $f(\mathbf{x}_0 + \Delta \mathbf{x})$  is minimized
- This is equivalent to minimizing  $f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$



## Taylor's Series to Gradient Descent

- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!} \Delta x$
- Or, in vector form:  $f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- Goal: Find  $\Delta \mathbf{x}$  such that  $f(\mathbf{x}_0 + \Delta \mathbf{x})$  is minimized
- This is equivalent to minimizing  $f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- This happens when vectors  $\nabla f(\mathbf{x}_0)$  and  $\Delta \mathbf{x}$  are at phase angle of  $180^\circ$

## Taylor's Series to Gradient Descent


- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!} \Delta x$
- Or, in vector form:  $f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- Goal: Find  $\Delta \mathbf{x}$  such that  $f(\mathbf{x}_0 + \Delta \mathbf{x})$  is minimized
- This is equivalent to minimizing  $f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- This happens when vectors  $\nabla f(\mathbf{x}_0)$  and  $\Delta \mathbf{x}$  are at phase angle of  $180^\circ$
- This happens when  $\Delta \mathbf{x} = -\alpha \nabla f(\mathbf{x}_0)$  where  $\alpha$  is a scalar

## Taylor's Series to Gradient Descent

- Then, we have:  $f(x_0 + \Delta x) \approx f(x_0) + \frac{f'(x_0)}{1!} \Delta x$
- Or, in vector form:  $f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- Goal: Find  $\Delta \mathbf{x}$  such that  $f(\mathbf{x}_0 + \Delta \mathbf{x})$  is minimized
- This is equivalent to minimizing  $f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$
- This happens when vectors  $\nabla f(\mathbf{x}_0)$  and  $\Delta \mathbf{x}$  are at phase angle of  $180^\circ$
- This happens when  $\Delta \mathbf{x} = -\alpha \nabla f(\mathbf{x}_0)$  where  $\alpha$  is a scalar
- This is the gradient descent algorithm:  $\mathbf{x}_1 = \mathbf{x}_0 - \alpha \nabla f(\mathbf{x}_0)$


## Effect of learning rate

Low learning rate  $\alpha = 0.01$  : Converges slowly




## Effect of learning rate

High learning rate  $\alpha = 0.8$ : Converges quickly, but might overshoot




## Effect of learning rate

Very high learning rate  $\alpha = 1.01$ : Diverges



## Effect of learning rate

Appropriate learning rate  $\alpha = 0.1$



## Gradient Descent for linear regression

---



## Some commonly confused terms

- **Loss function** is usually a function defined on a data point, prediction and label, and measures the penalty.

## Some commonly confused terms

- **Loss function** is usually a function defined on a data point, prediction and label, and measures the penalty.
- square loss  $l(f(x_i; \theta), y_i) = (f(x_i; \theta) - y_i)^2$ , used in linear regression

## Some commonly confused terms

- **Loss function** is usually a function defined on a data point, prediction and label, and measures the penalty.
- square loss  $l(f(x_i; \theta), y_i) = (f(x_i; \theta) - y_i)^2$ , used in linear regression
- **Cost function** is usually more general. It might be a sum of loss functions over your training set plus some model complexity penalty (regularization). For example:

## Some commonly confused terms

- **Loss function** is usually a function defined on a data point, prediction and label, and measures the penalty.
- square loss  $l(f(x_i; \theta), y_i) = (f(x_i; \theta) - y_i)^2$ , used in linear regression
- **Cost function** is usually more general. It might be a sum of loss functions over your training set plus some model complexity penalty (regularization). For example:
- Mean Squared Error  $\text{MSE}(\theta) = \frac{1}{n} \sum_{i=1}^n (f(x_i; \theta) - y_i)^2$

## Some commonly confused terms

- **Loss function** is usually a function defined on a data point, prediction and label, and measures the penalty.
- square loss  $l(f(x_i; \theta), y_i) = (f(x_i; \theta) - y_i)^2$ , used in linear regression
- **Cost function** is usually more general. It might be a sum of loss functions over your training set plus some model complexity penalty (regularization). For example:
- Mean Squared Error  $\text{MSE}(\theta) = \frac{1}{n} \sum_{i=1}^n (f(x_i; \theta) - y_i)^2$
- **Objective function** is the most general term for any function that you optimize during training.

## Gradient Descent : Example

Learn  $y = \theta_0 + \theta_1 x$  on following dataset, using gradient descent where initially  $(\theta_0, \theta_1) = (4, 0)$  and step-size,  $\alpha = 0.1$ , for 2 iterations.

<b>x</b>	<b>y</b>
1	1
2	2
3	3

## Gradient Descent : Example

Our predictor,  $\hat{y} = \theta_0 + \theta_1 x$

Error for  $i^{th}$  datapoint,  $\epsilon_i = y_i - \hat{y}_i$

$$\epsilon_1 = 1 - \theta_0 - \theta_1$$

$$\epsilon_2 = 2 - \theta_0 - 2\theta_1$$

$$\epsilon_3 = 3 - \theta_0 - 3\theta_1$$

$$\text{MSE} = \frac{\epsilon_1^2 + \epsilon_2^2 + \epsilon_3^2}{3} = \frac{14 + 3\theta_0^2 + 14\theta_1^2 - 12\theta_0 - 28\theta_1 + 12\theta_0\theta_1}{3}$$

## Difference between SSE and MSE

$\sum \epsilon_i^2$  increases as the number of examples increase

So, we use MSE

$$MSE = \frac{1}{n} \sum \epsilon_i^2$$

Here  $n$  denotes the number of samples



## Gradient Descent : Example

$$\frac{\partial \text{MSE}}{\partial \theta_0} = \frac{2 \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)(-1)}{n} = \frac{2 \sum_{i=1}^n \epsilon_i(-1)}{n}$$

$$\frac{\partial \text{MSE}}{\partial \theta_1} = \frac{2 \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)(-x_i)}{n} = \frac{2 \sum_{i=1}^n \epsilon_i(-x_i)}{n}$$

# Gradient Descent : Example

## Iteration 1

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

## Gradient Descent : Example

### Iteration 1

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 4 - 0.2 \frac{((1-(4+0))(-1) + (2-(4+0))(-1) + (3-(4+0))(-1))}{3}$$

$$\theta_0 = 3.6$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

# Gradient Descent : Example

## Iteration 1

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 4 - 0.2 \frac{((1-(4+0))(-1) + (2-(4+0))(-1) + (3-(4+0))(-1))}{3}$$

$$\theta_0 = 3.6$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

$$\theta_1 = 0 - 0.2 \frac{((1-(4+0))(-1) + (2-(4+0))(-2) + (3-(4+0))(-3))}{3}$$

$$\theta_1 = -0.67$$

# Gradient Descent : Example

## Iteration 2

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

# Gradient Descent : Example

## Iteration 2

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 3.6 - 0.2 \frac{((1 - (3.6 - 0.67))(-1) + (2 - (3.6 - 0.67 \times 2))(-1) + (3 - (3.6 - 0.67 \times 3))(-1))}{3}$$

$$\theta_0 = 3.54$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

# Gradient Descent : Example

## Iteration 2

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 =$$

$$3.6 - 0.2 \frac{((1-(3.6-0.67))(-1)+(2-(3.6-0.67 \times 2))(-1)+(3-(3.6-0.67 \times 3))(-1))}{3}$$

$$\theta_0 = 3.54$$



$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

$$\theta_0 =$$

$$3.6 - 0.2 \frac{((1-(3.6-0.67))(-1)+(2-(3.6-0.67 \times 2))(-2)+(3-(3.6-0.67 \times 3))(-3))}{3}$$



$$\theta_0 = -0.55$$

# Gradient Descent : Example (Iteration 0)









## Gradient Descent : Example (Iteration 2)





## Gradient Descent : Example (Iteration 4)





## Gradient Descent : Example (Iteration 6)





## Gradient Descent : Example (Iteration 8)





# Gradient Descent : Example (Iteration 10)





## Gradient Descent : Example (Iteration 12)



# Gradient Descent : Example (Iteration 14)


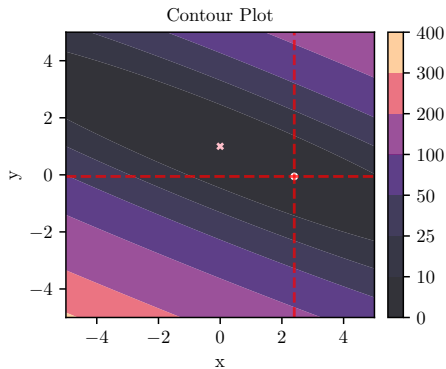


## Gradient Descent : Example (Iteration 16)








## Gradient Descent : Example (Iteration 18)




## Gradient Descent : Example (Iteration 20)





## Gradient Descent : Example (Iteration 22)





## Gradient Descent : Example (Iteration 24)





## Gradient Descent : Example (Iteration 26)





## Gradient Descent : Example (Iteration 28)



## Gradient Descent : Example (Iteration 30)





## Gradient Descent : Example (Iteration 32)









## Gradient Descent : Example (Iteration 34)





## Gradient Descent : Example (Iteration 36)



## Gradient Descent : Example (Iteration 38)



## Gradient Descent : Example (Iteration 40)



## Iteration vs Epochs for gradient descent

- Iteration: Each time you update the parameters of the model

## Iteration vs Epochs for gradient descent

- Iteration: Each time you update the parameters of the model
- Epoch: Each time you have seen all the set of examples

## Gradient Descent (GD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$

## Gradient Descent (GD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$



# Gradient Descent (GD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$

# Gradient Descent (GD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Predict  $\hat{\mathbf{y}} = \text{pred}(\mathbf{X}, \theta)$

## Gradient Descent (GD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Predict  $\hat{\mathbf{y}} = \text{pred}(\mathbf{X}, \theta)$
  - Compute loss:  $J(\theta) = \text{loss}(\mathbf{y}, \hat{\mathbf{y}})$

## Gradient Descent (GD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Predict  $\hat{\mathbf{y}} = \text{pred}(\mathbf{X}, \theta)$
  - Compute loss:  $J(\theta) = \text{loss}(\mathbf{y}, \hat{\mathbf{y}})$
  - Compute gradient:  $\nabla J(\theta) = \text{grad}(J)(\theta)$

# Gradient Descent (GD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Predict  $\hat{\mathbf{y}} = \text{pred}(\mathbf{X}, \theta)$
  - Compute loss:  $J(\theta) = \text{loss}(\mathbf{y}, \hat{\mathbf{y}})$
  - Compute gradient:  $\nabla J(\theta) = \text{grad}(J)(\theta)$
  - Update:  $\theta = \theta - \alpha \nabla J(\theta)$

## Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$

## Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$

# Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$



## Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$

# Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - For  $i$  in  $[1, n]$

# Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - For  $i$  in  $[1, n]$ 
    - Predict  $\hat{\mathbf{y}}_i = \text{pred}(\mathbf{x}_i, \theta)$

# Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - For  $i$  in  $[1, n]$ 
    - Predict  $\hat{\mathbf{y}}_i = \text{pred}(\mathbf{x}_i, \theta)$
    - Compute loss:  $J(\theta) = \text{loss}(y_i, \hat{\mathbf{y}}_i)$

# Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - For  $i$  in  $[1, n]$ 
    - Predict  $\hat{\mathbf{y}}_i = \text{pred}(\mathbf{x}_i, \theta)$
    - Compute loss:  $J(\theta) = \text{loss}(y_i, \hat{\mathbf{y}}_i)$
    - Compute gradient:  $\nabla J(\theta) = \text{grad}(J)(\theta)$

# Stochastic Gradient Descent (SGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - For  $i$  in  $[1, n]$ 
    - Predict  $\hat{\mathbf{y}}_i = \text{pred}(\mathbf{x}_i, \theta)$
    - Compute loss:  $J(\theta) = \text{loss}(y_i, \hat{\mathbf{y}}_i)$
    - Compute gradient:  $\nabla J(\theta) = \text{grad}(J)(\theta)$
    - Update:  $\theta = \theta - \alpha \nabla J(\theta)$

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$



## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - Batches = `make_batches( $\mathcal{D}$ ,  $B$ )`

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - Batches =  $\text{make\_batches}(\mathcal{D}, B)$
  - For  $b$  in Batches

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - Batches =  $\text{make\_batches}(\mathcal{D}, B)$
  - For  $b$  in Batches
    - $\mathbf{X}_b, \mathbf{y}_b = b$

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - Batches =  $\text{make\_batches}(\mathcal{D}, B)$
  - For  $b$  in Batches
    - $\mathbf{X}_b, \mathbf{y}_b = b$
    - Predict  $\hat{\mathbf{y}}_b = \text{pred}(\mathbf{X}_b, \theta)$

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - Batches =  $\text{make\_batches}(\mathcal{D}, B)$
  - For  $b$  in Batches
    - $\mathbf{X}_b, \mathbf{y}_b = b$
    - Predict  $\hat{\mathbf{y}}_b = \text{pred}(\mathbf{X}_b, \theta)$
    - Compute loss:  $J(\theta) = \text{loss}(\mathbf{y}_b, \hat{\mathbf{y}}_b)$

## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - Batches = `make_batches( $\mathcal{D}, B$ )`
  - For  $b$  in Batches
    - $\mathbf{X}_b, \mathbf{y}_b = b$
    - Predict  $\hat{\mathbf{y}}_b = \text{pred}(\mathbf{X}_b, \theta)$
    - Compute loss:  $J(\theta) = \text{loss}(\mathbf{y}_b, \hat{\mathbf{y}}_b)$
    - Compute gradient:  $\nabla J(\theta) = \text{grad}(J)(\theta)$



## Mini-Batch Gradient Descent (MBGD)

- Dataset:  $\mathcal{D} = \{(\mathbf{X}, \mathbf{y})\}$  of size  $n$
- Initialize  $\theta$
- For epoch  $e$  in  $[1, E]$ 
  - Shuffle  $\mathcal{D}$
  - Batches = make\_batches( $\mathcal{D}, B$ )
  - For  $b$  in Batches
    - $\mathbf{X}_b, \mathbf{y}_b = b$
    - Predict  $\hat{\mathbf{y}}_b = \text{pred}(\mathbf{X}_b, \theta)$
    - Compute loss:  $J(\theta) = \text{loss}(\mathbf{y}_b, \hat{\mathbf{y}}_b)$
    - Compute gradient:  $\nabla J(\theta) = \text{grad}(J)(\theta)$
    - Update:  $\theta = \theta - \alpha \nabla J(\theta)$

# Gradient Descent vs SGD

## Vanilla Gradient Descent

- in Vanilla (Batch) gradient descent: We update params after going through all the data

# Gradient Descent vs SGD

## Vanilla Gradient Descent

- in Vanilla (Batch) gradient descent: We update params after going through all the data
- Smooth curve for Iteration vs Cost

# Gradient Descent vs SGD

## Vanilla Gradient Descent

- in Vanilla (Batch) gradient descent: We update params after going through all the data
- Smooth curve for Iteration vs Cost
- For a single update, it needs to compute the gradient over all the samples. Hence takes more time

# Gradient Descent vs SGD

## Vanilla Gradient Descent

- in Vanilla (Batch) gradient descent: We update params after going through all the data
- Smooth curve for Iteration vs Cost
- For a single update, it needs to compute the gradient over all the samples. Hence takes more time

# Gradient Descent vs SGD

## Vanilla Gradient Descent

- in Vanilla (Batch) gradient descent: We update params after going through all the data
- Smooth curve for Iteration vs Cost
- For a single update, it needs to compute the gradient over all the samples. Hence takes more time

## Stochastic Gradient Descent

- In SGD, we update parameters after seeing each each point

# Gradient Descent vs SGD

## Vanilla Gradient Descent

- in Vanilla (Batch) gradient descent: We update params after going through all the data
- Smooth curve for Iteration vs Cost
- For a single update, it needs to compute the gradient over all the samples. Hence takes more time

## Stochastic Gradient Descent

- In SGD, we update parameters after seeing each each point
- Noisier curve for iteration vs cost

# Gradient Descent vs SGD

## Vanilla Gradient Descent

- in Vanilla (Batch) gradient descent: We update params after going through all the data
- Smooth curve for Iteration vs Cost
- For a single update, it needs to compute the gradient over all the samples. Hence takes more time

## Stochastic Gradient Descent

- In SGD, we update parameters after seeing each each point
- Noisier curve for iteration vs cost
- For a single update, it computes the gradient over one example. Hence lesser time



## Stochastic Gradient Descent : Example

Learn  $y = \theta_0 + \theta_1 x$  on following dataset, using SGD where initially  $(\theta_0, \theta_1) = (4, 0)$  and step-size,  $\alpha = 0.1$ , for 1 epoch (3 iterations).

<b>x</b>	<b>y</b>
2	2
3	3
1	1

## Stochastic Gradient Descent : Example

Our predictor,  $\hat{y} = \theta_0 + \theta_1 x$

Error for  $i^{th}$  datapoint,  $e_i = y_i - \hat{y}_i$

$$\epsilon_1 = 2 - \theta_0 - 2\theta_1$$

$$\epsilon_2 = 3 - \theta_0 - 3\theta_1$$


$$\epsilon_3 = 1 - \theta_0 - \theta_1$$

While using SGD, we compute the MSE using only 1 datapoint per iteration.

So MSE is  $\epsilon_1^2$  for iteration 1 and  $\epsilon_2^2$  for iteration 2.

# Stochastic Gradient Descent : Example

Contour plot of the cost functions for the three datapoints



**For Iteration  $i$**

$$\frac{\partial MSE}{\partial \theta_0} = 2 (y_i - \theta_0 - \theta_1 x_i) (-1) = 2\epsilon_i (-1)$$

$$\frac{\partial MSE}{\partial \theta_1} = 2 (y_i - \theta_0 - \theta_1 x_i) (-x_i) = 2\epsilon_i (-x_i)$$

# Stochastic Gradient Descent : Example

## Iteration 1

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

# Stochastic Gradient Descent : Example

## Iteration 1

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 4 - 0.1 \times 2 \times (2 - (4 + 0))(-1)$$

$$\theta_0 = 3.6$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

# Stochastic Gradient Descent : Example

## Iteration 1

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 4 - 0.1 \times 2 \times (2 - (4 + 0))(-1)$$

$$\theta_0 = 3.6$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

$$\theta_1 = 0 - 0.1 \times 2 \times (2 - (4 + 0))(-2)$$

$$\theta_1 = -0.8$$

# Stochastic Gradient Descent : Example

## Iteration 2

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$



## Iteration 2

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 3.6 - 0.1 \times 2 \times (3 - (3.6 - 0.8 \times 3))(-1)$$

$$\theta_0 = 3.96$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

## Stochastic Gradient Descent : Example

### Iteration 2

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 3.6 - 0.1 \times 2 \times (3 - (3.6 - 0.8 \times 3))(-1)$$

$$\theta_0 = 3.96$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

$$\theta_1 = -0.8 - 0.1 \times 2 \times (3 - (3.6 - 0.8 \times 3))(-3)$$

$$\theta_1 = 0.28$$

### Iteration 3

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

### Iteration 3

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 3.96 - 0.1 \times 2 \times (1 - (3.96 + 0.28 \times 1))(-1)$$

$$\theta_0 = 3.312$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

## Stochastic Gradient Descent : Example

### Iteration 3

$$\theta_0 = \theta_0 - \alpha \frac{\partial MSE}{\partial \theta_0}$$

$$\theta_0 = 3.96 - 0.1 \times 2 \times (1 - (3.96 + 0.28 \times 1))(-1)$$

$$\theta_0 = 3.312$$

$$\theta_1 = \theta_1 - \alpha \frac{\partial MSE}{\partial \theta_1}$$

$$\theta_1 = 0.28 - 0.1 \times 2 \times (1 - (3.96 + 0.28 \times 1))(-1)$$

$$\theta_1 = -0.368$$

**Stochastic gradient is an unbiased estimator of the true gradient**

---

Based on Estimation Theory and Machine Learning by Florian Hartmann

- Let us say we have a dataset  $\mathcal{D}$  containing input output pairs  $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$

Based on Estimation Theory and Machine Learning by Florian Hartmann

- Let us say we have a dataset  $\mathcal{D}$  containing input output pairs  $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- We can define overall loss as:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \text{loss}(f(x_i, \theta), y_i)$$



Based on Estimation Theory and Machine Learning by Florian Hartmann

- Let us say we have a dataset  $\mathcal{D}$  containing input output pairs  $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- We can define overall loss as:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \text{loss}(f(x_i, \theta), y_i)$$

- loss can be any loss function such as squared loss, cross-entropy loss etc.

$$\text{loss}(f(x_i, \theta), y_i) = (f(x_i, \theta) - y_i)^2$$

- The true gradient of the loss function is given by:

$$\begin{aligned}\nabla L &= \nabla \frac{1}{n} \sum_{i=1}^n \text{loss}(f(x_i), y_i) \\ &= \frac{1}{n} \sum_{i=1}^n \nabla \text{loss}(f(x_i), y_i)\end{aligned}$$

## True Gradient

- The true gradient of the loss function is given by:

$$\begin{aligned}\nabla L &= \nabla \frac{1}{n} \sum_{i=1}^n \text{loss}(f(x_i), y_i) \\ &= \frac{1}{n} \sum_{i=1}^n \nabla \text{loss}(f(x_i), y_i)\end{aligned}$$

- The above is a consequence of linearity of the gradient operator.

## Estimator for the true gradient

- In practice, we do not have access to the true gradient

## Estimator for the true gradient

- In practice, we do not have access to the true gradient
- We can only estimate the true gradient using a subset of the data

## Estimator for the true gradient

- In practice, we do not have access to the true gradient
- We can only estimate the true gradient using a subset of the data
- For SGD, we use a single example to estimate the true gradient, for mini-batch gradient descent, we use a mini-batch of examples to estimate the true gradient

## Estimator for the true gradient

- In practice, we do not have access to the true gradient
- We can only estimate the true gradient using a subset of the data
- For SGD, we use a single example to estimate the true gradient, for mini-batch gradient descent, we use a mini-batch of examples to estimate the true gradient
- Let us say we have a sample:  $(x, y)$

## Estimator for the true gradient

- In practice, we do not have access to the true gradient
- We can only estimate the true gradient using a subset of the data
- For SGD, we use a single example to estimate the true gradient, for mini-batch gradient descent, we use a mini-batch of examples to estimate the true gradient
- Let us say we have a sample:  $(x, y)$
- The estimated gradient is given by:

$$\nabla \tilde{L} = \nabla \text{loss}(f(x), y)$$



## Bias of the estimator

- One measure for the quality of an estimator  $\tilde{X}$  is its bias or how far off its estimate is on average from the true value  $X$  :

$$\text{bias}(X) = \mathbb{E}[\tilde{X}] - X$$

## Bias of the estimator

- One measure for the quality of an estimator  $\tilde{X}$  is its bias or how far off its estimate is on average from the true value  $X$  :

$$\text{bias}(X) = \mathbb{E}[\tilde{X}] - X$$

- Using the rules of expectation, we can show that the expected value of the estimated gradient is the true gradient:

$$\begin{aligned}\mathbb{E}[\nabla \tilde{L}] &= \sum_{i=1}^n \frac{1}{n} \nabla \text{loss}(f(x_i), y_i) \\ &= \frac{1}{n} \nabla \sum_{i=1}^n \text{loss}(f(x_i), y_i) \\ &= \nabla L\end{aligned}$$

## Bias of the estimator

- One measure for the quality of an estimator  $\tilde{X}$  is its bias or how far off its estimate is on average from the true value  $X$  :

$$\text{bias}(X) = \mathbb{E}[\tilde{X}] - X$$

- Using the rules of expectation, we can show that the expected value of the estimated gradient is the true gradient:

$$\begin{aligned}\mathbb{E}[\nabla \tilde{L}] &= \sum_{i=1}^n \frac{1}{n} \nabla \text{loss}(f(x_i), y_i) \\ &= \frac{1}{n} \nabla \sum_{i=1}^n \text{loss}(f(x_i), y_i) \\ &= \nabla L\end{aligned}$$

- Thus, the estimated gradient is an unbiased estimator of the true gradient

## **Time Complexity: Gradient Descent vs Normal Equation for Linear Regression**

---

# Normal Equation

- Consider  $\mathbf{X} \in \mathbb{R}^{n \times d}$

# Normal Equation

- Consider  $\mathbf{X} \in \mathbb{R}^{n \times d}$
- $n$  examples and  $d$  dimensions

# Normal Equation

- Consider  $\mathbf{X} \in \mathbb{R}^{n \times d}$
- $n$  examples and  $d$  dimensions
- What is the time complexity of solving the normal equation  $\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$ ?

## Normal Equation

- $\mathbf{X}$  has dimensions  $n \times d$ ,  $\mathbf{X}^T$  has dimensions  $d \times n$



## Normal Equation

- $\mathbf{X}$  has dimensions  $n \times d$ ,  $\mathbf{X}^T$  has dimensions  $d \times n$
- $\mathbf{X}^T \mathbf{X}$  is a matrix product of matrices of size:  $d \times n$  and  $n \times d$ , which is  $\mathcal{O}(d^2 n)$

## Normal Equation

- $\mathbf{X}$  has dimensions  $n \times d$ ,  $\mathbf{X}^T$  has dimensions  $d \times n$
- $\mathbf{X}^T \mathbf{X}$  is a matrix product of matrices of size:  $d \times n$  and  $n \times d$ , which is  $\mathcal{O}(d^2 n)$
- Inversion of  $\mathbf{X}^T \mathbf{X}$  is an inversion of a  $d \times d$  matrix, which is  $\mathcal{O}(d^3)$

## Normal Equation

- $\mathbf{X}$  has dimensions  $n \times d$ ,  $\mathbf{X}^T$  has dimensions  $d \times n$
- $\mathbf{X}^T \mathbf{X}$  is a matrix product of matrices of size:  $d \times n$  and  $n \times d$ , which is  $\mathcal{O}(d^2 n)$
- Inversion of  $\mathbf{X}^T \mathbf{X}$  is an inversion of a  $d \times d$  matrix, which is  $\mathcal{O}(d^3)$
- $\mathbf{X}^T \mathbf{y}$  is a matrix vector product of size  $d \times n$  and  $n \times 1$ , which is  $\mathcal{O}(dn)$

## Normal Equation

- $\mathbf{X}$  has dimensions  $n \times d$ ,  $\mathbf{X}^T$  has dimensions  $d \times n$
- $\mathbf{X}^T \mathbf{X}$  is a matrix product of matrices of size:  $d \times n$  and  $n \times d$ , which is  $\mathcal{O}(d^2 n)$
- Inversion of  $\mathbf{X}^T \mathbf{X}$  is an inversion of a  $d \times d$  matrix, which is  $\mathcal{O}(d^3)$
- $\mathbf{X}^T \mathbf{y}$  is a matrix vector product of size  $d \times n$  and  $n \times 1$ , which is  $\mathcal{O}(dn)$
- $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$  is a matrix product of a  $d \times d$  matrix and  $d \times 1$  matrix, which is  $\mathcal{O}(d^2)$

## Normal Equation

- $\mathbf{X}$  has dimensions  $n \times d$ ,  $\mathbf{X}^T$  has dimensions  $d \times n$
- $\mathbf{X}^T \mathbf{X}$  is a matrix product of matrices of size:  $d \times n$  and  $n \times d$ , which is  $\mathcal{O}(d^2 n)$
- Inversion of  $\mathbf{X}^T \mathbf{X}$  is an inversion of a  $d \times d$  matrix, which is  $\mathcal{O}(d^3)$
- $\mathbf{X}^T \mathbf{y}$  is a matrix vector product of size  $d \times n$  and  $n \times 1$ , which is  $\mathcal{O}(dn)$
- $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$  is a matrix product of a  $d \times d$  matrix and  $d \times 1$  matrix, which is  $\mathcal{O}(d^2)$
- Overall complexity:  $\mathcal{O}(d^2 n) + \mathcal{O}(d^3) + \mathcal{O}(dn) + \mathcal{O}(d^2) = \mathcal{O}(d^2 n) + \mathcal{O}(d^3)$

## Normal Equation

- $\mathbf{X}$  has dimensions  $n \times d$ ,  $\mathbf{X}^T$  has dimensions  $d \times n$
- $\mathbf{X}^T \mathbf{X}$  is a matrix product of matrices of size:  $d \times n$  and  $n \times d$ , which is  $\mathcal{O}(d^2 n)$
- Inversion of  $\mathbf{X}^T \mathbf{X}$  is an inversion of a  $d \times d$  matrix, which is  $\mathcal{O}(d^3)$
- $\mathbf{X}^T \mathbf{y}$  is a matrix vector product of size  $d \times n$  and  $n \times 1$ , which is  $\mathcal{O}(dn)$
- $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$  is a matrix product of a  $d \times d$  matrix and  $d \times 1$  matrix, which is  $\mathcal{O}(d^2)$
- Overall complexity:  $\mathcal{O}(d^2 n) + \mathcal{O}(d^3) + \mathcal{O}(dn) + \mathcal{O}(d^2) = \mathcal{O}(d^2 n) + \mathcal{O}(d^3)$
- Scales cubic in the number of columns/features of  $\mathbf{X}$

# Gradient Descent

Start with random values of  $\theta_0$  and  $\theta_1$

Till convergence

- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} (\sum \epsilon_i^2)$

# Gradient Descent

Start with random values of  $\theta_0$  and  $\theta_1$

Till convergence

- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} (\sum \epsilon_i^2)$
- $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} (\sum \epsilon_i^2)$



# Gradient Descent

Start with random values of  $\theta_0$  and  $\theta_1$

Till convergence

- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} (\sum \epsilon_i^2)$
- $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} (\sum \epsilon_i^2)$
- Question: Can you write the above for  $d$  dimensional data in vectorised form?

# Gradient Descent

Start with random values of  $\theta_0$  and  $\theta_1$

Till convergence

- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} (\sum \epsilon_i^2)$
- $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} (\sum \epsilon_i^2)$
- Question: Can you write the above for  $d$  dimensional data in vectorised form?
- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} (\mathbf{y} - \mathbf{X}\theta)^\top (\mathbf{y} - \mathbf{X}\theta)$   
 $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} (\mathbf{y} - \mathbf{X}\theta)^\top (\mathbf{y} - \mathbf{X}\theta)$   
 $\vdots$   
 $\theta_d = \theta_d - \alpha \frac{\partial}{\partial \theta_d} (\mathbf{y} - \mathbf{X}\theta)^\top (\mathbf{y} - \mathbf{X}\theta)$

# Gradient Descent

Start with random values of  $\theta_0$  and  $\theta_1$

Till convergence

- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} (\sum \epsilon_i^2)$
- $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} (\sum \epsilon_i^2)$
- Question: Can you write the above for  $d$  dimensional data in vectorised form?
- $\theta_0 = \theta_0 - \alpha \frac{\partial}{\partial \theta_0} (\mathbf{y} - \mathbf{X}\theta)^\top (\mathbf{y} - \mathbf{X}\theta)$   
 $\theta_1 = \theta_1 - \alpha \frac{\partial}{\partial \theta_1} (\mathbf{y} - \mathbf{X}\theta)^\top (\mathbf{y} - \mathbf{X}\theta)$   
 $\vdots$   
 $\theta_d = \theta_d - \alpha \frac{\partial}{\partial \theta_d} (\mathbf{y} - \mathbf{X}\theta)^\top (\mathbf{y} - \mathbf{X}\theta)$
- $\theta = \theta - \alpha \frac{\partial}{\partial \theta} (\mathbf{y} - \mathbf{X}\theta)^\top (\mathbf{y} - \mathbf{X}\theta)$

# Gradient Descent

$$\begin{aligned}& \frac{\partial}{\partial \boldsymbol{\theta}} (\mathbf{y} - \mathbf{X}\boldsymbol{\theta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}) \\&= \frac{\partial}{\partial \boldsymbol{\theta}} (\mathbf{y}^\top - \boldsymbol{\theta}^\top \mathbf{X}^\top) (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}) \\&= \frac{\partial}{\partial \boldsymbol{\theta}} (\mathbf{y}^\top \mathbf{y} - \boldsymbol{\theta}^\top \mathbf{X}^\top \mathbf{y} - \mathbf{y}^\top \mathbf{X}\boldsymbol{\theta} + \boldsymbol{\theta}^\top \mathbf{X}^\top \mathbf{X}\boldsymbol{\theta}) \\&= -2\mathbf{X}^\top \mathbf{y} + 2\mathbf{X}^\top \mathbf{X}\boldsymbol{\theta} \\&= 2\mathbf{X}^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})\end{aligned}$$

## Gradient Descent

We can write the vectorised update equation as follows, for each iteration

$$\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$$

## Gradient Descent

We can write the vectorised update equation as follows, for each iteration

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^{\top}(\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

For  $t$  iterations, what is the computational complexity of our gradient descent solution?

## Gradient Descent

We can write the vectorised update equation as follows, for each iteration

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

For  $t$  iterations, what is the computational complexity of our gradient descent solution?

Hint, rewrite the above as:  $\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top \mathbf{X} \boldsymbol{\theta} + \alpha \mathbf{X}^\top \mathbf{y}$

## Gradient Descent

We can write the vectorised update equation as follows, for each iteration

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

For  $t$  iterations, what is the computational complexity of our gradient descent solution?

Hint, rewrite the above as:  $\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top \mathbf{X} \boldsymbol{\theta} + \alpha \mathbf{X}^\top \mathbf{y}$

Complexity of computing  $\mathbf{X}^\top \mathbf{y}$  is  $\mathcal{O}(dn)$



## Gradient Descent

We can write the vectorised update equation as follows, for each iteration

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

For  $t$  iterations, what is the computational complexity of our gradient descent solution?

Hint, rewrite the above as:  $\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top \mathbf{X}\boldsymbol{\theta} + \alpha \mathbf{X}^\top \mathbf{y}$

Complexity of computing  $\mathbf{X}^\top \mathbf{y}$  is  $\mathcal{O}(dn)$

Complexity of computing  $\alpha \mathbf{X}^\top \mathbf{y}$  once we have  $\mathbf{X}^\top \mathbf{y}$  is  $\mathcal{O}(d)$  since  $\mathbf{X}^\top \mathbf{y}$  has  $d$  entries

## Gradient Descent

We can write the vectorised update equation as follows, for each iteration

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

For  $t$  iterations, what is the computational complexity of our gradient descent solution?

Hint, rewrite the above as:  $\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top \mathbf{X} \boldsymbol{\theta} + \alpha \mathbf{X}^\top \mathbf{y}$

Complexity of computing  $\mathbf{X}^\top \mathbf{y}$  is  $\mathcal{O}(dn)$

Complexity of computing  $\alpha \mathbf{X}^\top \mathbf{y}$  once we have  $\mathbf{X}^\top \mathbf{y}$  is  $\mathcal{O}(d)$  since  $\mathbf{X}^\top \mathbf{y}$  has  $d$  entries

Complexity of computing  $\mathbf{X}^\top \mathbf{X}$  is  $\mathcal{O}(d^2n)$  and then multiplying with  $\alpha$  is  $\mathcal{O}(d^2)$

# Gradient Descent

We can write the vectorised update equation as follows, for each iteration

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

For  $t$  iterations, what is the computational complexity of our gradient descent solution?

Hint, rewrite the above as:  $\boldsymbol{\theta} = \boldsymbol{\theta} - \alpha \mathbf{X}^\top \mathbf{X} \boldsymbol{\theta} + \alpha \mathbf{X}^\top \mathbf{y}$

Complexity of computing  $\mathbf{X}^\top \mathbf{y}$  is  $\mathcal{O}(dn)$

Complexity of computing  $\alpha \mathbf{X}^\top \mathbf{y}$  once we have  $\mathbf{X}^\top \mathbf{y}$  is  $\mathcal{O}(d)$  since  $\mathbf{X}^\top \mathbf{y}$  has  $d$  entries

Complexity of computing  $\mathbf{X}^\top \mathbf{X}$  is  $\mathcal{O}(d^2n)$  and then multiplying with  $\alpha$  is  $\mathcal{O}(d^2)$

All of the above need only be calculated once!

# Gradient Descent

For each of the  $t$  iterations, we now need to first multiply  $\alpha \mathbf{X}^\top \mathbf{X}$  with  $\boldsymbol{\theta}$  which is matrix multiplication of a  $d \times d$  matrix with a  $d \times 1$ , which is  $\mathcal{O}(d^2)$

For each of the  $t$  iterations, we now need to first multiply  $\alpha \mathbf{X}^\top \mathbf{X}$  with  $\boldsymbol{\theta}$  which is matrix multiplication of a  $d \times d$  matrix with a  $d \times 1$ , which is  $\mathcal{O}(d^2)$

The remaining subtraction/addition can be done in  $\mathcal{O}(d)$  for each iteration.

For each of the  $t$  iterations, we now need to first multiply  $\alpha \mathbf{X}^\top \mathbf{X}$  with  $\boldsymbol{\theta}$  which is matrix multiplication of a  $d \times d$  matrix with a  $d \times 1$ , which is  $\mathcal{O}(d^2)$

The remaining subtraction/addition can be done in  $\mathcal{O}(d)$  for each iteration.

What is overall computational complexity?

## Gradient Descent

For each of the  $t$  iterations, we now need to first multiply  $\alpha \mathbf{X}^\top \mathbf{X}$  with  $\boldsymbol{\theta}$  which is matrix multiplication of a  $d \times d$  matrix with a  $d \times 1$ , which is  $\mathcal{O}(d^2)$

The remaining subtraction/addition can be done in  $\mathcal{O}(d)$  for each iteration.

What is overall computational complexity?

$$\mathcal{O}(td^2) + \mathcal{O}(d^2n) = \mathcal{O}((t+n)d^2)$$



## Gradient Descent (Alternative)

## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$

## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$
- Computing  $\mathbf{X}\theta - \mathbf{y}$  is  $\mathcal{O}(n)$

## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$
- Computing  $\mathbf{X}\theta - \mathbf{y}$  is  $\mathcal{O}(n)$
- Computing  $\alpha \mathbf{X}^\top$  is  $\mathcal{O}(nd)$

## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$
- Computing  $\mathbf{X}\theta - \mathbf{y}$  is  $\mathcal{O}(n)$
- Computing  $\alpha \mathbf{X}^\top$  is  $\mathcal{O}(nd)$
- Computing  $\alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(nd)$

## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$
- Computing  $\mathbf{X}\theta - \mathbf{y}$  is  $\mathcal{O}(n)$
- Computing  $\alpha \mathbf{X}^\top$  is  $\mathcal{O}(nd)$
- Computing  $\alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(nd)$
- Computing  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(n)$

## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$
- Computing  $\mathbf{X}\theta - \mathbf{y}$  is  $\mathcal{O}(n)$
- Computing  $\alpha \mathbf{X}^\top$  is  $\mathcal{O}(nd)$
- Computing  $\alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(nd)$
- Computing  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(n)$

## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$
- Computing  $\mathbf{X}\theta - \mathbf{y}$  is  $\mathcal{O}(n)$
- Computing  $\alpha \mathbf{X}^\top$  is  $\mathcal{O}(nd)$
- Computing  $\alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(nd)$
- Computing  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(n)$

What is overall computational complexity?



## Gradient Descent (Alternative)

If we do not rewrite the expression  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$

For each iteration, we have:

- Computing  $\mathbf{X}\theta$  is  $\mathcal{O}(nd)$
- Computing  $\mathbf{X}\theta - \mathbf{y}$  is  $\mathcal{O}(n)$
- Computing  $\alpha \mathbf{X}^\top$  is  $\mathcal{O}(nd)$
- Computing  $\alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(nd)$
- Computing  $\theta = \theta - \alpha \mathbf{X}^\top (\mathbf{X}\theta - \mathbf{y})$  is  $\mathcal{O}(n)$

What is overall computational complexity?

$\mathcal{O}(ndt)$