

# Decision Trees

---

Nipun Batra and teaching staff

IIT Gandhinagar

August 8, 2025

# Table of Contents

1. Introduction and Motivation
2. Discrete Input, Discrete Output
3. Discrete Input, Real Output
4. Real Input Discrete Output
5. Real Input Real Output
6. Pruning and Overfitting
7. Summary and Key Takeaways
8. Weighted Entropy

# The need for interpretability

## How to maintain trust in AI


Beyond developing initial trust, however, creators of AI also must work to maintain that trust. Siau and Wang suggest seven ways of "developing continuous trust" beyond the initial phases of product development:

- Usability and reliability. AI "should be designed to operate easily and intuitively," Siau and Wang write. "There should be no unexpected downtime or crashes."
- Collaboration and communication. AI developers want to create systems that perform autonomously, without human involvement. Developers must focus on creating AI applications that smoothly and easily collaborate and communicate with humans.
- Sociability and bonding. Building social activities into AI applications is one way to strengthen trust. A robotic dog that can recognize its owner and show affection is one example, Siau and Wang write.
- Security and privacy protection. AI applications rely on large data sets, so ensuring privacy and security will be crucial to establishing trust in the applications.
- Interpretability. Just as transparency is instrumental in building initial trust, interpretability – or the ability for a machine to explain its conclusions or actions – will help sustain trust


# Training Data

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

# Learning a Complicated Neural Network



# Learnt Decision Tree



# Medical Diagnosis using Decision Trees



Source: Improving medical decision trees by combining relevant health-care criteria

# Leo Breiman (1928-2005)



## Leo Breiman 1928-2005

Professor of Statistics, UC Berkeley  
Verified email at stat.berkeley.edu - Homepage  
Data Analysis Statistics Machine Learning


100%

TITLE	CITED BY	YEAR
<b>Random forests</b> L. Breiman Machine learning 45 (1), 5-32	53816	2001
<b>Classification and Regression Trees</b> L. Breiman, J.H. Friedman, R.A. Olshen, C.J. Stone CRC Press, New York	49900	1999
<b>Classification and regression trees</b> L. Breiman Chapman & Hall/CRC	43992	1984
<b>Bagging predictors</b> L. Breiman Machine learning 24 (2), 123-140	22742	1996
<b>Statistical Modeling: The Two Cultures</b> L. Breiman	2785	2003
<b>Statistical modeling: The two cultures (with comments and a rejoinder by the author)</b> L. Breiman Statistical Science 18 (3), 199-231	2772	2001
<b>Estimates of central transformations for multiple regression and correlation</b>	3096	1985

Cited by

VIEW ALL

	All	Since 2015
Citations	142867	69736
h-index	51	33
i10-index	80	46





# Leo Breiman: Revolutionary Contributions to ML

## Key Points

### Major Algorithmic Breakthroughs:

- **CART (1984):** Classification and Regression Trees

# Leo Breiman: Revolutionary Contributions to ML

## Key Points

### Major Algorithmic Breakthroughs:

- **CART (1984):** Classification and Regression Trees
- **Bagging (1994):** Bootstrap Aggregating

# Leo Breiman: Revolutionary Contributions to ML

## Key Points

### Major Algorithmic Breakthroughs:

- **CART (1984)**: Classification and Regression Trees
- **Bagging (1994)**: Bootstrap Aggregating
- **Random Forests (2001)**: Ensemble of decision trees

# Leo Breiman: Revolutionary Contributions to ML

## Key Points

### Major Algorithmic Breakthroughs:

- **CART (1984)**: Classification and Regression Trees
- **Bagging (1994)**: Bootstrap Aggregating
- **Random Forests (2001)**: Ensemble of decision trees
- **Two Cultures (2001)**: Data modeling vs. algorithmic modeling

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time
  - Example: Given a sudoku solution, verify it's correct



# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time
  - Example: Given a sudoku solution, verify it's correct
- **NP-Complete**: Hardest problems in NP

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time
  - Example: Given a sudoku solution, verify it's correct
- **NP-Complete**: Hardest problems in NP
  - Both in NP and at least as hard as any NP problem

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time
  - Example: Given a sudoku solution, verify it's correct
- **NP-Complete**: Hardest problems in NP
  - Both in NP and at least as hard as any NP problem
  - Example: Boolean satisfiability (SAT)

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time
  - Example: Given a sudoku solution, verify it's correct
- **NP-Complete**: Hardest problems in NP
  - Both in NP and at least as hard as any NP problem
  - Example: Boolean satisfiability (SAT)
- **NP-Hard**: At least as hard as NP-Complete problems

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time
  - Example: Given a sudoku solution, verify it's correct
- **NP-Complete**: Hardest problems in NP
  - Both in NP and at least as hard as any NP problem
  - Example: Boolean satisfiability (SAT)
- **NP-Hard**: At least as hard as NP-Complete problems
  - May not be in NP (solutions might not be verifiable quickly)

# Computational Complexity Classes: A Quick Primer

## Definition: Key Complexity Classes

- **P**: Problems solvable in polynomial time
  - Example: Sorting  $n$  numbers in  $O(n \log n)$  time
- **NP**: Problems where solutions can be *verified* in polynomial time
  - Example: Given a sudoku solution, verify it's correct
- **NP-Complete**: Hardest problems in NP
  - Both in NP and at least as hard as any NP problem
  - Example: Boolean satisfiability (SAT)
- **NP-Hard**: At least as hard as NP-Complete problems
  - May not be in NP (solutions might not be verifiable quickly)
  - Example: Optimization versions of NP-Complete problems

# Finding the Optimal Decision Tree

Volume 5, number 1

INFORMATION PROCESSING LETTERS

May 1976

## CONSTRUCTING OPTIMAL BINARY DECISION TREES IS NP-COMPLETE\*

Laurent HYAFIL

*IRIA – Laboria, 78150 Rocquencourt, France*

and

Ronald L. RIVEST

*Dept. of Electrical Engineering and Computer Science, M.I.T., Cambridge, Massachusetts 02139, USA*

Received 7 November 1975, revised version received 26 January 1976

Binary decision trees, computational complexity, NP-complete

**The Problem:** Given training data, find the decision tree with highest accuracy

# Optimal Decision Trees are NP-Complete

## Important: Computational Complexity

### Finding optimal decision tree is NP-Complete

- **Verification:** Given a tree, check its accuracy quickly ✓



# Optimal Decision Trees are NP-Complete

## Important: Computational Complexity

### Finding optimal decision tree is NP-Complete

- **Verification:** Given a tree, check its accuracy quickly ✓
- **Construction:** Exponentially many trees to check ✗

# Optimal Decision Trees are NP-Complete

## Important: Computational Complexity

### Finding optimal decision tree is NP-Complete

- **Verification:** Given a tree, check its accuracy quickly ✓
- **Construction:** Exponentially many trees to check ✗

# Optimal Decision Trees are NP-Complete

## Important: Computational Complexity

### Finding optimal decision tree is NP-Complete

- **Verification:** Given a tree, check its accuracy quickly ✓
- **Construction:** Exponentially many trees to check ✗

### Example: What This Means

- No efficient algorithm exists (unless  $P = NP$ )

# Optimal Decision Trees are NP-Complete

## Important: Computational Complexity

### Finding optimal decision tree is NP-Complete

- **Verification:** Given a tree, check its accuracy quickly ✓
- **Construction:** Exponentially many trees to check ✗

### Example: What This Means

- No efficient algorithm exists (unless  $P = NP$ )
- Must use heuristics like greedy algorithms

# Optimal Decision Trees are NP-Complete

## Important: Computational Complexity

### Finding optimal decision tree is NP-Complete

- **Verification:** Given a tree, check its accuracy quickly ✓
- **Construction:** Exponentially many trees to check ✗

### Example: What This Means

- No efficient algorithm exists (unless  $P = NP$ )
- Must use heuristics like greedy algorithms
- ID3, C4.5, CART use greedy approaches

# Optimal Decision Trees are NP-Complete

## Important: Computational Complexity

### Finding optimal decision tree is NP-Complete

- **Verification:** Given a tree, check its accuracy quickly ✓
- **Construction:** Exponentially many trees to check ✗

### Example: What This Means

- No efficient algorithm exists (unless  $P = NP$ )
- Must use heuristics like greedy algorithms
- ID3, C4.5, CART use greedy approaches
- Good solutions, but no optimality guarantee

# Greedy Algorithm

Core idea: At each level, choose an attribute that gives **biggest estimated** performance gain!




Image source: analyticsvidhya

Greedy  $\neq$  Optimal

Towards biggest estimated performance gain



## Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

## Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

## Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

- For examples, we have 9 Yes, 5 No

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

- For examples, we have 9 Yes, 5 No

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!



# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!
- Key insight: Problem is “easier” when there is less disagreement

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!
- Key insight: Problem is “easier” when there is less disagreement

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!
- Key insight: Problem is “easier” when there is less disagreement
- Need some statistical measure of “disagreement”

# Claude Shannon (1948): The Birth of Information Theory

## Definition: The Big Idea

Information is inversely related to probability. **Rare events are more informative!**

# Claude Shannon (1948): The Birth of Information Theory

## Definition: The Big Idea

Information is inversely related to probability. **Rare events are more informative!**

**Think about it:** Which headline tells you more?

- "The sun rose this morning"

# Claude Shannon (1948): The Birth of Information Theory

## Definition: The Big Idea

Information is inversely related to probability. **Rare events are more informative!**

**Think about it:** Which headline tells you more?

- "The sun rose this morning"
- "It snowed in Death Valley in July"

# Claude Shannon (1948): The Birth of Information Theory

## Definition: The Big Idea

Information is inversely related to probability. **Rare events are more informative!**

**Think about it:** Which headline tells you more?

- "The sun rose this morning"
- "It snowed in Death Valley in July"



# Claude Shannon (1948): The Birth of Information Theory

## Definition: The Big Idea

Information is inversely related to probability. **Rare events are more informative!**

**Think about it:** Which headline tells you more?

- "The sun rose this morning"
- "It snowed in Death Valley in July"

The second one! Because it's **unexpected**.

# Claude Shannon (1948): The Birth of Information Theory

## Definition: The Big Idea

Information is inversely related to probability. **Rare events are more informative!**

**Think about it:** Which headline tells you more?

- "The sun rose this morning"
- "It snowed in Death Valley in July"

The second one! Because it's **unexpected**.

**Shannon's insight:** The amount of information in an event should be inversely proportional to its probability.

# Measuring Surprise: Step by Step

## Shannon's Information Formula:

$$I(x) = -\log_2 p(x)$$

# Measuring Surprise: Step by Step

## Shannon's Information Formula:

$$I(x) = -\log_2 p(x)$$

## Why the negative log?

- Probabilities are between 0 and 1

# Measuring Surprise: Step by Step

## Shannon's Information Formula:

$$I(x) = -\log_2 p(x)$$

## Why the negative log?

- Probabilities are between 0 and 1
- $\log_2$  of values  $< 1$  gives negative numbers

# Measuring Surprise: Step by Step

## Shannon's Information Formula:

$$I(x) = -\log_2 p(x)$$

## Why the negative log?

- Probabilities are between 0 and 1
- $\log_2$  of values  $< 1$  gives negative numbers
- We want information to be positive

# Measuring Surprise: Step by Step

## Shannon's Information Formula:

$$I(x) = -\log_2 p(x)$$

## Why the negative log?

- Probabilities are between 0 and 1
- $\log_2$  of values  $< 1$  gives negative numbers
- We want information to be positive
- Hence the negative sign!

# Measuring Surprise: Step by Step

## Shannon's Information Formula:

$$I(x) = -\log_2 p(x)$$

## Why the negative log?

- Probabilities are between 0 and 1
- $\log_2$  of values  $< 1$  gives negative numbers
- We want information to be positive
- Hence the negative sign!



# Measuring Surprise: Step by Step

## Shannon's Information Formula:

$$I(x) = -\log_2 p(x)$$

## Why the negative log?

- Probabilities are between 0 and 1
- $\log_2$  of values  $< 1$  gives negative numbers
- We want information to be positive
- Hence the negative sign!

**Why base 2?** So information is measured in **bits**.

# Calculating Surprise: Detailed Examples

## **Example 1: Summer weather in Phoenix**

- Sunny day:  $p = 0.9$

# Calculating Surprise: Detailed Examples

## **Example 1: Summer weather in Phoenix**

- Sunny day:  $p = 0.9$

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

## Example 2: Snow in Phoenix in July

- Probability:  $p = 0.0001$  (extremely rare!)



# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

## Example 2: Snow in Phoenix in July

- Probability:  $p = 0.0001$  (extremely rare!)

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

## Example 2: Snow in Phoenix in July

- Probability:  $p = 0.0001$  (extremely rare!)
- Information:  $I = -\log_2(0.0001) = -(-13.29) = 13.29$  bits

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

## Example 2: Snow in Phoenix in July

- Probability:  $p = 0.0001$  (extremely rare!)
- Information:  $I = -\log_2(0.0001) = -(-13.29) = 13.29$  bits

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

## Example 2: Snow in Phoenix in July

- Probability:  $p = 0.0001$  (extremely rare!)
- Information:  $I = -\log_2(0.0001) = -(-13.29) = 13.29$  bits
- **High surprise** - this would be shocking news!

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

## Example 2: Snow in Phoenix in July

- Probability:  $p = 0.0001$  (extremely rare!)
- Information:  $I = -\log_2(0.0001) = -(-13.29) = 13.29$  bits
- **High surprise** - this would be shocking news!

# Calculating Surprise: Detailed Examples

## Example 1: Summer weather in Phoenix

- Sunny day:  $p = 0.9$
- Information:  $I = -\log_2(0.9) = -(-0.152) = 0.152$  bits
- **Low surprise** - we expect sunny weather

## Example 2: Snow in Phoenix in July

- Probability:  $p = 0.0001$  (extremely rare!)
- Information:  $I = -\log_2(0.0001) = -(-13.29) = 13.29$  bits
- **High surprise** - this would be shocking news!

**Notice:** Rare events carry  $\sim 90\times$  more information!

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$



# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$
- Snowy:  $p = 0.05$

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$
- Snowy:  $p = 0.05$

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$
- Snowy:  $p = 0.05$

**Problem:** Each day gives different amounts of information!

- If it's rainy:  $I = -\log_2(0.5) = 1.0$  bit

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$
- Snowy:  $p = 0.05$

**Problem:** Each day gives different amounts of information!

- If it's rainy:  $I = -\log_2(0.5) = 1.0$  bit
- If it's sunny:  $I = -\log_2(0.15) = 2.74$  bits

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$
- Snowy:  $p = 0.05$

**Problem:** Each day gives different amounts of information!

- If it's rainy:  $I = -\log_2(0.5) = 1.0$  bit
- If it's sunny:  $I = -\log_2(0.15) = 2.74$  bits
- If it's snowy:  $I = -\log_2(0.05) = 4.32$  bits

# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$
- Snowy:  $p = 0.05$

**Problem:** Each day gives different amounts of information!

- If it's rainy:  $I = -\log_2(0.5) = 1.0$  bit
- If it's sunny:  $I = -\log_2(0.15) = 2.74$  bits
- If it's snowy:  $I = -\log_2(0.05) = 4.32$  bits



# From Single Events to Distributions

**Question:** What if we have multiple possible outcomes?

**Example:** Weather in Seattle (4 possibilities)

- Rainy:  $p = 0.5$
- Cloudy:  $p = 0.3$
- Sunny:  $p = 0.15$
- Snowy:  $p = 0.05$

**Problem:** Each day gives different amounts of information!

- If it's rainy:  $I = -\log_2(0.5) = 1.0$  bit
- If it's sunny:  $I = -\log_2(0.15) = 2.74$  bits
- If it's snowy:  $I = -\log_2(0.05) = 4.32$  bits

**Solution:** Take the **expected** (average) information!

# Entropy: Expected Information

## Definition: Entropy Formula

$$H(X) = \mathbb{E}[I(X)] = - \sum_i p(x_i) \log_2 p(x_i)$$

**Entropy** = Expected amount of information per observation

# Entropy: Expected Information

## Definition: Entropy Formula

$$H(X) = \mathbb{E}[I(X)] = - \sum_i p(x_i) \log_2 p(x_i)$$

**Entropy** = Expected amount of information per observation

### Seattle weather calculation:

$$H = -p(\text{rain}) \log_2 p(\text{rain}) - p(\text{cloudy}) \log_2 p(\text{cloudy}) \quad (1)$$

$$- p(\text{sunny}) \log_2 p(\text{sunny}) - p(\text{snow}) \log_2 p(\text{snow}) \quad (2)$$

$$= -0.5 \log_2(0.5) - 0.3 \log_2(0.3) - 0.15 \log_2(0.15) - 0.05 \log_2(0.05) \quad (3)$$

$$= 0.5(1.0) + 0.3(1.74) + 0.15(2.74) + 0.05(4.32) \quad (4)$$

$$= 0.5 + 0.52 + 0.41 + 0.22 = \mathbf{1.65} \text{ bits} \quad (5)$$

# Entropy: Expected Information

## Definition: Entropy Formula

$$H(X) = \mathbb{E}[I(X)] = - \sum_i p(x_i) \log_2 p(x_i)$$

**Entropy** = Expected amount of information per observation

### Seattle weather calculation:

$$H = -p(\text{rain}) \log_2 p(\text{rain}) - p(\text{cloudy}) \log_2 p(\text{cloudy}) \quad (1)$$

$$- p(\text{sunny}) \log_2 p(\text{sunny}) - p(\text{snow}) \log_2 p(\text{snow}) \quad (2)$$

$$= 0.5(1.0) + 0.3(1.74) + 0.15(2.74) + 0.05(4.32) \quad (4)$$

$$= 0.5 + 0.52 + 0.41 + 0.22 = \mathbf{1.65} \text{ bits} \quad (5)$$

# Entropy: Expected Information

## Definition: Entropy Formula

$$H(X) = \mathbb{E}[I(X)] = - \sum_i p(x_i) \log_2 p(x_i)$$

**Entropy** = Expected amount of information per observation

### Seattle weather calculation:

$$H = -p(\text{rain}) \log_2 p(\text{rain}) - p(\text{cloudy}) \log_2 p(\text{cloudy}) \quad (1)$$

$$- p(\text{sunny}) \log_2 p(\text{sunny}) - p(\text{snow}) \log_2 p(\text{snow}) \quad (2)$$

$$= 0.5 + 0.52 + 0.41 + 0.22 = \mathbf{1.65} \text{ bits} \quad (5)$$

# Entropy: Expected Information

## Definition: Entropy Formula

$$H(X) = \mathbb{E}[I(X)] = - \sum_i p(x_i) \log_2 p(x_i)$$

**Entropy** = Expected amount of information per observation

### Seattle weather calculation:

$$H = -p(\text{rain}) \log_2 p(\text{rain}) - p(\text{cloudy}) \log_2 p(\text{cloudy}) \quad (1)$$

$$- p(\text{sunny}) \log_2 p(\text{sunny}) - p(\text{snow}) \log_2 p(\text{snow}) \quad (2)$$

$$= -0.5 \log_2(0.5) - 0.3 \log_2(0.3) - 0.15 \log_2(0.15) - 0.05 \log_2(0.05) \quad (3)$$

$$= 0.5(1.0) + 0.3(1.74) + 0.15(2.74) + 0.05(4.32) \quad (4)$$

$$= 0.5 + 0.52 + 0.41 + 0.22 = \mathbf{1.65} \text{ bits} \quad (5)$$

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )



# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

## Case 2: Maximum uncertainty

- Fair coin: Heads/Tails equally likely ( $p = 0.5$  each)

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

## Case 2: Maximum uncertainty

- Fair coin: Heads/Tails equally likely ( $p = 0.5$  each)

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

## Case 2: Maximum uncertainty

- Fair coin: Heads/Tails equally likely ( $p = 0.5$  each)
- $H = -0.5 \log_2(0.5) - 0.5 \log_2(0.5) = 0.5(1) + 0.5(1) = \mathbf{1.0}$  bit

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

## Case 2: Maximum uncertainty

- Fair coin: Heads/Tails equally likely ( $p = 0.5$  each)
- $H = -0.5 \log_2(0.5) - 0.5 \log_2(0.5) = 0.5(1) + 0.5(1) = \mathbf{1.0}$  bit



# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

## Case 2: Maximum uncertainty

- Fair coin: Heads/Tails equally likely ( $p = 0.5$  each)
- $H = -0.5 \log_2(0.5) - 0.5 \log_2(0.5) = 0.5(1) + 0.5(1) = \mathbf{1.0}$  bit
- **Maximum entropy** = Maximum surprise = Completely unpredictable

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

## Case 2: Maximum uncertainty

- Fair coin: Heads/Tails equally likely ( $p = 0.5$  each)
- $H = -0.5 \log_2(0.5) - 0.5 \log_2(0.5) = 0.5(1) + 0.5(1) = \mathbf{1.0}$  bit
- **Maximum entropy** = Maximum surprise = Completely unpredictable

# Entropy Intuition: Extreme Cases

## Case 1: Completely predictable

- Desert: Always sunny ( $p = 1.0$ )
- $H = -1.0 \log_2(1.0) = -1.0 \times 0 = \mathbf{0}$  bits
- **Zero entropy** = No surprise = Completely predictable

## Case 2: Maximum uncertainty

- Fair coin: Heads/Tails equally likely ( $p = 0.5$  each)
- $H = -0.5 \log_2(0.5) - 0.5 \log_2(0.5) = 0.5(1) + 0.5(1) = \mathbf{1.0}$  bit
- **Maximum entropy** = Maximum surprise = Completely unpredictable

**Key insight:** Entropy ranges from 0 (certain) to  $\log_2(n)$  (uniform over  $n$  outcomes)

# Entropy in Decision Trees: The Connection

**Why do we care about entropy in ML?**

# Entropy in Decision Trees: The Connection

Why do we care about entropy in ML?

## Example: Decision Tree Goal

We want to split data into **pure** subsets where we can make confident predictions.

# Entropy in Decision Trees: The Connection

## Why do we care about entropy in ML?

### Example: Decision Tree Goal

We want to split data into **pure** subsets where we can make confident predictions.

- **Pure node:** All examples same class → Low entropy → Good split

# Entropy in Decision Trees: The Connection

## Why do we care about entropy in ML?

### Example: Decision Tree Goal

We want to split data into **pure** subsets where we can make confident predictions.

- **Pure node:** All examples same class → Low entropy → Good split

# Entropy in Decision Trees: The Connection

## Why do we care about entropy in ML?

### Example: Decision Tree Goal

We want to split data into **pure** subsets where we can make confident predictions.

- **Pure node:** All examples same class → Low entropy → Good split
- **Mixed node:** Examples from different classes → High entropy → Bad split



# Entropy in Decision Trees: The Connection

## Why do we care about entropy in ML?

### Example: Decision Tree Goal

We want to split data into **pure** subsets where we can make confident predictions.

- **Pure node:** All examples same class → Low entropy → Good split
- **Mixed node:** Examples from different classes → High entropy → Bad split

# Entropy in Decision Trees: The Connection

## Why do we care about entropy in ML?

### Example: Decision Tree Goal

We want to split data into **pure** subsets where we can make confident predictions.

- **Pure node:** All examples same class → Low entropy → Good split
- **Mixed node:** Examples from different classes → High entropy → Bad split

**Strategy:** Choose splits that **reduce entropy** the most!

# Entropy in Decision Trees: The Connection

## Why do we care about entropy in ML?

### Example: Decision Tree Goal

We want to split data into **pure** subsets where we can make confident predictions.

- **Pure node:** All examples same class → Low entropy → Good split
- **Mixed node:** Examples from different classes → High entropy → Bad split

**Strategy:** Choose splits that **reduce entropy** the most!  
This is exactly what **Information Gain** measures.

# Entropy


Statistical measure to characterize the (im)purity of examples

# Entropy

Statistical measure to characterize the (im)purity of examples

$$H(X) = - \sum_{i=1}^k p(x_i) \log_2 p(x_i)$$

**Notebook:** entropy.html



Towards biggest estimated performance gain

## Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

## Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No



## Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

- Can we use Outlook as the root node?

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy		Play
D1	Sunny	Hot	High	Weak		No
D2	Sunny	Hot	High	Strong		No
D3	Overcast	Hot	High	Weak		Yes
D4	Rain	Mild	High	Weak		Yes
D5	Rain	Cool	Normal	Weak		Yes
D6	Rain	Cool	Normal	Strong		No
D7	Overcast	Cool	Normal	Strong		Yes
D8	Sunny	Mild	High	Weak		No
D9	Sunny	Cool	Normal	Weak		Yes
D10	Rain	Mild	Normal	Weak		Yes
D11	Sunny	Mild	Normal	Strong		Yes
D12	Overcast	Mild	High	Strong		Yes
D13	Overcast	Hot	Normal	Weak		Yes
D14	Rain	Mild	High	Strong		No

- Can we use Outlook as the root node?

# Towards biggest estimated performance gain

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

- Can we use Outlook as the root node?
- When Outlook is overcast, we always Play and thus no “disagreement”

# Information Gain

Reduction in entropy by partitioning examples ( $S$ ) on attribute  $A$

$$\text{Gain}(S, A) \equiv \text{Entropy}(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} \text{Entropy}(S_v)$$

# Pop Quiz #1

## Quick Question!

What does entropy measure in the context of decision trees?

A) The depth of the tree

## Pop Quiz #2

### Quick Question!

What does entropy measure in the context of decision trees?

- A) The depth of the tree
- B) The impurity or "disagreement" in a set of examples

## Pop Quiz #3

### Quick Question!

What does entropy measure in the context of decision trees?

- A) The depth of the tree
- B) The impurity or "disagreement" in a set of examples
- C) The number of features in the dataset



## Pop Quiz #4

### Quick Question!

What does entropy measure in the context of decision trees?

- A) The depth of the tree
- B) The impurity or "disagreement" in a set of examples
- C) The number of features in the dataset
- D) The accuracy of the tree

## Pop Quiz #5

### Quick Question!

What does entropy measure in the context of decision trees?

- A) The depth of the tree
- B) The impurity or "disagreement" in a set of examples
- C) The number of features in the dataset
- D) The accuracy of the tree

## Pop Quiz #6

### Quick Question!

What does entropy measure in the context of decision trees?

- A) The depth of the tree
- B) The impurity or "disagreement" in a set of examples
- C) The number of features in the dataset
- D) The accuracy of the tree

**Answer: B) The impurity or "disagreement" in a set of examples** - Higher entropy means more mixed classes, lower entropy means more pure subsets.

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin
  - $A \leftarrow$  attribute from Attributes which best classifies Examples



## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin
  - $A \leftarrow$  attribute from Attributes which best classifies Examples
  - $\text{Root} \leftarrow A$

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin
  - $A \leftarrow$  attribute from Attributes which best classifies Examples
  - $\text{Root} \leftarrow A$
  - For each value ( $v$ ) of  $A$

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin
  - $A \leftarrow$  attribute from Attributes which best classifies Examples
  - $\text{Root} \leftarrow A$
  - For each value ( $v$ ) of  $A$ 
    - Add new tree branch :  $A = v$

# ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin
  - $A \leftarrow$  attribute from Attributes which best classifies Examples
  - $\text{Root} \leftarrow A$
  - For each value ( $v$ ) of  $A$ 
    - Add new tree branch :  $A = v$
    - $\text{Examples}_v$ : subset of examples that  $A = v$

## ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin
  - $A \leftarrow$  attribute from Attributes which best classifies Examples
  - $\text{Root} \leftarrow A$
  - For each value ( $v$ ) of  $A$ 
    - Add new tree branch :  $A = v$
    - $\text{Examples}_v$ : subset of examples that  $A = v$
    - If  $\text{Examples}_v$  is empty: add leaf with label = most common value of Target Attribute

# ID3 (Examples, Target Attribute, Attributes)

- Create a root node for tree
- If all examples are  $+/-$ , return root with label =  $+/-$
- If attributes = empty, return root with most common value of Target Attribute in Examples
- Begin
  - $A \leftarrow$  attribute from Attributes which best classifies Examples
  - $\text{Root} \leftarrow A$
  - For each value ( $v$ ) of  $A$ 
    - Add new tree branch :  $A = v$
    - $\text{Examples}_v$ : subset of examples that  $A = v$
    - If  $\text{Examples}_v$  is empty: add leaf with label = most common value of Target Attribute
    - Else: ID3 ( $\text{Examples}_v$ , Target attribute, Attributes -  $A$ )

# Learnt Decision Tree

Root Node (empty)

# Training Data

Day	Outlook	Temp	Humidity	Windy	Play
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No



## Entropy calculated

We have 14 examples in  $S$ : 5 No, 9 Yes

$$\begin{aligned}\text{Entropy}(S) &= -p_{\text{No}} \log_2 p_{\text{No}} - p_{\text{Yes}} \log_2 p_{\text{Yes}} \\ &= -\frac{5}{14} \log_2 \left( \frac{5}{14} \right) - \frac{9}{14} \log_2 \left( \frac{9}{14} \right) = 0.940\end{aligned}$$

## Information Gain for Outlook

Outlook	Play
Sunny	No
Sunny	No
Overcast	Yes
Rain	Yes
Rain	Yes
Rain	No
Overcast	Yes
Sunny	No
Sunny	Yes
Rain	Yes
Sunny	Yes
Overcast	Yes
Overcast	Yes
Rain	No

# Information Gain for Outlook

## Information Gain for Outlook

Outlook	Play
Sunny	No
Sunny	No
Sunny	No
Sunny	Yes
Sunny	Yes

We have 2 Yes, 3

$$\begin{aligned} \text{No Entropy} = & \\ & -\frac{3}{5} \log_2 \left( \frac{3}{5} \right) - \\ & \frac{2}{5} \log_2 \left( \frac{2}{5} \right) = 0.971 \end{aligned}$$

## Information Gain for Outlook

Outlook	Play
Sunny	No
Sunny	No
Sunny	No
Sunny	Yes
Sunny	Yes

We have 2 Yes, 3

$$\begin{aligned} \text{No Entropy} = & \\ & -\frac{3}{5} \log_2 \left( \frac{3}{5} \right) - \\ & \frac{2}{5} \log_2 \left( \frac{2}{5} \right) = 0.971 \end{aligned}$$

## Information Gain for Outlook

Outlook	Play
Sunny	No
Sunny	No
Sunny	No
Sunny	Yes
Sunny	Yes

We have 2 Yes, 3

$$\begin{aligned} \text{No Entropy} = & \\ & -\frac{3}{5} \log_2 \left( \frac{3}{5} \right) - \\ & \frac{2}{5} \log_2 \left( \frac{2}{5} \right) = 0.971 \end{aligned}$$

Outlook	Play
Overcast	Yes
Overcast	Yes
Overcast	Yes
Overcast	Yes

We have 4 Yes, 0

No Entropy = 0  
(pure subset)

## Information Gain for Outlook

Outlook	Play
Sunny	No
Sunny	No
Sunny	No
Sunny	Yes
Sunny	Yes

We have 2 Yes, 3

$$\begin{aligned} \text{No Entropy} = & \\ & -\frac{3}{5} \log_2 \left( \frac{3}{5} \right) - \\ & \frac{2}{5} \log_2 \left( \frac{2}{5} \right) = 0.971 \end{aligned}$$

Outlook	Play
Overcast	Yes
Overcast	Yes
Overcast	Yes
Overcast	Yes

We have 4 Yes, 0

$$\begin{aligned} \text{No Entropy} = & 0 \\ & \text{(pure subset)} \end{aligned}$$

## Information Gain for Outlook

Outlook	Play
Sunny	No
Sunny	No
Sunny	No
Sunny	Yes
Sunny	Yes

We have 2 Yes, 3  
No Entropy =  
 $-\frac{3}{5} \log_2 \left(\frac{3}{5}\right) -$   
 $\frac{2}{5} \log_2 \left(\frac{2}{5}\right) = 0.971$

Outlook	Play
Overcast	Yes
Overcast	Yes
Overcast	Yes
Overcast	Yes

We have 4 Yes, 0  
No Entropy = 0  
(pure subset)

Outlook	Play
Rain	Yes
Rain	Yes
Rain	No
Rain	Yes
Rain	No

We have 3 Yes, 2  
No Entropy =  
 $-\frac{3}{5} \log_2 \left(\frac{3}{5}\right) -$   
 $\frac{2}{5} \log_2 \left(\frac{2}{5}\right) = 0.971$




# Information Gain

$$\text{Gain}(S, \text{Outlook}) = \text{Entropy}(S) - \sum_{v \in \{\text{Rain, Sunny, Overcast}\}} \frac{|S_v|}{|S|} \text{Entropy}(S_v)$$


$$\text{Gain}(S, \text{Outlook}) = \text{Entropy}(S) - \frac{5}{14} \text{Entropy}(S_{\text{Sunny}}) - \frac{4}{14} \text{Entropy}(S_{\text{Overcast}})$$

$$= 0.940 - \frac{5}{14} \times 0.971 - \frac{4}{14} \times 0 - \frac{5}{14} \times 0.971 = 0.940 - 0.347 - 0 - 0.347 = 0.2$$

# Information Gain



# Learnt Decision Tree



## Calling ID3 on Outlook=Sunny

Day	Temp	Humidity	Windy	Play
D1	Hot	High	Weak	No
D2	Hot	High	Strong	No
D8	Mild	High	Weak	No
D9	Cool	Normal	Weak	Yes
D11	Mild	Normal	Strong	Yes

## Calling ID3 on Outlook=Sunny

Day	Temp	Humidity	Windy	Play
D1	Hot	High	Weak	No
D2	Hot	High	Strong	No
D8	Mild	High	Weak	No
D9	Cool	Normal	Weak	Yes
D11	Mild	Normal	Strong	Yes

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$

## Calling ID3 on Outlook=Sunny

Day	Temp	Humidity	Windy	Play
D1	Hot	High	Weak	No
D2	Hot	High	Strong	No
D8	Mild	High	Weak	No
D9	Cool	Normal	Weak	Yes
D11	Mild	Normal	Strong	Yes

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$

## Calling ID3 on Outlook=Sunny

Day	Temp	Humidity	Windy	Play
D1	Hot	High	Weak	No
D2	Hot	High	Strong	No
D8	Mild	High	Weak	No
D9	Cool	Normal	Weak	Yes
D11	Mild	Normal	Strong	Yes

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Humidity}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(2 \text{ Yes}, 0 \text{ No}) - (3/5) * \text{Entropy}(0 \text{ Yes}, 3 \text{ No})$   
 $\Rightarrow$  **maximum possible for the set**

## Calling ID3 on Outlook=Sunny

Day	Temp	Humidity	Windy	Play
D1	Hot	High	Weak	No
D2	Hot	High	Strong	No
D8	Mild	High	Weak	No
D9	Cool	Normal	Weak	Yes
D11	Mild	Normal	Strong	Yes

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Humidity}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(2 \text{ Yes}, 0 \text{ No}) - (3/5) * \text{Entropy}(0 \text{ Yes}, 3 \text{ No})$   
 $\Rightarrow$  **maximum possible for the set**




## Calling ID3 on Outlook=Sunny

Day	Temp	Humidity	Windy	Play
D1	Hot	High	Weak	No
D2	Hot	High	Strong	No
D8	Mild	High	Weak	No
D9	Cool	Normal	Weak	Yes
D11	Mild	Normal	Strong	Yes

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Humidity}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(2 \text{ Yes}, 0 \text{ No}) - (3/5) * \text{Entropy}(0 \text{ Yes}, 3 \text{ No})$   
 $\implies$  **maximum possible for the set**
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Windy}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (3/5) * \text{Entropy}(1 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No})$

# Learnt Decision Tree




## Calling ID3 on (Outlook=Rain)

Day	Temp	Humidity	Windy	Play
D4	Mild	High	Weak	Yes
D5	Cool	Normal	Weak	Yes
D6	Cool	Normal	Strong	No
D10	Mild	Normal	Weak	Yes
D14	Mild	High	Strong	No


- The attribute Windy gives the highest information gain

# Learnt Decision Tree




# Prediction for Decision Tree

Just walk down the tree!



# Prediction for Decision Tree

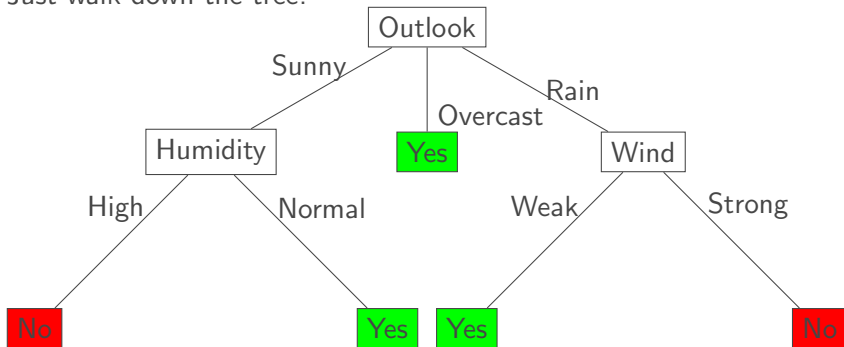
Just walk down the tree!



Prediction for <High Humidity, Strong Wind, Sunny Outlook, Hot Temp> is ?

# Prediction for Decision Tree

Just walk down the tree!



Prediction for <High Humidity, Strong Wind, Sunny Outlook, Hot Temp> is ?  
No

# Limiting Tree Depth

## Definition: Depth-Limited Trees

When depth limit is reached, assign the **most common class** in that path as the leaf node prediction.



# Limiting Tree Depth

## Definition: Depth-Limited Trees

When depth limit is reached, assign the **most common class** in that path as the leaf node prediction.

- **Depth-0 tree** (no decisions):

# Limiting Tree Depth

## Definition: Depth-Limited Trees

When depth limit is reached, assign the **most common class** in that path as the leaf node prediction.

- **Depth-0 tree** (no decisions):
  - Always predict the most common class

# Limiting Tree Depth

## Definition: Depth-Limited Trees

When depth limit is reached, assign the **most common class** in that path as the leaf node prediction.

- **Depth-0 tree** (no decisions):
  - Always predict the most common class
  - For our dataset: Always predict **Yes**

# Limiting Tree Depth

## Definition: Depth-Limited Trees

When depth limit is reached, assign the **most common class** in that path as the leaf node prediction.


- **Depth-0 tree** (no decisions):
  - Always predict the most common class
  - For our dataset: Always predict **Yes**

# Limiting Tree Depth

## Definition: Depth-Limited Trees

When depth limit is reached, assign the **most common class** in that path as the leaf node prediction.

- **Depth-0 tree** (no decisions):
  - Always predict the most common class
  - For our dataset: Always predict **Yes**
- **Depth-1 tree** (single decision):



## Pop Quiz #7

### Quick Question!

In the tennis dataset, why did "Outlook" have the highest information gain?

A) It was the first feature in the dataset

## Pop Quiz #8

### Quick Question!

In the tennis dataset, why did "Outlook" have the highest information gain?

- A) It was the first feature in the dataset
- B) When Outlook=Overcast, all examples have Play=Yes (pure subset)

## Pop Quiz #9

### Quick Question!

In the tennis dataset, why did "Outlook" have the highest information gain?

- A) It was the first feature in the dataset
- B) When Outlook=Overcast, all examples have Play=Yes (pure subset)
- C) It has the most possible values



## Pop Quiz #10

### Quick Question!

In the tennis dataset, why did "Outlook" have the highest information gain?

- A) It was the first feature in the dataset
- B) When Outlook=Overcast, all examples have Play=Yes (pure subset)
- C) It has the most possible values
- D) It was chosen randomly

## Pop Quiz #11

### Quick Question!

In the tennis dataset, why did "Outlook" have the highest information gain?

- A) It was the first feature in the dataset
- B) When Outlook=Overcast, all examples have Play=Yes (pure subset)
- C) It has the most possible values
- D) It was chosen randomly

## Pop Quiz #12

### Quick Question!

In the tennis dataset, why did "Outlook" have the highest information gain?

- A) It was the first feature in the dataset
- B) When Outlook=Overcast, all examples have Play=Yes (pure subset)
- C) It has the most possible values
- D) It was chosen randomly

**Answer: B) When Outlook=Overcast, all examples have Play=Yes** - This creates a pure subset with entropy=0, maximizing information gain.

## Modified Dataset

Day	Outlook	Temp	Humidity	Wind	Minutes Played
D1	Sunny	Hot	High	Weak	20
D2	Sunny	Hot	High	Strong	24
D3	Overcast	Hot	High	Weak	40
D4	Rain	Mild	High	Weak	50
D5	Rain	Cool	Normal	Weak	60
D6	Rain	Cool	Normal	Strong	10
D7	Overcast	Cool	Normal	Strong	4
D8	Sunny	Mild	High	Weak	10
D9	Sunny	Cool	Normal	Weak	60
D10	Rain	Mild	Normal	Weak	40
D11	Sunny	Mild	High	Strong	45
D12	Overcast	Mild	High	Strong	40
D13	Overcast	Hot	Normal	Weak	35
D14	Rain	Mild	High	Strong	20

# Measure of Impurity for Regression?

# Measure of Impurity for Regression?

# Measure of Impurity for Regression?

- Any guesses?

# Measure of Impurity for Regression?

- Any guesses?



# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error

# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error

# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$

# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$

# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?

# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?

# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?
- **MSE Reduction** (not Information Gain!)

# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error
- $MSE(S) = 311.34$
- What about splitting criterion for regression?
- **MSE Reduction** (not Information Gain!)



# Measure of Impurity for Regression?

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?
- **MSE Reduction** (not Information Gain!)
- $\text{MSE Reduction} = \text{MSE}(S) - \sum_v \frac{|S_v|}{|S|} \text{MSE}(S_v)$

## Gain by splitting on Wind

## Gain by splitting on Wind

Wind	Minutes Played
Weak	20
Strong	24
Weak	40
Weak	50
Weak	60
Strong	10
Strong	4
Weak	10
Weak	60
Weak	40
Strong	45
Strong	40
Weak	35
Strong	20

$$\text{MSE}(S)=311.34$$

# Gain by splitting on Wind

Wind	Minutes Played
Weak	20
Strong	24
Weak	40
Weak	50
Weak	60
Strong	10
Strong	4
Weak	10
Weak	60
Weak	40
Strong	45
Strong	40
Weak	35
Strong	20

$$\text{MSE}(S) = 311.34$$

Wind	Minutes Played
Weak	20
Weak	40
Weak	50
Weak	60
Weak	10
Weak	60
Weak	40
Weak	35

$$\text{MSE}(S_{\text{Wind=Weak}}) = 277, \text{ Weight} = \frac{8}{14}$$

Wind	Minutes Played
Strong	24
Strong	10
Strong	4
Strong	45
Strong	40
Strong	20

$$\text{MSE}(S_{\text{Wind=Strong}}) = 218, \text{ Weight} = \frac{6}{14}$$

# MSE Reduction Calculation

**Correct calculation for Wind split:**

$$\begin{aligned}\text{MSE Reduction} &= \text{MSE}(S) - \text{Weighted Average MSE} \\ &= 311.34 - \left[ \frac{8}{14} \times 277 + \frac{6}{14} \times 218 \right] = 311.34 - [158.857 + 93.429] = 311.34 - 252.286 = 59.054\end{aligned}$$

**Key insight:** MSE Reduction  $> 0$  means the split improves our model!

**For regression:** Use MSE Reduction, NOT Information Gain!

## Pop Quiz #13

### Quick Question!

For regression trees, what criterion do we use instead of Information Gain?

A) Information Gain

## Pop Quiz #14

### Quick Question!

For regression trees, what criterion do we use instead of Information Gain?

- A) Information Gain
- B) Gini Impurity

## Pop Quiz #15

### Quick Question!

For regression trees, what criterion do we use instead of Information Gain?

- A) Information Gain
- B) Gini Impurity
- C) Mean Squared Error (MSE) Reduction



## Pop Quiz #16

### Quick Question!

For regression trees, what criterion do we use instead of Information Gain?

- A) Information Gain
- B) Gini Impurity
- C) Mean Squared Error (MSE) Reduction
- D) Accuracy

## Pop Quiz #17

### Quick Question!

For regression trees, what criterion do we use instead of Information Gain?

- A) Information Gain
- B) Gini Impurity
- C) Mean Squared Error (MSE) Reduction
- D) Accuracy

## Pop Quiz #18

### Quick Question!


For regression trees, what criterion do we use instead of Information Gain?

- A) Information Gain
- B) Gini Impurity
- C) Mean Squared Error (MSE) Reduction
- D) Accuracy

**Answer: C) Mean Squared Error (MSE) Reduction** -  
For regression, we minimize MSE instead of maximizing information gain.


# MSE Reduction for Regression Trees

**Notebook:** decision-tree-real-output.html



# Learnt Tree

Assume a tree  
like this is  
learnt ...



	Day	Outlook	Temp	Humidity	Wind	Minutes Played
2	D3	Overcast	Hot	High	Weak	40
12	D13	Overcast	Hot	Normal	Weak	35

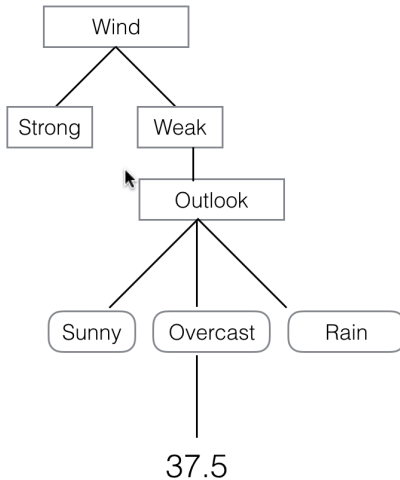
# Learnt Tree

## Method 1

Mins

Played=(40+35)

/2



## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- How do you find splits?

# Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- How do you find splits?
- Sort by attribute



## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- How do you find splits?
- Sort by attribute
- Find potential split points (midpoints).

# Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- How do you find splits?
- Sort by attribute
- Find potential split points (midpoints).
- For the above example, we have 5 potential splits: 44, 54, 66, 76, 85

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- How do you find splits?
- Sort by attribute
- Find potential split points (midpoints).
- For the above example, we have 5 potential splits: 44, 54, 66, 76, 85
- Calculate the weighted impurity for each split

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- How do you find splits?
- Sort by attribute
- Find potential split points (midpoints).
- For the above example, we have 5 potential splits: 44, 54, 66, 76, 85
- Calculate the weighted impurity for each split
- Choose the split with the lowest impurity

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 44

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 44
- LHS has 1 No and 0 Yes; RHS has 3 Yes and 2 No

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 44
- LHS has 1 No and 0 Yes; RHS has 3 Yes and 2 No
- Entropy for LHS = 0, Entropy for RHS = 0.971

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 44
- LHS has 1 No and 0 Yes; RHS has 3 Yes and 2 No
- Entropy for LHS = 0, Entropy for RHS = 0.971
- Weighted Entropy =  $0.971 * 5/6 = 0.808$



## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 54

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 54
- LHS has 2 No and 0 Yes; RHS has 3 Yes and 1 No

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 54
- LHS has 2 No and 0 Yes; RHS has 3 Yes and 1 No
- Entropy for LHS = 0, Entropy for RHS = 0.811

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 54
- LHS has 2 No and 0 Yes; RHS has 3 Yes and 1 No
- Entropy for LHS = 0, Entropy for RHS = 0.811
- Weighted Entropy =  $0.811 * 4/6 = 0.541$

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 66

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 66
- LHS has 2 No and 1 Yes; RHS has 2 Yes and 1 No

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 66
- LHS has 2 No and 1 Yes; RHS has 2 Yes and 1 No
- Entropy for LHS = 0.918, Entropy for RHS = 0.918

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 66
- LHS has 2 No and 1 Yes; RHS has 2 Yes and 1 No
- Entropy for LHS = 0.918, Entropy for RHS = 0.918
- Weighted Entropy =  $0.918 \cdot 3/6 + 0.918 \cdot 3/6 = 0.918$



## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 76

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 76
- LHS has 2 No and 2 Yes; RHS has 1 Yes and 1 No

## Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 76
- LHS has 2 No and 2 Yes; RHS has 1 Yes and 1 No
- Entropy for LHS = 1, Entropy for RHS = 1

## Finding splits


Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

- Consider split at 76
- LHS has 2 No and 2 Yes; RHS has 1 Yes and 1 No
- Entropy for LHS = 1, Entropy for RHS = 1
- Weighted Entropy =  $1*4/6 + 1*2/6 = 1$

# Finding splits

Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No


**Notebook:** [decision-tree-real-input-discrete-output.html](#)



# Finding splits


Day	Temperature	PlayTennis
D1	40	No
D2	48	No
D3	60	Yes
D4	72	Yes
D5	80	Yes
D6	90	No

**Notebook:** [decision-tree-real-input-discrete-output.html](#)




## Example (DT of depth 1)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)



## Example (DT of depth 2)


**Notebook:** [decision-tree-real-input-discrete-output.html](#)






## Example (DT of depth 3)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)




## Example (DT of depth 4)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)




## Example (DT of depth 5)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)




## Example (DT of depth 6)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)




## Example (DT of depth 7)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)




## Example (DT of depth 8)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)



## Example (DT of depth 9)

**Notebook:** [decision-tree-real-input-discrete-output.html](#)



## Pop Quiz #19

### Quick Question!

When finding splits for continuous features, how do we determine candidate split points?

A) Use all feature values as split points



## Pop Quiz #20

### Quick Question!

When finding splits for continuous features, how do we determine candidate split points?

- A) Use all feature values as split points
- B) Use midpoints between consecutive sorted feature values

## Pop Quiz #21

### Quick Question!

When finding splits for continuous features, how do we determine candidate split points?

- A) Use all feature values as split points
- B) Use midpoints between consecutive sorted feature values
- C) Use random values within the feature range

## Pop Quiz #22

### Quick Question!

When finding splits for continuous features, how do we determine candidate split points?

- A) Use all feature values as split points
- B) Use midpoints between consecutive sorted feature values
- C) Use random values within the feature range
- D) Use only the minimum and maximum values

## Pop Quiz #23

### Quick Question!

When finding splits for continuous features, how do we determine candidate split points?

- A) Use all feature values as split points
- B) Use midpoints between consecutive sorted feature values
- C) Use random values within the feature range
- D) Use only the minimum and maximum values

## Pop Quiz #24

### Quick Question!

When finding splits for continuous features, how do we determine candidate split points?


- A) Use all feature values as split points
- B) Use midpoints between consecutive sorted feature values
- C) Use random values within the feature range
- D) Use only the minimum and maximum values

**Answer: B) Use midpoints between consecutive sorted feature values** - This ensures we test all meaningful boundaries between different class regions.

# Example 1

Let us consider the dataset given below


**Notebook:** [decision-tree-real-input-real-output.html](#)



## Example 1

What would be the prediction for decision tree with depth 0?

**Notebook:** [decision-tree-real-input-real-output.html](#)




# Example 1

Prediction for decision tree with depth 0.

Horizontal dashed line shows the predicted  $Y$  value. It is the average of  $Y$  values of all datapoints.

**Notebook:** [decision-tree-real-input-real-output.html](#)






# Example 1

What would be the decision tree with depth 1?


**Notebook:** [decision-tree-real-input-real-output.html](#)



# Example 1

Decision tree with depth 1


**Notebook:** [decision-tree-real-input-real-output.html](#)



# Example 1

## The Decision Boundary


**Notebook:** [decision-tree-real-input-real-output.html](#)



# Example 1

What would be the decision tree with depth 2 ?


**Notebook:** [decision-tree-real-input-real-output.html](#)



# Example 1

Decision tree with depth 2

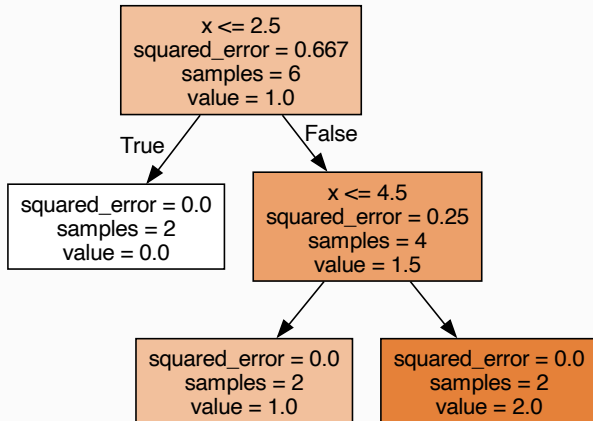
**Notebook:** [decision-tree-real-input-real-output.html](#)



# Example 1

## The Decision Boundary

**Notebook:** [decision-tree-real-input-real-output.html](#)



# Objective Function for Regression Trees

Feature is denoted by  $X$  and target by  $Y$ .

Let the split be at  $X = s$ .

Define regions:  $R_1 = \{x : x \leq s\}$  and  $R_2 = \{x : x > s\}$ .

# Objective Function for Regression Trees

Feature is denoted by  $X$  and target by  $Y$ .

Let the split be at  $X = s$ .

Define regions:  $R_1 = \{x : x \leq s\}$  and  $R_2 = \{x : x > s\}$ .

For each region, compute the mean prediction:

$$c_1 = \frac{1}{|R_1|} \sum_{x_i \in R_1} y_i$$

$$c_2 = \frac{1}{|R_2|} \sum_{x_i \in R_2} y_i$$



# Objective Function for Regression Trees

Feature is denoted by  $X$  and target by  $Y$ .

Let the split be at  $X = s$ .

Define regions:  $R_1 = \{x : x \leq s\}$  and  $R_2 = \{x : x > s\}$ .

For each region, compute the mean prediction:

$$c_1 = \frac{1}{|R_1|} \sum_{x_i \in R_1} y_i$$

$$c_2 = \frac{1}{|R_2|} \sum_{x_i \in R_2} y_i$$

The loss function is:

$$\text{Loss}(s) = \sum_{x_i \in R_1} (y_i - c_1)^2 + \sum_{x_i \in R_2} (y_i - c_2)^2$$

# Objective Function for Regression Trees

Feature is denoted by  $X$  and target by  $Y$ .

Let the split be at  $X = s$ .

Define regions:  $R_1 = \{x : x \leq s\}$  and  $R_2 = \{x : x > s\}$ .

For each region, compute the mean prediction:

$$c_1 = \frac{1}{|R_1|} \sum_{x_i \in R_1} y_i$$

$$c_2 = \frac{1}{|R_2|} \sum_{x_i \in R_2} y_i$$

The loss function is:

$$\text{Loss}(s) = \sum_{x_i \in R_1} (y_i - c_1)^2 + \sum_{x_i \in R_2} (y_i - c_2)^2$$

Our objective is to find the optimal split:

## Algorithm: Finding the Optimal Split

1. Sort all data points  $(x_i, y_i)$  in increasing order of  $x_i$ .

## Algorithm: Finding the Optimal Split

1. Sort all data points  $(x_i, y_i)$  in increasing order of  $x_i$ .
2. Evaluate the loss function for all candidate splits:

$$s = \frac{x_i + x_{i+1}}{2} \text{ for } i = 1, 2, \dots, n - 1$$

3. Select the split  $s^*$  that minimizes the loss function.


## A Question!

Draw a regression tree for  $Y = \sin(X)$ ,  $0 \leq X \leq 2\pi$

## A Question!

Dataset of  $Y = \sin(X)$ ,  $0 \leq X \leq 7$  with 10,000 points

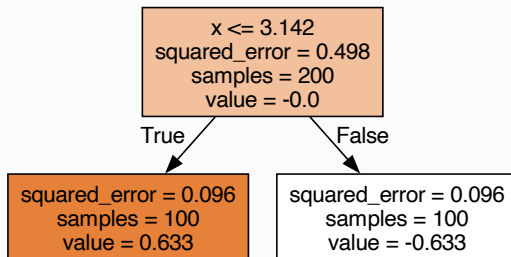
**Notebook:** [decision-tree-real-input-real-output.html](#)



# A Question!

Regression tree of depth 1


**Notebook:** [decision-tree-real-input-real-output.html](#)



# A Question!

Decision Boundary

**Notebook:** [decision-tree-real-input-real-output.html](#)






## A Question!

Regression tree with no depth limit is too big to fit in a slide. It has of depth 4. The decision boundaries are in figure below.

**Notebook:** [decision-tree-real-input-real-output.html](#)



## Pop Quiz #25

### Quick Question!

What is the prediction function for a regression tree leaf node?

- A) The median of target values in that region

## Pop Quiz #26

### Quick Question!

What is the prediction function for a regression tree leaf node?

- A) The median of target values in that region
- B) The mode of target values in that region

## Pop Quiz #27

### Quick Question!

What is the prediction function for a regression tree leaf node?

- A) The median of target values in that region
- B) The mode of target values in that region
- C) The mean of target values in that region

## Pop Quiz #28

### Quick Question!

What is the prediction function for a regression tree leaf node?

- A) The median of target values in that region
- B) The mode of target values in that region
- C) The mean of target values in that region
- D) A linear function of the features

## Pop Quiz #29

### Quick Question!

What is the prediction function for a regression tree leaf node?

- A) The median of target values in that region
- B) The mode of target values in that region
- C) The mean of target values in that region
- D) A linear function of the features

## Pop Quiz #30

### Quick Question!

What is the prediction function for a regression tree leaf node?

- A) The median of target values in that region
- B) The mode of target values in that region
- C) The mean of target values in that region
- D) A linear function of the features

**Answer: C) The mean of target values in that region**

- Each leaf predicts the average target value of training samples that reach that leaf.

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex



# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy
  - Very deep trees with many leaves

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy
  - Very deep trees with many leaves
  - Rules that are too specific to training data

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy
  - Very deep trees with many leaves
  - Rules that are too specific to training data
- **Solution:** Pruning to control model complexity

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)

**Advantages:** Simple, computationally efficient

**Disadvantages:** May stop too early, miss good splits later



# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $j$   $N$  samples

**Advantages:** Simple, computationally efficient

**Disadvantages:** May stop too early, miss good splits later

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $j$   $N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples

**Advantages:** Simple, computationally efficient

**Disadvantages:** May stop too early, miss good splits later

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $j$   $N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples
- **Maximum features:** Consider only subset of features at each split

**Advantages:** Simple, computationally efficient

**Disadvantages:** May stop too early, miss good splits later

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples
- **Maximum features:** Consider only subset of features at each split
- **Minimum impurity decrease:** Only split if improvement  $\geq$  threshold

**Advantages:** Simple, computationally efficient

**Disadvantages:** May stop too early, miss good splits later

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance
3. Remove branches that don't improve validation accuracy



# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance
3. Remove branches that don't improve validation accuracy
4. Repeat until no beneficial removals remain

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data
  2. Use validation set to evaluate subtree performance
  3. Remove branches that don't improve validation accuracy
  4. Repeat until no beneficial removals remain
- **Cost Complexity Pruning:** Minimize  $\text{Error} + \alpha \times \text{Tree Size}$

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data
  2. Use validation set to evaluate subtree performance
  3. Remove branches that don't improve validation accuracy
  4. Repeat until no beneficial removals remain
- **Cost Complexity Pruning:** Minimize Error +  $\alpha \times \text{Tree Size}$
- **Advantages:** More thorough, can recover from early stopping mistakes

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data
  2. Use validation set to evaluate subtree performance
  3. Remove branches that don't improve validation accuracy
  4. Repeat until no beneficial removals remain
- **Cost Complexity Pruning:** Minimize Error +  $\alpha \times \text{Tree Size}$
- **Advantages:** More thorough, can recover from early stopping mistakes
- **Disadvantages:** More computationally expensive

# Cost Complexity Pruning Algorithm

**Systematic approach to find optimal tree size:**

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$

# Cost Complexity Pruning Algorithm

**Systematic approach to find optimal tree size:**

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set

# Cost Complexity Pruning Algorithm

**Systematic approach to find optimal tree size:**

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)

# Cost Complexity Pruning Algorithm

**Systematic approach to find optimal tree size:**

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)



# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )
  2. Gradually increase  $\alpha$

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )
  2. Gradually increase  $\alpha$
  3. At each  $\alpha$ , prune branches that increase cost

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )
  2. Gradually increase  $\alpha$
  3. At each  $\alpha$ , prune branches that increase cost
  4. Select  $\alpha$  with best cross-validation performance

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)



# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)
  - Risk of underfitting

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)
  - Risk of underfitting
- **Optimal pruning:** Balances bias and variance

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)
  - Risk of underfitting
- **Optimal pruning:** Balances bias and variance
- **Cross-validation:** Essential for finding this balance

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters



# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters** (sklearn):

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100
  - `min_samples_leaf`: Try 5-50

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100
  - `min_samples_leaf`: Try 5-50
  - `ccp_alpha`: Use for cost complexity pruning

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100
  - `min_samples_leaf`: Try 5-50
  - `ccp_alpha`: Use for cost complexity pruning
- **Domain knowledge:** Consider interpretability requirements



# Summary

- Interpretability an important goal

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”
- Issues:


# Summary


- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”
- Issues:
  - Can overfit easily!


# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”
- Issues:
  - Can overfit easily!
  - Empirically not as powerful as other methods











$$\text{Entropy} = -P(+)\log_2 P(+)-P(-)\log_2 P(-)$$

$$P(+)=\frac{0.1+0.1+0.3}{1}=0.5,\quad P(-)=\frac{0.3+0.1+0.1}{1}=0.5$$




Candidate Line:  $X_2 = 4(X_1^*)$



Entropy of  $X1 \leq X1^* = E_{S(X1 < X1^*)}$

$$P(+) = \frac{0.1 + 0.1}{0.1 + 0.1 + 0.3} = \frac{2}{5}$$


$$P(-) = \frac{3}{5}$$



Entropy of  $X_1 > X_1^* = E_S(X_1 > X_1^*)$

$$P(+)=\frac{3}{5}$$

$$P(-)=\frac{2}{5}$$



$$\text{IG}(X_1 = X_1^*) = E_S - \frac{0.5}{1} \cdot E_{S(X_1 < X_1^*)} - \frac{0.5}{1} \cdot E_{S(X_1 > X_1^*)}$$