

# Decision Trees

---

Nipun Batra and teaching staff

IIT Gandhinagar

July 30, 2025

# Table of Contents

1. Discrete Input, Real Output
2. Real Input Real Output
3. Pruning and Overfitting
4. Summary and Key Takeaways
5. Weighted Entropy

# Information Gain Intuition

- For examples, we have 9 Yes, 5 No

# Information Gain Intuition

- For examples, we have 9 Yes, 5 No

# Information Gain Intuition

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?

# Information Gain Intuition

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?

# Information Gain Intuition

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!

# Information Gain Intuition

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!



# Information Gain Intuition

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!
- Key insight: Problem is “easier” when there is less disagreement

# Information Gain Intuition

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!
- Key insight: Problem is “easier” when there is less disagreement


# Information Gain Intuition

- For examples, we have 9 Yes, 5 No
- Would it be trivial if we had 14 Yes or 14 No?
- Yes!
- Key insight: Problem is “easier” when there is less disagreement
- Need some statistical measure of “disagreement”

# Entropy Formula

$$H(X) = - \sum_{i=1}^k p(x_i) \log_2 p(x_i)$$

**Notebook:** entropy.html



# Root Node Selection

- Can we use Outlook as the root node?

# Root Node Selection

- Can we use Outlook as the root node?

# Root Node Selection

- Can we use Outlook as the root node?
- When Outlook is overcast, we always Play and thus no "disagreement"

# What Does Entropy Measure?

**Answer: B) The impurity or “disagreement” in a set of examples** - Higher entropy means more mixed classes, lower entropy means more pure subsets.



# Entropy Calculation Examples

# Entropy Calculation Examples

Outlook	Play
Overcast	Yes
Overcast	Yes
Overcast	Yes
Overcast	Yes

We have 4 Yes, 0

No Entropy = 0  
(pure subset)

# Entropy Calculation Examples

Outlook	Play
Overcast	Yes
Overcast	Yes
Overcast	Yes
Overcast	Yes

We have 4 Yes, 0

No Entropy = 0  
(pure subset)

# Entropy Calculation Examples

Outlook	Play
Overcast	Yes
Overcast	Yes
Overcast	Yes
Overcast	Yes

We have 4 Yes, 0

No Entropy = 0  
(pure subset)

Outlook	Play
Rain	Yes
Rain	Yes
Rain	No
Rain	Yes
Rain	No

We have 3 Yes, 2

No Entropy =  
 $-\frac{3}{5} \log_2 \left(\frac{3}{5}\right) -$   
 $\frac{2}{5} \log_2 \left(\frac{2}{5}\right) = 0.971$

# Information Gain Calculations

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$

# Information Gain Calculations

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$

# Information Gain Calculations

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Humidity}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(2 \text{ Yes}, 0 \text{ No}) - (3/5) * \text{Entropy}(0 \text{ Yes}, 3 \text{ No}) \implies \textbf{maximum possible for the set}$

# Information Gain Calculations

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Humidity}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(2 \text{ Yes}, 0 \text{ No}) - (3/5) * \text{Entropy}(0 \text{ Yes}, 3 \text{ No}) \implies \textbf{maximum possible for the set}$



# Information Gain Calculations

- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Temp}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(0 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No}) - (1/5) * \text{Entropy}(1 \text{ Yes}, 0 \text{ No})$
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Humidity}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (2/5) * \text{Entropy}(2 \text{ Yes}, 0 \text{ No}) - (3/5) * \text{Entropy}(0 \text{ Yes}, 3 \text{ No}) \implies \textbf{maximum possible for the set}$
- $\text{Gain}(S_{\text{Outlook}=\text{Sunny}}, \text{Windy}) = \text{Entropy}(2 \text{ Yes}, 3 \text{ No}) - (3/5) * \text{Entropy}(1 \text{ Yes}, 2 \text{ No}) - (2/5) * \text{Entropy}(1 \text{ Yes}, 1 \text{ No})$

# Prediction Example

Prediction for <High Humidity, Strong Wind, Sunny Outlook, Hot Temp> is ?

# Prediction Example

Prediction for <High Humidity, Strong Wind, Sunny Outlook, Hot Temp> is ?  
No

# Depth-Limited Trees

Apply the same rules, except when depth limit is reached, the leaf node is assigned the most common occurring value in that path.

# Depth-Limited Trees

Apply the same rules, except when depth limit is reached, the leaf node is assigned the most common occurring value in that path.

What is depth-0 tree (no decision) for the examples?

# Depth-Limited Trees

Apply the same rules, except when depth limit is reached, the leaf node is assigned the most common occurring value in that path.

What is depth-0 tree (no decision) for the examples?

Always predicting Yes

# Depth-Limited Trees

Apply the same rules, except when depth limit is reached, the leaf node is assigned the most common occurring value in that path.

What is depth-0 tree (no decision) for the examples?

Always predicting Yes

What is depth-1 tree (no decision) for the examples?


# Depth-Limited Trees

Apply the same rules, except when depth limit is reached, the leaf node is assigned the most common occurring value in that path.

What is depth-0 tree (no decision) for the examples?

Always predicting Yes

What is depth-1 tree (no decision) for the examples?





# Why Outlook is Good Root?

**Answer: B) When Outlook=Overcast, all examples have Play=Yes** - This creates a pure subset with entropy=0, maximizing information gain.

# Regression Trees

- Any guesses?

# Regression Trees

- Any guesses?

# Regression Trees

- Any guesses?

# Regression Trees

- Any guesses?
- Mean Squared Error

# Regression Trees

- Any guesses?
- Mean Squared Error

# Regression Trees

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$

# Regression Trees

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$



# Regression Trees

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?

# Regression Trees

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?

# Regression Trees

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?
- **MSE Reduction** (not Information Gain!)

# Regression Trees

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?
- **MSE Reduction** (not Information Gain!)

# Regression Trees

- Any guesses?
- Mean Squared Error
- $\text{MSE}(S) = 311.34$
- What about splitting criterion for regression?
- **MSE Reduction** (not Information Gain!)
- $\text{MSE Reduction} = \text{MSE}(S) - \sum_v \frac{|S_v|}{|S|} \text{MSE}(S_v)$

# Regression Splitting Criterion

**Answer: C) Mean Squared Error (MSE) Reduction** - For regression, we minimize MSE instead of maximizing information gain.

# Continuous Features

**Answer: B) Use midpoints between consecutive sorted feature values** - This ensures we test all meaningful boundaries between different class regions.

# Leaf Node Predictions

**Answer: C) The mean of target values in that region -**  
Each leaf predicts the average target value of training samples that reach that leaf.



# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy



# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy
  - Very deep trees with many leaves

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy
  - Very deep trees with many leaves
  - Rules that are too specific to training data

# The Problem: Overfitting in Decision Trees

- **Unpruned trees:** Can grow very deep and complex
- **Perfect training accuracy:** Each leaf contains single training example
- **But:** Poor generalization to new data
- **Symptoms:**
  - High training accuracy, low test accuracy
  - Very deep trees with many leaves
  - Rules that are too specific to training data
- **Solution:** Pruning to control model complexity

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples



# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples
- **Maximum features:** Consider only subset of features at each split

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples
- **Maximum features:** Consider only subset of features at each split

# Pre-pruning (Early Stopping)

**Stop growing tree before it becomes too complex:**

- **Maximum depth:** Limit tree depth (e.g., `max_depth = 5`)
- **Minimum samples per split:** Don't split if node has  $< N$  samples
- **Minimum samples per leaf:** Ensure each leaf has  $\geq M$  samples
- **Maximum features:** Consider only subset of features at each split
- **Minimum impurity decrease:** Only split if improvement  $>$  threshold

**Advantages:** Simple, computationally efficient

**Disadvantages:** May stop too early, miss good splits later

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data



# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data
  2. Use validation set to evaluate subtree performance

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance
3. Remove branches that don't improve validation accuracy

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance
3. Remove branches that don't improve validation accuracy
4. Repeat until no beneficial removals remain

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance
3. Remove branches that don't improve validation accuracy
4. Repeat until no beneficial removals remain

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance
3. Remove branches that don't improve validation accuracy
4. Repeat until no beneficial removals remain

- **Cost Complexity Pruning:** Minimize

Error +  $\alpha \times$  Tree Size

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**

1. Grow complete tree on training data
2. Use validation set to evaluate subtree performance
3. Remove branches that don't improve validation accuracy
4. Repeat until no beneficial removals remain

- **Cost Complexity Pruning:** Minimize

Error +  $\alpha \times$  Tree Size

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data
  2. Use validation set to evaluate subtree performance
  3. Remove branches that don't improve validation accuracy
  4. Repeat until no beneficial removals remain
- **Cost Complexity Pruning:** Minimize  
 $\text{Error} + \alpha \times \text{Tree Size}$
- **Advantages:** More thorough, can recover from early stopping mistakes

# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data
  2. Use validation set to evaluate subtree performance
  3. Remove branches that don't improve validation accuracy
  4. Repeat until no beneficial removals remain
- **Cost Complexity Pruning:** Minimize  
 $\text{Error} + \alpha \times \text{Tree Size}$
- **Advantages:** More thorough, can recover from early stopping mistakes



# Post-pruning (Tree Simplification)

**Grow full tree, then remove unnecessary branches:**

- **Algorithm:**
  1. Grow complete tree on training data
  2. Use validation set to evaluate subtree performance
  3. Remove branches that don't improve validation accuracy
  4. Repeat until no beneficial removals remain
- **Cost Complexity Pruning:** Minimize  
 $\text{Error} + \alpha \times \text{Tree Size}$
- **Advantages:** More thorough, can recover from early stopping mistakes
- **Disadvantages:** More computationally expensive

# Cost Complexity Pruning Algorithm

**Systematic approach to find optimal tree size:**

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$

# Cost Complexity Pruning Algorithm

**Systematic approach to find optimal tree size:**

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set

# Cost Complexity Pruning Algorithm

**Systematic approach to find optimal tree size:**

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )



# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )
  2. Gradually increase  $\alpha$

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )
  2. Gradually increase  $\alpha$
  3. At each  $\alpha$ , prune branches that increase cost

# Cost Complexity Pruning Algorithm

## Systematic approach to find optimal tree size:

- **Cost function:**  $R_{\alpha}(T) = R(T) + \alpha|T|$ 
  - $R(T)$ : Misclassification error on validation set
  - $|T|$ : Number of terminal nodes (tree size)
  - $\alpha$ : Complexity parameter (penalty for larger trees)
- **Process:**
  1. Start with full tree ( $\alpha = 0$ )
  2. Gradually increase  $\alpha$
  3. At each  $\alpha$ , prune branches that increase cost
  4. Select  $\alpha$  with best cross-validation performance

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting



# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)
  - Risk of underfitting

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)
  - Risk of underfitting
- **Optimal pruning:** Balances bias and variance

# Bias-Variance Trade-off in Trees

- **Unpruned trees:**
  - Low bias (can fit complex patterns)
  - High variance (sensitive to training data changes)
  - Prone to overfitting
- **Heavily pruned trees:**
  - High bias (may miss important patterns)
  - Low variance (more stable predictions)
  - Risk of underfitting
- **Optimal pruning:** Balances bias and variance
- **Cross-validation:** Essential for finding this balance

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters



# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters** (sklearn):

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters** (sklearn):
  - `max_depth`: Start with 3-10

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters** (sklearn):
  - `max_depth`: Start with 3-10



# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100
  - `min_samples_leaf`: Try 5-50

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100
  - `min_samples_leaf`: Try 5-50
  - `ccp_alpha`: Use for cost complexity pruning

# Practical Pruning Guidelines

- **Start simple:** Begin with restrictive pre-pruning parameters
- **Cross-validation:** Always use CV to select pruning parameters
- **Validation curves:** Plot training/validation error vs. tree complexity
- **Common parameters (sklearn):**
  - `max_depth`: Start with 3-10
  - `min_samples_split`: Try 10-100
  - `min_samples_leaf`: Try 5-50
  - `ccp_alpha`: Use for cost complexity pruning
- **Domain knowledge:** Consider interpretability requirements

# Summary

- Interpretability an important goal

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard



# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”

# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”

# Summary


- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”
- Issues:


# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”
- Issues:
  - Can overfit easily!


# Summary

- Interpretability an important goal
- Decision trees: well known interpretable models
- Learning optimal tree is hard
- Greedy approach:
- Recursively split to maximize “performance gain”
- Issues:
  - Can overfit easily!
  - Empirically not as powerful as other methods











$$\text{Entropy} = -P(+) \log_2 P(+) - P(-) \log_2 P(-)$$

$$P(+) = \frac{0.1 + 0.1 + 0.3}{1} = 0.5, \quad P(-) = \frac{0.3 + 0.1 + 0.1}{1} = 0.5$$




Candidate Line:  $X_1 = 4(X_1^*)$



Entropy of  $X_1 \leq X_1^* = E_{S(X_1 < X_1^*)}$

$$P(+)=\frac{0.1+0.1}{0.1+0.1+0.3}=\frac{2}{5}$$


$$P(-)=\frac{3}{5}$$



Entropy of  $X_1 > X_1^* = E_{S(X_1 > X_1^*)}$

$$P(+) = \frac{3}{5}$$

$$P(-) = \frac{2}{5}$$



$$\text{IG}(X_1 = X_1^*) = E_S - \frac{0.5}{1} \cdot E_{S(X_1 < X_1^*)} - \frac{0.5}{1} \cdot E_{S(X_1 > X_1^*)}$$