

Matrix Factorization for Movie Recommendation Systems

Nipun Batra

IIT Gandhinagar

August 1, 2025

Today's Learning Journey

- **The Problem:** Why do we need recommendation systems?

Today's Learning Journey

- **The Problem:** Why do we need recommendation systems?
- **Matrix View:** How ratings become a mathematical problem

Today's Learning Journey

- **The Problem:** Why do we need recommendation systems?
- **Matrix View:** How ratings become a mathematical problem
- **Key Insight:** Matrix factorization as the solution

Today's Learning Journey

- **The Problem:** Why do we need recommendation systems?
- **Matrix View:** How ratings become a mathematical problem
- **Key Insight:** Matrix factorization as the solution
- **Step-by-Step:** Building intuition with examples

Today's Learning Journey

- **The Problem:** Why do we need recommendation systems?
- **Matrix View:** How ratings become a mathematical problem
- **Key Insight:** Matrix factorization as the solution
- **Step-by-Step:** Building intuition with examples
- **Algorithms:** ALS vs Gradient Descent

Today's Learning Journey

- **The Problem:** Why do we need recommendation systems?
- **Matrix View:** How ratings become a mathematical problem
- **Key Insight:** Matrix factorization as the solution
- **Step-by-Step:** Building intuition with examples
- **Algorithms:** ALS vs Gradient Descent
- **Practice:** Hands-on understanding

The Movie Recommendation Challenge

The Movie Recommendation Challenge

Real-World Scenario:

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users,
15K+ titles

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products
- Spotify: 400M+ users, 70M+ songs

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products
- Spotify: 400M+ users, 70M+ songs
- Most ratings are missing!

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products
- Spotify: 400M+ users, 70M+ songs
- Most ratings are missing!

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products
- Spotify: 400M+ users, 70M+ songs
- Most ratings are missing!

Think About It:

- You've rated 100 movies out of 15,000

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products
- Spotify: 400M+ users, 70M+ songs
- Most ratings are missing!

Think About It:

- You've rated 100 movies out of 15,000
- Your friend has similar but different tastes

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products
- Spotify: 400M+ users, 70M+ songs
- Most ratings are missing!


Think About It:

- You've rated 100 movies out of 15,000
- Your friend has similar but different tastes
- How do we predict what

The Movie Recommendation Challenge

Real-World Scenario:

- Netflix: 200M+ users, 15K+ titles
- Amazon: 300M+ users, millions of products
- Spotify: 400M+ users, 70M+ songs
- Most ratings are missing!



Think About It:

- You've rated 100 movies out of 15,000
- Your friend has similar but different tastes
- How do we predict what

Pop Quiz 1: Understanding the Scale

Quick Question

If Netflix has 200 million users and 15,000 movies, how many possible ratings exist?

Pop Quiz 1: Understanding the Scale

Quick Question

If Netflix has 200 million users and 15,000 movies, how many possible ratings exist?

Answer: $200 \times 10^6 \times 15 \times 10^3 = 3 \times 10^{12}$ possible ratings!

Pop Quiz 1: Understanding the Scale

Quick Question

If Netflix has 200 million users and 15,000 movies, how many possible ratings exist?

Answer: $200 \times 10^6 \times 15 \times 10^3 = 3 \times 10^{12}$ possible ratings!

But typical users rate only 20-100 movies. What percentage of the matrix is filled?

Pop Quiz 1: Understanding the Scale

Quick Question

If Netflix has 200 million users and 15,000 movies, how many possible ratings exist?

Answer: $200 \times 10^6 \times 15 \times 10^3 = 3 \times 10^{12}$ possible ratings!

But typical users rate only 20-100 movies. What percentage of the matrix is filled?

Answer: $\frac{100}{15000} = 0.67\%$ - extremely sparse!

Mathematical Problem Setup

The Rating Matrix $A \in \mathbb{R}^{N \times M}$:

Mathematical Problem Setup

The Rating Matrix $\mathbf{A} \in \mathbb{R}^{N \times M}$:

$$\mathbf{A} = \begin{bmatrix} a_{11} & ? & a_{13} & ? & \cdots \\ ? & a_{22} & ? & a_{24} & \cdots \\ a_{31} & ? & ? & a_{34} & \cdots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

Mathematical Problem Setup

The Rating Matrix $\mathbf{A} \in \mathbb{R}^{N \times M}$:

$$\mathbf{A} = \begin{bmatrix} a_{11} & ? & a_{13} & ? & \dots \\ ? & a_{22} & ? & a_{24} & \dots \\ a_{31} & ? & ? & a_{34} & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

- **Rows:** Users u_1, u_2, \dots, u_N

Mathematical Problem Setup

The Rating Matrix $\mathbf{A} \in \mathbb{R}^{N \times M}$:

$$\mathbf{A} = \begin{bmatrix} a_{11} & ? & a_{13} & ? & \dots \\ ? & a_{22} & ? & a_{24} & \dots \\ a_{31} & ? & ? & a_{34} & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

- **Rows:** Users u_1, u_2, \dots, u_N
- **Columns:** Movies m_1, m_2, \dots, m_M

Mathematical Problem Setup

The Rating Matrix $\mathbf{A} \in \mathbb{R}^{N \times M}$:

$$\mathbf{A} = \begin{bmatrix} a_{11} & ? & a_{13} & ? & \dots \\ ? & a_{22} & ? & a_{24} & \dots \\ a_{31} & ? & ? & a_{34} & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

- **Rows:** Users u_1, u_2, \dots, u_N
- **Columns:** Movies m_1, m_2, \dots, m_M
- **Entries:** $a_{ij} \in \{1, 2, 3, 4, 5\}$ (when observed)

Mathematical Problem Setup

The Rating Matrix $\mathbf{A} \in \mathbb{R}^{N \times M}$:

$$\mathbf{A} = \begin{bmatrix} a_{11} & ? & a_{13} & ? & \dots \\ ? & a_{22} & ? & a_{24} & \dots \\ a_{31} & ? & ? & a_{34} & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

- **Rows:** Users u_1, u_2, \dots, u_N
- **Columns:** Movies m_1, m_2, \dots, m_M
- **Entries:** $a_{ij} \in \{1, 2, 3, 4, 5\}$ (when observed)
- **Challenge:** Predict missing entries ?

Mathematical Problem Setup

The Rating Matrix $\mathbf{A} \in \mathbb{R}^{N \times M}$:

$$\mathbf{A} = \begin{bmatrix} a_{11} & ? & a_{13} & ? & \dots \\ ? & a_{22} & ? & a_{24} & \dots \\ a_{31} & ? & ? & a_{34} & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

- **Rows:** Users u_1, u_2, \dots, u_N
- **Columns:** Movies m_1, m_2, \dots, m_M
- **Entries:** $a_{ij} \in \{1, 2, 3, 4, 5\}$ (when observed)
- **Challenge:** Predict missing entries ?
- **Notation:** $\Omega = \{(i, j) : a_{ij} \text{ is observed}\}$

Concrete Example: Our Movie Dataset

Let's work with a small, concrete example:

Concrete Example: Our Movie Dataset

Let's work with a small, concrete example:

User	Sholay	Swades	Batman	Interstellar	Shawshank
Alice	5	4	2	3	2
Bob	?	5	1	4	?
Carol	4	?	1	5	?

Concrete Example: Our Movie Dataset

Let's work with a small, concrete example:

User	Sholay	Swades	Batman	Interstellar	Shawshank
Alice	5	4	2	3	2
Bob	?	5	1	4	?
Carol	4	?	1	5	?

Observations:

- Alice loves Bollywood films (Sholay, Swades)

Concrete Example: Our Movie Dataset

Let's work with a small, concrete example:

User	Sholay	Swades	Batman	Interstellar	Shawshank
Alice	5	4	2	3	2
Bob	?	5	1	4	?
Carol	4	?	1	5	?

Observations:

- Alice loves Bollywood films (Sholay, Swades)
- Carol enjoys Sci-Fi (Interstellar)

Concrete Example: Our Movie Dataset

Let's work with a small, concrete example:

User	Sholay	Swades	Batman	Interstellar	Shawshank
Alice	5	4	2	3	2
Bob	?	5	1	4	?
Carol	4	?	1	5	?

Observations:

- Alice loves Bollywood films (Sholay, Swades)
- Carol enjoys Sci-Fi (Interstellar)
- Can we predict Bob's rating for Sholay?

Concrete Example: Our Movie Dataset

Let's work with a small, concrete example:

User	Sholay	Swades	Batman	Interstellar	Shawshank
Alice	5	4	2	3	2
Bob	?	5	1	4	?
Carol	4	?	1	5	?

Observations:

- Alice loves Bollywood films (Sholay, Swades)
- Carol enjoys Sci-Fi (Interstellar)
- Can we predict Bob's rating for Sholay?
- Can we predict Carol's rating for Swades?

Before We Dive In: A Simple Question

Why do you like the movies you like?

Before We Dive In: A Simple Question

Why do you like the movies you like?

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)
- Director (Christopher Nolan, Rajkumar Hirani)

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)
- Director (Christopher Nolan, Rajkumar Hirani)
- Language (Hindi, English, Tamil)

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)
- Director (Christopher Nolan, Rajkumar Hirani)
- Language (Hindi, English, Tamil)
- Era (90s classics, modern CGI)

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)
- Director (Christopher Nolan, Rajkumar Hirani)
- Language (Hindi, English, Tamil)
- Era (90s classics, modern CGI)

Key Insight:

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)
- Director (Christopher Nolan, Rajkumar Hirani)
- Language (Hindi, English, Tamil)
- Era (90s classics, modern CGI)

Key Insight:

- Your taste = combination of preferences

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)
- Director (Christopher Nolan, Rajkumar Hirani)
- Language (Hindi, English, Tamil)
- Era (90s classics, modern CGI)

Key Insight:

- Your taste = combination of preferences
- Movie appeal = combination of features

Before We Dive In: A Simple Question

Why do you like the movies you like?

Maybe because of:

- Genre (Action, Romance, Comedy)
- Star cast (Shah Rukh Khan, Tom Cruise)
- Director (Christopher Nolan, Rajkumar Hirani)
- Language (Hindi, English, Tamil)
- Era (90s classics, modern CGI)

Key Insight:

- Your taste = combination of preferences
- Movie appeal = combination of features
- **But we don't know these explicitly!**

The Core Insight: Hidden Patterns

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

Intuition: Think of latent features as "hidden DNA" of movies and users!

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

Intuition: Think of latent features as "hidden DNA" of movies and users!

For Movies:

- Bollywood vs Hollywood

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

Intuition: Think of latent features as "hidden DNA" of movies and users!

For Movies:

- Bollywood vs Hollywood
- Action vs Drama

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

Intuition: Think of latent features as "hidden DNA" of movies and users!

For Movies:

- Bollywood vs Hollywood
- Action vs Drama
- Comedy vs Serious

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

Intuition: Think of latent features as "hidden DNA" of movies and users!

For Movies:

- Bollywood vs Hollywood
- Action vs Drama
- Comedy vs Serious
- Runtime (Short vs Long)

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

Intuition: Think of latent features as "hidden DNA" of movies and users!


For Movies:

- Bollywood vs Hollywood
- Action vs Drama
- Comedy vs Serious
- Runtime (Short vs Long)
- Year (Classic vs Modern)

The Core Insight: Hidden Patterns

Hypothesis: User preferences and movie characteristics can be captured by a small number of **latent features**.

Intuition: Think of latent features as "hidden DNA" of movies and users!



Latent Features

For Movies:

- Bollywood vs Hollywood
- Action vs Drama
- Comedy vs Serious
- Runtime (Short vs Long)
- Year (Classic vs Modern)

Step 1: Define Movie Features Explicitly

Let's manually define features for our 5 movies:

Step 1: Define Movie Features Explicitly

Let's manually define features for our 5 movies:

Movie	Bollywood	Sci-Fi	Drama
Sholay	0.95	0.10	0.85
Swades	1.00	0.20	0.90
Batman	0.05	0.80	0.30
Interstellar	0.05	0.95	0.70
Shawshank	0.05	0.15	0.95

Step 1: Define Movie Features Explicitly

Let's manually define features for our 5 movies:

Movie	Bollywood	Sci-Fi	Drama
Sholay	0.95	0.10	0.85
Swades	1.00	0.20	0.90
Batman	0.05	0.80	0.30
Interstellar	0.05	0.95	0.70
Shawshank	0.05	0.15	0.95

Movie Feature Matrix $\mathbf{H} \in \mathbb{R}^{3 \times 5}$:

$$\mathbf{H} = \begin{bmatrix} 0.95 & 1.00 & 0.05 & 0.05 & 0.05 \\ 0.10 & 0.20 & 0.80 & 0.95 & 0.15 \\ 0.85 & 0.90 & 0.30 & 0.70 & 0.95 \end{bmatrix}$$

Step 2: What About User Preferences?

User Feature Matrix $\mathbf{W} \in \mathbb{R}^{3 \times 3}$ represents user preferences:

Step 2: What About User Preferences?

User Feature Matrix $\mathbf{W} \in \mathbb{R}^{3 \times 3}$ represents user preferences:

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix}$$

Step 2: What About User Preferences?

User Feature Matrix $\mathbf{W} \in \mathbb{R}^{3 \times 3}$ represents user preferences:

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix}$$

Where row i represents user i 's affinity for:

- w_{i1} : Bollywood preference

Step 2: What About User Preferences?

User Feature Matrix $\mathbf{W} \in \mathbb{R}^{3 \times 3}$ represents user preferences:

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix}$$

Where row i represents user i 's affinity for:

- w_{i1} : Bollywood preference
- w_{i2} : Sci-Fi preference

Step 2: What About User Preferences?

User Feature Matrix $\mathbf{W} \in \mathbb{R}^{3 \times 3}$ represents user preferences:

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix}$$

Where row i represents user i 's affinity for:

- w_{i1} : Bollywood preference
- w_{i2} : Sci-Fi preference
- w_{i3} : Drama preference

Step 2: What About User Preferences?

User Feature Matrix $\mathbf{W} \in \mathbb{R}^{3 \times 3}$ represents user preferences:

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix}$$

Where row i represents user i 's affinity for:

- w_{i1} : Bollywood preference
- w_{i2} : Sci-Fi preference
- w_{i3} : Drama preference

Step 2: What About User Preferences?

User Feature Matrix $\mathbf{W} \in \mathbb{R}^{3 \times 3}$ represents user preferences:

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix}$$

Where row i represents user i 's affinity for:

- w_{i1} : Bollywood preference
- w_{i2} : Sci-Fi preference
- w_{i3} : Drama preference

Key Question: How do we learn these w_{ij} values from observed ratings?

Step 3: The Matrix Factorization Idea

Core Hypothesis: Rating = User preferences · Movie features

Step 3: The Matrix Factorization Idea

Core Hypothesis: Rating = User preferences · Movie features

$$a_{ij} \approx \mathbf{w}_i^T \mathbf{h}_j = \sum_{k=1}^r w_{ik} h_{kj}$$

Step 3: The Matrix Factorization Idea

Core Hypothesis: Rating = User preferences · Movie features

$$a_{ij} \approx \mathbf{w}_i^T \mathbf{h}_j = \sum_{k=1}^r w_{ik} h_{kj}$$

In Matrix Form:

$$\mathbf{A} \approx \mathbf{WH}$$

Step 3: The Matrix Factorization Idea

Core Hypothesis: Rating = User preferences · Movie features

$$a_{ij} \approx \mathbf{w}_i^T \mathbf{h}_j = \sum_{k=1}^r w_{ik} h_{kj}$$

In Matrix Form:

$$\mathbf{A} \approx \mathbf{WH}$$

$$\mathbf{A}_{3 \times 5} = \begin{bmatrix} 5 & 4 & 2 & 3 & 2 \\ ? & 5 & 1 & 4 & ? \\ 4 & ? & 1 & 5 & ? \end{bmatrix} \approx$$

$$\begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix} \begin{bmatrix} 0.95 & 1.00 & 0.05 & 0.05 & 0.05 \\ 0.10 & 0.20 & 0.80 & 0.95 & 0.15 \\ 0.85 & 0.90 & 0.30 & 0.70 & 0.95 \end{bmatrix} =$$
$$\mathbf{W}_{3 \times 3} \mathbf{H}_{3 \times 5}$$

Step 4: Understanding the Calculation

Let's think step by step...

Step 4: Understanding the Calculation

Let's think step by step...

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}
- How much does she like Comedy? w_{13}

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}
- How much does she like Comedy? w_{13}

Sholay's DNA:

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}
- How much does she like Comedy? w_{13}

Sholay's DNA:

- Bollywood-ness: 0.95
(very high!)

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}
- How much does she like Comedy? w_{13}

Sholay's DNA:

- Bollywood-ness: 0.95 (very high!)
- Action-ness: 0.10 (low)

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}
- How much does she like Comedy? w_{13}

Sholay's DNA:

- Bollywood-ness: 0.95 (very high!)
- Action-ness: 0.10 (low)
- Comedy-ness: 0.85 (high)

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}
- How much does she like Comedy? w_{13}

Sholay's DNA:

- Bollywood-ness: 0.95 (very high!)
- Action-ness: 0.10 (low)
- Comedy-ness: 0.85 (high)

Step 4: Understanding the Calculation

Let's think step by step...

Alice's Profile:

- How much does she like Bollywood? w_{11}
- How much does she like Action? w_{12}
- How much does she like Comedy? w_{13}

Sholay's DNA:

- Bollywood-ness: 0.95 (very high!)
- Action-ness: 0.10 (low)
- Comedy-ness: 0.85 (high)

The Magic Formula:

Alice's rating = Alice's preferences · Sholay's features

Step 4: Detailed Calculation Example

Let's compute Alice's predicted rating for Sholay:

Step 4: Detailed Calculation Example

Let's compute Alice's predicted rating for Sholay:

Alice's preferences: $\mathbf{w}_1 = [w_{11}, w_{12}, w_{13}]$ **Sholay's features:** $\mathbf{h}_1 = [0.95, 0.10, 0.85]^T$

Step 4: Detailed Calculation Example

Let's compute Alice's predicted rating for Sholay:

Alice's preferences: $\mathbf{w}_1 = [w_{11}, w_{12}, w_{13}]$ **Sholay's features:** $\mathbf{h}_1 = [0.95, 0.10, 0.85]^T$

$$\hat{a}_{11} = \mathbf{w}_1^T \mathbf{h}_1 \quad (1)$$

$$= w_{11} \cdot 0.95 + w_{12} \cdot 0.10 + w_{13} \cdot 0.85 \quad (2)$$

Step 4: Detailed Calculation Example

Let's compute Alice's predicted rating for Sholay:

Alice's preferences: $\mathbf{w}_1 = [w_{11}, w_{12}, w_{13}]$ **Sholay's features:** $\mathbf{h}_1 = [0.95, 0.10, 0.85]^T$

$$\hat{a}_{11} = \mathbf{w}_1^T \mathbf{h}_1 \quad (1)$$

$$= w_{11} \cdot 0.95 + w_{12} \cdot 0.10 + w_{13} \cdot 0.85 \quad (2)$$

Goal: Find w_{11}, w_{12}, w_{13} such that $\hat{a}_{11} \approx 5$ (Alice's actual rating)

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of A ?

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Answers:

1. $\mathbf{A} \in \mathbb{R}^{N \times M}$

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Answers:

1. $\mathbf{A} \in \mathbb{R}^{N \times M}$
2. $\mathbf{W} \in \mathbb{R}^{N \times r}$

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Answers:

1. $\mathbf{A} \in \mathbb{R}^{N \times M}$
2. $\mathbf{W} \in \mathbb{R}^{N \times r}$
3. $\mathbf{H} \in \mathbb{R}^{r \times M}$

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Answers:

1. $\mathbf{A} \in \mathbb{R}^{N \times M}$
2. $\mathbf{W} \in \mathbb{R}^{N \times r}$
3. $\mathbf{H} \in \mathbb{R}^{r \times M}$
4. Total parameters: $Nr + rM = r(N + M)$

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Answers:

1. $\mathbf{A} \in \mathbb{R}^{N \times M}$
2. $\mathbf{W} \in \mathbb{R}^{N \times r}$
3. $\mathbf{H} \in \mathbb{R}^{r \times M}$
4. Total parameters: $Nr + rM = r(N + M)$

Pop Quiz 2: Matrix Dimensions

Dimension Check

If we have N users, M movies, and r latent features:

1. What are the dimensions of \mathbf{A} ?
2. What are the dimensions of \mathbf{W} ?
3. What are the dimensions of \mathbf{H} ?
4. How many parameters do we need to learn?

Answers:

1. $\mathbf{A} \in \mathbb{R}^{N \times M}$
2. $\mathbf{W} \in \mathbb{R}^{N \times r}$
3. $\mathbf{H} \in \mathbb{R}^{r \times M}$
4. Total parameters: $Nr + rM = r(N + M)$

Key Insight: If $r \ll \min(N, M)$, we have huge param-

The Optimization Problem

Objective: Minimize prediction error on observed ratings only

The Optimization Problem

Objective: Minimize prediction error on observed ratings only

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

The Optimization Problem

Objective: Minimize prediction error on observed ratings only

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

In Matrix Notation:

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \|P_{\Omega}(\mathbf{A} - \mathbf{WH})\|_F^2$$

The Optimization Problem

Objective: Minimize prediction error on observed ratings only

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

In Matrix Notation:

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \|P_{\Omega}(\mathbf{A} - \mathbf{WH})\|_F^2$$

Where:

- $P_{\Omega}(\cdot)$: projection onto observed entries

The Optimization Problem

Objective: Minimize prediction error on observed ratings only

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

In Matrix Notation:

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \|P_{\Omega}(\mathbf{A} - \mathbf{WH})\|_F^2$$

Where:

- $P_{\Omega}(\cdot)$: projection onto observed entries
- $\|\cdot\|_F$: Frobenius norm

The Optimization Problem

Objective: Minimize prediction error on observed ratings only

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

In Matrix Notation:

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \|P_{\Omega}(\mathbf{A} - \mathbf{WH})\|_F^2$$

Where:

- $P_{\Omega}(\cdot)$: projection onto observed entries
- $\|\cdot\|_F$: Frobenius norm
- Ω : set of observed (i,j) pairs

Why This is Challenging

Problem Characteristics:

- **Non-convex:** Multiple local minima exist

Why This is Challenging

Problem Characteristics:

- **Non-convex:** Multiple local minima exist
- **Bilinear:** Linear in \mathbf{W} when \mathbf{H} fixed, and vice versa

Why This is Challenging

Problem Characteristics:

- **Non-convex:** Multiple local minima exist
- **Bilinear:** Linear in \mathbf{W} when \mathbf{H} fixed, and vice versa
- **Large-scale:** Millions of users and items

Why This is Challenging

Problem Characteristics:

- **Non-convex:** Multiple local minima exist
- **Bilinear:** Linear in \mathbf{W} when \mathbf{H} fixed, and vice versa
- **Large-scale:** Millions of users and items
- **Sparse:** Only 0.1-1% of entries observed

Why This is Challenging

Problem Characteristics:

- **Non-convex:** Multiple local minima exist
- **Bilinear:** Linear in \mathbf{W} when \mathbf{H} fixed, and vice versa
- **Large-scale:** Millions of users and items
- **Sparse:** Only 0.1-1% of entries observed

Why This is Challenging

Problem Characteristics:

- **Non-convex:** Multiple local minima exist
- **Bilinear:** Linear in \mathbf{W} when \mathbf{H} fixed, and vice versa
- **Large-scale:** Millions of users and items
- **Sparse:** Only 0.1-1% of entries observed

Key Insight: While non-convex jointly, it's convex in each matrix individually!

ALS: The Big Picture

Alternating Least Squares Strategy:

ALS: The Big Picture

Alternating Least Squares Strategy:

1. **Initialize:** $\mathbf{W}^{(0)}$ and $\mathbf{H}^{(0)}$ randomly

ALS: The Big Picture

Alternating Least Squares Strategy:

1. **Initialize:** $\mathbf{W}^{(0)}$ and $\mathbf{H}^{(0)}$ randomly
2. **Repeat until convergence:**

ALS: The Big Picture

Alternating Least Squares Strategy:

1. **Initialize:** $W^{(0)}$ and $H^{(0)}$ randomly
2. **Repeat until convergence:**
 - 1) **Fix H , solve for W :** Each row independently

ALS: The Big Picture

Alternating Least Squares Strategy:

1. **Initialize:** $W^{(0)}$ and $H^{(0)}$ randomly
2. **Repeat until convergence:**
 - 1) **Fix H , solve for W :** Each row independently
 - 2) **Fix W , solve for H :** Each column independently

ALS: The Big Picture

Alternating Least Squares Strategy:

1. **Initialize:** $W^{(0)}$ and $H^{(0)}$ randomly
2. **Repeat until convergence:**
 - 1) **Fix H , solve for W :** Each row independently
 - 2) **Fix W , solve for H :** Each column independently
3. Each subproblem is a standard least squares problem!

ALS: The Big Picture


Alternating Least Squares Strategy:

1. **Initialize:** $W^{(0)}$ and $H^{(0)}$ randomly
2. **Repeat until convergence:**
 - 1) **Fix H , solve for W :** Each row independently
 - 2) **Fix W , solve for H :** Each column independently
3. Each subproblem is a standard least squares problem!

ALS: The Big Picture

Alternating Least Squares Strategy:

1. **Initialize:** $W^{(0)}$ and $H^{(0)}$ randomly
2. **Repeat until convergence:**
 - 1) **Fix H , solve for W :** Each row independently
 - 2) **Fix W , solve for H :** Each column independently
3. Each subproblem is a standard least squares problem!



ALS Step 1: Updating User Features

Fix H , solve for each user i independently:

ALS Step 1: Updating User Features

Fix H , solve for each user i independently:

$$\text{minimize}_{\mathbf{w}_i} \sum_{j:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

ALS Step 1: Updating User Features

Fix H , solve for each user i independently:

$$\text{minimize}_{\mathbf{w}_i} \sum_{j:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Matrix Form for User i : Let $\Omega_i = \{j : (i,j) \in \Omega\}$ (movies rated by user i)

ALS Step 1: Updating User Features

Fix \mathbf{H} , solve for each user i independently:

$$\text{minimize}_{\mathbf{w}_i} \sum_{j:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Matrix Form for User i : Let $\Omega_i = \{j : (i,j) \in \Omega\}$ (movies rated by user i)

$$\mathbf{y}_i = [a_{i,j_1}, a_{i,j_2}, \dots, a_{i,j_{|\Omega_i|}}]^T \quad (3)$$

$$\mathbf{X}_i = [\mathbf{h}_{j_1}, \mathbf{h}_{j_2}, \dots, \mathbf{h}_{j_{|\Omega_i|}}]^T \quad (4)$$

ALS Step 1: Updating User Features

Fix \mathbf{H} , solve for each user i independently:

$$\text{minimize}_{\mathbf{w}_i} \sum_{j:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Matrix Form for User i : Let $\Omega_i = \{j : (i, j) \in \Omega\}$ (movies rated by user i)

$$\mathbf{y}_i = [a_{i,j_1}, a_{i,j_2}, \dots, a_{i,j_{|\Omega_i|}}]^T \quad (3)$$

$$\mathbf{X}_i = [\mathbf{h}_{j_1}, \mathbf{h}_{j_2}, \dots, \mathbf{h}_{j_{|\Omega_i|}}]^T \quad (4)$$

Least Squares Solution:

$$\boxed{\mathbf{w}_i^* = (\mathbf{X}_i^T \mathbf{X}_i)^{-1} \mathbf{X}_i^T \mathbf{y}_i}$$

ALS Step 1: Concrete Example

Update Alice's preferences (w_1):

Alice rated: Sholay(5), Swades(4), Batman(2),
Interstellar(3), Shawshank(2)

ALS Step 1: Concrete Example

Update Alice's preferences (w_1):

Alice rated: Sholay(5), Swades(4), Batman(2),
Interstellar(3), Shawshank(2)

$$y_1 = \begin{bmatrix} 5 \\ 4 \\ 2 \\ 3 \\ 2 \end{bmatrix} \quad (5)$$

$$X_1 = \begin{bmatrix} 0.95 & 0.10 & 0.85 \\ 1.00 & 0.20 & 0.90 \\ 0.05 & 0.80 & 0.30 \\ 0.05 & 0.95 & 0.70 \\ 0.05 & 0.15 & 0.95 \end{bmatrix} \quad (6)$$

ALS Step 1: Concrete Example

Update Alice's preferences (w_1):

Alice rated: Sholay(5), Swades(4), Batman(2),
Interstellar(3), Shawshank(2)

$$y_1 = \begin{bmatrix} 5 \\ 4 \\ 2 \\ 3 \\ 2 \end{bmatrix} \quad (5)$$

$$X_1 = \begin{bmatrix} 0.95 & 0.10 & 0.85 \\ 1.00 & 0.20 & 0.90 \\ 0.05 & 0.80 & 0.30 \\ 0.05 & 0.95 & 0.70 \\ 0.05 & 0.15 & 0.95 \end{bmatrix} \quad (6)$$

Solution: $w_1^* = (X_1^T X_1)^{-1} X_1^T y_1$

This gives us Alice's feature preferences!

ALS Step 2: Updating Movie Features

Fix W , solve for each movie j independently:

ALS Step 2: Updating Movie Features

Fix W , solve for each movie j independently:

$$\text{minimize}_{\mathbf{h}_j} \sum_{i:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

ALS Step 2: Updating Movie Features

Fix W , solve for each movie j independently:

$$\text{minimize}_{\mathbf{h}_j} \sum_{i:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Matrix Form for Movie j : Let $\Omega_j = \{i : (i,j) \in \Omega\}$ (users who rated movie j)

ALS Step 2: Updating Movie Features

Fix W , solve for each movie j independently:

$$\text{minimize}_{\mathbf{h}_j} \sum_{i:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Matrix Form for Movie j : Let $\Omega_j = \{i : (i,j) \in \Omega\}$ (users who rated movie j)

$$\mathbf{y}_j = [a_{i_1,j}, a_{i_2,j}, \dots, a_{i_{|\Omega_j|},j}]^T \quad (7)$$

$$\mathbf{X}_j = [\mathbf{w}_{i_1}, \mathbf{w}_{i_2}, \dots, \mathbf{w}_{i_{|\Omega_j|}}]^T \quad (8)$$

ALS Step 2: Updating Movie Features

Fix W , solve for each movie j independently:

$$\text{minimize}_{\mathbf{h}_j} \sum_{i:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Matrix Form for Movie j : Let $\Omega_j = \{i : (i,j) \in \Omega\}$ (users who rated movie j)

$$\mathbf{y}_j = [a_{i_1,j}, a_{i_2,j}, \dots, a_{i_{|\Omega_j|},j}]^T \quad (7)$$

$$\mathbf{X}_j = [\mathbf{w}_{i_1}, \mathbf{w}_{i_2}, \dots, \mathbf{w}_{i_{|\Omega_j|}}]^T \quad (8)$$

Least Squares Solution:

$$\mathbf{h}_j^* = (\mathbf{X}_j^T \mathbf{X}_j)^{-1} \mathbf{X}_j^T \mathbf{y}_j$$

ALS: Complete Algorithm

Algorithm 1: [

H] **Input:** Rating matrix \mathbf{A} , rank r , max iterations T

1. **Initialize:** $\mathbf{W}^{(0)} \in \mathbb{R}^{N \times r}$, $\mathbf{H}^{(0)} \in \mathbb{R}^{r \times M}$ randomly

Output: $\mathbf{W}^{(T)}$ $\mathbf{H}^{(T)}$

ALS: Complete Algorithm

Algorithm 2: [

H] **Input:** Rating matrix \mathbf{A} , rank r , max iterations T

1. **Initialize:** $\mathbf{W}^{(0)} \in \mathbb{R}^{N \times r}$, $\mathbf{H}^{(0)} \in \mathbb{R}^{r \times M}$ randomly
2. **For** $t = 1, 2, \dots, T$:

Output: $\mathbf{W}^{(T)}$ $\mathbf{H}^{(T)}$

ALS: Complete Algorithm

Algorithm 3: [

H] **Input:** Rating matrix \mathbf{A} , rank r , max iterations T

1. **Initialize:** $\mathbf{W}^{(0)} \in \mathbb{R}^{N \times r}$, $\mathbf{H}^{(0)} \in \mathbb{R}^{r \times M}$ randomly
2. **For** $t = 1, 2, \dots, T$:
 - 1) **Update Users:** For each user $i = 1, \dots, N$:

$$\mathbf{w}_i^{(t)} = (\mathbf{X}_i^T \mathbf{X}_i)^{-1} \mathbf{X}_i^T \mathbf{y}_i$$

Output: $\mathbf{W}^{(T)}$ $\mathbf{H}^{(T)}$

ALS: Complete Algorithm

Algorithm 4: [

H] **Input:** Rating matrix \mathbf{A} , rank r , max iterations T

1. **Initialize:** $\mathbf{W}^{(0)} \in \mathbb{R}^{N \times r}$, $\mathbf{H}^{(0)} \in \mathbb{R}^{r \times M}$ randomly
2. **For** $t = 1, 2, \dots, T$:

1) **Update Users:** For each user $i = 1, \dots, N$:

$$\mathbf{w}_i^{(t)} = (\mathbf{X}_i^T \mathbf{X}_i)^{-1} \mathbf{X}_i^T \mathbf{y}_i$$

2) **Update Movies:** For each movie $j = 1, \dots, M$:

$$\mathbf{h}_j^{(t)} = (\mathbf{X}_j^T \mathbf{X}_j)^{-1} \mathbf{X}_j^T \mathbf{y}_j$$

Output: $\mathbf{W}^{(T)}$ $\mathbf{H}^{(T)}$

ALS: Complete Algorithm

Algorithm 5: [

H] **Input:** Rating matrix \mathbf{A} , rank r , max iterations T

1. **Initialize:** $\mathbf{W}^{(0)} \in \mathbb{R}^{N \times r}$, $\mathbf{H}^{(0)} \in \mathbb{R}^{r \times M}$ randomly
2. **For** $t = 1, 2, \dots, T$:

1) **Update Users:** For each user $i = 1, \dots, N$:

$$\mathbf{w}_i^{(t)} = (\mathbf{X}_i^T \mathbf{X}_i)^{-1} \mathbf{X}_i^T \mathbf{y}_i$$

2) **Update Movies:** For each movie $j = 1, \dots, M$:

$$\mathbf{h}_j^{(t)} = (\mathbf{X}_j^T \mathbf{X}_j)^{-1} \mathbf{X}_j^T \mathbf{y}_j$$

3. **Check Convergence:** Stop if
 $\|\mathbf{W}^{(t)} \mathbf{H}^{(t)} - \mathbf{W}^{(t-1)} \mathbf{H}^{(t-1)}\|_F < \epsilon$

Output: $\mathbf{W}^{(T)}$ $\mathbf{H}^{(T)}$

Gradient Descent Approach

Simultaneous Updates: Update both \mathbf{W} and \mathbf{H} together

Gradient Descent Approach

Simultaneous Updates: Update both \mathbf{W} and \mathbf{H} together
Objective Function:

$$L(\mathbf{W}, \mathbf{H}) = \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Gradient Descent Approach

Simultaneous Updates: Update both \mathbf{W} and \mathbf{H} together
Objective Function:

$$L(\mathbf{W}, \mathbf{H}) = \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2$$

Gradients:

$$\frac{\partial L}{\partial \mathbf{w}_i} = -2 \sum_{j:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j) \mathbf{h}_j \quad (9)$$

$$\frac{\partial L}{\partial \mathbf{h}_j} = -2 \sum_{i:(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j) \mathbf{w}_i \quad (10)$$

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

1. Make a guess about their rating

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

1. Make a guess about their rating
2. See their actual rating

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

1. Make a guess about their rating
2. See their actual rating
3. Adjust your understanding

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

1. Make a guess about their rating
2. See their actual rating
3. Adjust your understanding
4. Repeat for next movie

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

SGD does exactly this!

1. Make a guess about their rating
2. See their actual rating
3. Adjust your understanding
4. Repeat for next movie

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

1. Make a guess about their rating
2. See their actual rating
3. Adjust your understanding
4. Repeat for next movie

SGD does exactly this!

- One rating at a time

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

1. Make a guess about their rating
2. See their actual rating
3. Adjust your understanding
4. Repeat for next movie

SGD does exactly this!

- One rating at a time
- Small adjustments

SGD: Learning Like a Human

Imagine you're learning someone's taste in movies...

Your Process:

1. Make a guess about their rating
2. See their actual rating
3. Adjust your understanding
4. Repeat for next movie

SGD does exactly this!

- One rating at a time
- Small adjustments
- Gradually improves

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

1. **Predict:** $\hat{a}_{ij} = \mathbf{w}_i^T \mathbf{h}_j$

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

1. **Predict:** $\hat{a}_{ij} = \mathbf{w}_i^T \mathbf{h}_j$
2. **Compute Error:** $e_{ij} = a_{ij} - \hat{a}_{ij}$

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

1. **Predict:** $\hat{a}_{ij} = \mathbf{w}_i^T \mathbf{h}_j$
2. **Compute Error:** $e_{ij} = a_{ij} - \hat{a}_{ij}$
3. **Update:**

$$\mathbf{w}_i \leftarrow \mathbf{w}_i + \alpha \cdot e_{ij} \cdot \mathbf{h}_j \quad (11)$$

$$\mathbf{h}_j \leftarrow \mathbf{h}_j + \alpha \cdot e_{ij} \cdot \mathbf{w}_i \quad (12)$$

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

1. **Predict:** $\hat{a}_{ij} = \mathbf{w}_i^T \mathbf{h}_j$
2. **Compute Error:** $e_{ij} = a_{ij} - \hat{a}_{ij}$
3. **Update:**

$$\mathbf{w}_i \leftarrow \mathbf{w}_i + \alpha \cdot e_{ij} \cdot \mathbf{h}_j \quad (11)$$

$$\mathbf{h}_j \leftarrow \mathbf{h}_j + \alpha \cdot e_{ij} \cdot \mathbf{w}_i \quad (12)$$

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

1. **Predict:** $\hat{a}_{ij} = \mathbf{w}_i^T \mathbf{h}_j$
2. **Compute Error:** $e_{ij} = a_{ij} - \hat{a}_{ij}$
3. **Update:**

$$\mathbf{w}_i \leftarrow \mathbf{w}_i + \alpha \cdot e_{ij} \cdot \mathbf{h}_j \quad (11)$$

$$\mathbf{h}_j \leftarrow \mathbf{h}_j + \alpha \cdot e_{ij} \cdot \mathbf{w}_i \quad (12)$$

Intuition:

- If $e_{ij} > 0$: Predicted rating too low \rightarrow Increase similarity

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

1. **Predict:** $\hat{a}_{ij} = \mathbf{w}_i^T \mathbf{h}_j$
2. **Compute Error:** $e_{ij} = a_{ij} - \hat{a}_{ij}$
3. **Update:**

$$\mathbf{w}_i \leftarrow \mathbf{w}_i + \alpha \cdot e_{ij} \cdot \mathbf{h}_j \quad (11)$$

$$\mathbf{h}_j \leftarrow \mathbf{h}_j + \alpha \cdot e_{ij} \cdot \mathbf{w}_i \quad (12)$$

Intuition:

- If $e_{ij} > 0$: Predicted rating too low \rightarrow Increase similarity
- If $e_{ij} < 0$: Predicted rating too high \rightarrow Decrease similarity

Stochastic Gradient Descent (SGD)

For each observed rating $(i, j) \in \Omega$:

1. **Predict:** $\hat{a}_{ij} = \mathbf{w}_i^T \mathbf{h}_j$
2. **Compute Error:** $e_{ij} = a_{ij} - \hat{a}_{ij}$
3. **Update:**

$$\mathbf{w}_i \leftarrow \mathbf{w}_i + \alpha \cdot e_{ij} \cdot \mathbf{h}_j \quad (11)$$

$$\mathbf{h}_j \leftarrow \mathbf{h}_j + \alpha \cdot e_{ij} \cdot \mathbf{w}_i \quad (12)$$

Intuition:

- If $e_{ij} > 0$: Predicted rating too low \rightarrow Increase similarity
- If $e_{ij} < 0$: Predicted rating too high \rightarrow Decrease similarity
- Learning rate α controls step size

SGD: Step-by-Step Example

Example: Alice rates Sholay as 5, but we predict 3.2

SGD: Step-by-Step Example

Example: Alice rates Sholay as 5, but we predict 3.2

$$\text{Current: } \mathbf{w}_1 = [0.4, 0.2, 0.3], \quad \mathbf{h}_1 = [0.95, 0.10, 0.85] \quad (13)$$

$$\text{Prediction: } \hat{a}_{11} = 0.4 \times 0.95 + 0.2 \times 0.10 + 0.3 \times 0.85 = 0.655 \quad (14)$$

$$\text{Error: } e_{11} = 5 - 0.655 = 4.345 \quad (15)$$

SGD: Step-by-Step Example

Example: Alice rates Sholay as 5, but we predict 3.2

$$\text{Current: } \mathbf{w}_1 = [0.4, 0.2, 0.3], \quad \mathbf{h}_1 = [0.95, 0.10, 0.85] \quad (13)$$

$$\text{Prediction: } \hat{a}_{11} = 0.4 \times 0.95 + 0.2 \times 0.10 + 0.3 \times 0.85 = 0.655 \quad (14)$$

$$\text{Error: } e_{11} = 5 - 0.655 = 4.345 \quad (15)$$

Updates with $\alpha = 0.01$:

$$\mathbf{w}_1 \leftarrow [0.4, 0.2, 0.3] + 0.01 \times 4.345 \times [0.95, 0.10, 0.85] \quad (16)$$

$$= [0.4413, 0.2043, 0.3369] \quad (17)$$

$$\mathbf{h}_1 \leftarrow [0.95, 0.10, 0.85] + 0.01 \times 4.345 \times [0.4, 0.2, 0.3] \quad (18)$$

$$= [0.9674, 0.1087, 0.8631] \quad (19)$$

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?
2. Should we increase or decrease the user-movie similarity?

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?
2. Should we increase or decrease the user-movie similarity?
3. If $\alpha = 0.1$, $\mathbf{w}_i = [0.8, 0.3]$, $\mathbf{h}_j = [0.6, 0.9]$, what are the updates?

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?
2. Should we increase or decrease the user-movie similarity?
3. If $\alpha = 0.1$, $\mathbf{w}_i = [0.8, 0.3]$, $\mathbf{h}_j = [0.6, 0.9]$, what are the updates?

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?
2. Should we increase or decrease the user-movie similarity?
3. If $\alpha = 0.1$, $\mathbf{w}_i = [0.8, 0.3]$, $\mathbf{h}_j = [0.6, 0.9]$, what are the updates?

Answers:

1. $e_{ij} = 2 - 4.5 = -2.5$

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?
2. Should we increase or decrease the user-movie similarity?
3. If $\alpha = 0.1$, $\mathbf{w}_i = [0.8, 0.3]$, $\mathbf{h}_j = [0.6, 0.9]$, what are the updates?

Answers:

1. $e_{ij} = 2 - 4.5 = -2.5$
2. Decrease similarity (negative error)

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?
2. Should we increase or decrease the user-movie similarity?
3. If $\alpha = 0.1$, $\mathbf{w}_i = [0.8, 0.3]$, $\mathbf{h}_j = [0.6, 0.9]$, what are the updates?

Answers:

1. $e_{ij} = 2 - 4.5 = -2.5$
2. Decrease similarity (negative error)
3. $\mathbf{w}_i \leftarrow [0.8, 0.3] + 0.1 \times (-2.5) \times [0.6, 0.9] = [0.65, 0.075]$

Pop Quiz 3: SGD Understanding

Quick Check

A user gives a rating of 2 to a movie, but our model predicts 4.5.

1. What is the error e_{ij} ?
2. Should we increase or decrease the user-movie similarity?
3. If $\alpha = 0.1$, $\mathbf{w}_i = [0.8, 0.3]$, $\mathbf{h}_j = [0.6, 0.9]$, what are the updates?

Answers:

1. $e_{ij} = 2 - 4.5 = -2.5$
2. Decrease similarity (negative error)
3. $\mathbf{w}_i \leftarrow [0.8, 0.3] + 0.1 \times (-2.5) \times [0.6, 0.9] = [0.65, 0.075]$
4. $\mathbf{h}_j \leftarrow [0.6, 0.9] + 0.1 \times (-2.5) \times [0.8, 0.3] = [0.4, 0.825]$

ALS vs SGD: Head-to-Head Comparison

Aspect	ALS	SGD
Updates	Alternating	Simultaneous
Convergence	Faster, more stable	Slower, can oscillate
Parallelization	Excellent	Limited
Memory	Higher	Lower
Implementation	Complex	Simple
Hyperparameters	Few (rank r)	Many (α , schedule)
Scalability	Very good	Good

ALS vs SGD: Head-to-Head Comparison

Aspect	ALS	SGD
Updates	Alternating	Simultaneous
Convergence	Faster, more stable	Slower, can oscillate
Parallelization	Excellent	Limited
Memory	Higher	Lower
Implementation	Complex	Simple
Hyperparameters	Few (rank r)	Many (α , schedule)
Scalability	Very good	Good

When to Use Which?

- **ALS:** Large-scale, production systems (Spark, distributed)

ALS vs SGD: Head-to-Head Comparison

Aspect	ALS	SGD
Updates	Alternating	Simultaneous
Convergence	Faster, more stable	Slower, can oscillate
Parallelization	Excellent	Limited
Memory	Higher	Lower
Implementation	Complex	Simple
Hyperparameters	Few (rank r)	Many (α , schedule)
Scalability	Very good	Good

When to Use Which?

- **ALS:** Large-scale, production systems (Spark, distributed)
- **SGD:** Online learning, real-time updates, research

Advanced Practical Considerations

Regularization: Prevent overfitting

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \lambda (\|\mathbf{W}\|_F^2 + \|\mathbf{H}\|_F^2)$$

Advanced Practical Considerations

Regularization: Prevent overfitting

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \lambda (\|\mathbf{W}\|_F^2 + \|\mathbf{H}\|_F^2)$$

Bias Terms: Account for global, user, and item biases

$$\hat{a}_{ij} = \mu + b_i + b_j + \mathbf{w}_i^T \mathbf{h}_j$$

Advanced Practical Considerations

Regularization: Prevent overfitting

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \lambda (\|\mathbf{W}\|_F^2 + \|\mathbf{H}\|_F^2)$$

Bias Terms: Account for global, user, and item biases

$$\hat{a}_{ij} = \mu + b_i + b_j + \mathbf{w}_i^T \mathbf{h}_j$$

Implicit Feedback: Binary observations (clicks, views)

$$\text{Confidence: } c_{ij} = 1 + \alpha \cdot \text{frequency}_{ij}$$

Advanced Practical Considerations

Regularization: Prevent overfitting

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \lambda (\|\mathbf{W}\|_F^2 + \|\mathbf{H}\|_F^2)$$

Bias Terms: Account for global, user, and item biases

$$\hat{a}_{ij} = \mu + b_i + b_j + \mathbf{w}_i^T \mathbf{h}_j$$

Implicit Feedback: Binary observations (clicks, views)

$$\text{Confidence: } c_{ij} = 1 + \alpha \cdot \text{frequency}_{ij}$$

Cold Start Problem: New users/items with no ratings

- Content-based features

Advanced Practical Considerations

Regularization: Prevent overfitting

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \lambda (\|\mathbf{W}\|_F^2 + \|\mathbf{H}\|_F^2)$$

Bias Terms: Account for global, user, and item biases

$$\hat{a}_{ij} = \mu + b_i + b_j + \mathbf{w}_i^T \mathbf{h}_j$$

Implicit Feedback: Binary observations (clicks, views)

$$\text{Confidence: } c_{ij} = 1 + \alpha \cdot \text{frequency}_{ij}$$

Cold Start Problem: New users/items with no ratings

- Content-based features
- Demographic information

Advanced Practical Considerations

Regularization: Prevent overfitting

$$\text{minimize}_{\mathbf{W}, \mathbf{H}} \sum_{(i,j) \in \Omega} (a_{ij} - \mathbf{w}_i^T \mathbf{h}_j)^2 + \lambda (\|\mathbf{W}\|_F^2 + \|\mathbf{H}\|_F^2)$$

Bias Terms: Account for global, user, and item biases

$$\hat{a}_{ij} = \mu + b_i + b_j + \mathbf{w}_i^T \mathbf{h}_j$$

Implicit Feedback: Binary observations (clicks, views)

$$\text{Confidence: } c_{ij} = 1 + \alpha \cdot \text{frequency}_{ij}$$

Cold Start Problem: New users/items with no ratings

- Content-based features
- Demographic information
- Hybrid approaches

Let's Build Intuition: Small Example

Our 3×3 rating matrix:

$$\mathbf{A} = \begin{bmatrix} 5 & ? & 2 \\ 4 & 4 & ? \\ ? & 5 & 1 \end{bmatrix}$$

Let's Build Intuition: Small Example

Our 3×3 rating matrix:

$$\mathbf{A} = \begin{bmatrix} 5 & ? & 2 \\ 4 & 4 & ? \\ ? & 5 & 1 \end{bmatrix}$$

Goal: Find $\mathbf{W} \in \mathbb{R}^{3 \times 2}$ and $\mathbf{H} \in \mathbb{R}^{2 \times 3}$ such that:

$$\mathbf{A} \approx \mathbf{WH} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \\ w_{31} & w_{32} \end{bmatrix} \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \end{bmatrix}$$

Let's Build Intuition: Small Example

Our 3×3 rating matrix:

$$\mathbf{A} = \begin{bmatrix} 5 & ? & 2 \\ 4 & 4 & ? \\ ? & 5 & 1 \end{bmatrix}$$

Goal: Find $\mathbf{W} \in \mathbb{R}^{3 \times 2}$ and $\mathbf{H} \in \mathbb{R}^{2 \times 3}$ such that:

$$\mathbf{A} \approx \mathbf{WH} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \\ w_{31} & w_{32} \end{bmatrix} \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \end{bmatrix}$$

Constraint: Only minimize error on observed entries!

Step-by-Step ALS Solution

Iteration 1: Initialize randomly

$$\mathbf{W}^{(0)} = \begin{bmatrix} 0.5 & 0.3 \\ 0.4 & 0.6 \\ 0.2 & 0.8 \end{bmatrix}, \quad \mathbf{H}^{(0)} = \begin{bmatrix} 1.0 & 0.5 & 0.2 \\ 0.3 & 1.2 & 0.8 \end{bmatrix}$$

Step-by-Step ALS Solution

Iteration 1: Initialize randomly

$$\mathbf{W}^{(0)} = \begin{bmatrix} 0.5 & 0.3 \\ 0.4 & 0.6 \\ 0.2 & 0.8 \end{bmatrix}, \quad \mathbf{H}^{(0)} = \begin{bmatrix} 1.0 & 0.5 & 0.2 \\ 0.3 & 1.2 & 0.8 \end{bmatrix}$$

Update User 1: Only use observed ratings (positions 1,3)

$$\mathbf{y}_1 = [5, 2]^T \quad (20)$$

$$\mathbf{X}_1 = \begin{bmatrix} 1.0 & 0.3 \\ 0.2 & 0.8 \end{bmatrix} \text{ (columns 1,3 of } \mathbf{H}^{(0)T}) \quad (21)$$

Step-by-Step ALS Solution

Iteration 1: Initialize randomly

$$\mathbf{W}^{(0)} = \begin{bmatrix} 0.5 & 0.3 \\ 0.4 & 0.6 \\ 0.2 & 0.8 \end{bmatrix}, \quad \mathbf{H}^{(0)} = \begin{bmatrix} 1.0 & 0.5 & 0.2 \\ 0.3 & 1.2 & 0.8 \end{bmatrix}$$

Update User 1: Only use observed ratings (positions 1,3)

$$\mathbf{y}_1 = [5, 2]^T \quad (20)$$

$$\mathbf{X}_1 = \begin{bmatrix} 1.0 & 0.3 \\ 0.2 & 0.8 \end{bmatrix} \text{ (columns 1,3 of } \mathbf{H}^{(0)T}) \quad (21)$$

Solve: $\mathbf{w}_1^{(1)} = (\mathbf{X}_1^T \mathbf{X}_1)^{-1} \mathbf{X}_1^T \mathbf{y}_1$

Continue for all users and movies...

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?
3. How to handle new users?

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?
3. How to handle new users?
4. How to handle the scale?

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?
3. How to handle new users?
4. How to handle the scale?

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?
3. How to handle new users?
4. How to handle the scale?

Suggested Solution:

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?
3. How to handle new users?
4. How to handle the scale?

Suggested Solution:

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?
3. How to handle new users?
4. How to handle the scale?

Suggested Solution:

Pop Quiz 4: Final Challenge

Master Check

You're Netflix's lead ML engineer. You have:

- 200M users, 15K movies
- 20B ratings (0.67% filled)
- Need real-time recommendations
- New users/movies arrive daily

Design your recommendation system:

1. Which algorithm: ALS or SGD? Why?
2. What rank r would you choose?
3. How to handle new users?
4. How to handle the scale?

Suggested Solution:

Key Insights Summary

1. **Sparsity** \Rightarrow **Factorization**: Sparse rating matrices can be approximated by low-rank factorizations

Key Insights Summary

1. **Sparsity \Rightarrow Factorization:** Sparse rating matrices can be approximated by low-rank factorizations
2. **Latent Features:** Users and items are characterized by latent factors (not manually defined!)

Key Insights Summary

1. **Sparsity \Rightarrow Factorization:** Sparse rating matrices can be approximated by low-rank factorizations
2. **Latent Features:** Users and items are characterized by latent factors (not manually defined!)
3. **Bilinear Problem:** Non-convex jointly, but convex individually \rightarrow Alternating optimization works well

Key Insights Summary

1. **Sparsity \Rightarrow Factorization:** Sparse rating matrices can be approximated by low-rank factorizations
2. **Latent Features:** Users and items are characterized by latent factors (not manually defined!)
3. **Bilinear Problem:** Non-convex jointly, but convex individually \rightarrow Alternating optimization works well
4. **Scale Matters:** Algorithm choice depends on data size and computational constraints

Key Insights Summary

1. **Sparsity \Rightarrow Factorization:** Sparse rating matrices can be approximated by low-rank factorizations
2. **Latent Features:** Users and items are characterized by latent factors (not manually defined!)
3. **Bilinear Problem:** Non-convex jointly, but convex individually \rightarrow Alternating optimization works well
4. **Scale Matters:** Algorithm choice depends on data size and computational constraints
5. **Real-World Complexity:** Regularization, bias terms, cold start, implicit feedback all matter

Key Insights Summary

1. **Sparsity \Rightarrow Factorization:** Sparse rating matrices can be approximated by low-rank factorizations
2. **Latent Features:** Users and items are characterized by latent factors (not manually defined!)
3. **Bilinear Problem:** Non-convex jointly, but convex individually \rightarrow Alternating optimization works well
4. **Scale Matters:** Algorithm choice depends on data size and computational constraints
5. **Real-World Complexity:** Regularization, bias terms, cold start, implicit feedback all matter

Key Insights Summary

1. **Sparsity \Rightarrow Factorization:** Sparse rating matrices can be approximated by low-rank factorizations
2. **Latent Features:** Users and items are characterized by latent factors (not manually defined!)
3. **Bilinear Problem:** Non-convex jointly, but convex individually \rightarrow Alternating optimization works well
4. **Scale Matters:** Algorithm choice depends on data size and computational constraints
5. **Real-World Complexity:** Regularization, bias terms, cold start, implicit feedback all matter

The Mathematical Beauty:

Collaborative Filtering = Matrix Factorization = Dimensionality Reduction

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations
- **Graph Neural Networks:** Leverage user-item interaction graphs

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations
- **Graph Neural Networks:** Leverage user-item interaction graphs
- **Multi-armed Bandits:** Exploration vs exploitation in recommendations

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations
- **Graph Neural Networks:** Leverage user-item interaction graphs
- **Multi-armed Bandits:** Exploration vs exploitation in recommendations

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations
- **Graph Neural Networks:** Leverage user-item interaction graphs
- **Multi-armed Bandits:** Exploration vs exploitation in recommendations

Applications Beyond Movies:

- E-commerce (Amazon, eBay)

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations
- **Graph Neural Networks:** Leverage user-item interaction graphs
- **Multi-armed Bandits:** Exploration vs exploitation in recommendations

Applications Beyond Movies:

- E-commerce (Amazon, eBay)
- Music streaming (Spotify, Apple Music)

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations
- **Graph Neural Networks:** Leverage user-item interaction graphs
- **Multi-armed Bandits:** Exploration vs exploitation in recommendations

Applications Beyond Movies:

- E-commerce (Amazon, eBay)
- Music streaming (Spotify, Apple Music)

Extensions and Advanced Topics

Beyond Basic Matrix Factorization:

- **Non-negative Matrix Factorization (NMF):**
Interpretable factors
- **Deep Matrix Factorization:** Neural networks for non-linear patterns
- **Factorization Machines:** Handle multi-way interactions
- **Variational Autoencoders:** Probabilistic approach to recommendations
- **Graph Neural Networks:** Leverage user-item interaction graphs
- **Multi-armed Bandits:** Exploration vs exploitation in recommendations

Applications Beyond Movies:

- E-commerce (Amazon, eBay)
- Music streaming (Spotify, Apple Music)

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations
5. SGD is better than ALS for all applications

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations
5. SGD is better than ALS for all applications

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations
5. SGD is better than ALS for all applications

Answers:

1. **False** - Works with implicit feedback too (clicks, views)

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations
5. SGD is better than ALS for all applications

Answers:

1. **False** - Works with implicit feedback too (clicks, views)

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations
5. SGD is better than ALS for all applications

Answers:

1. **False** - Works with implicit feedback too (clicks, views)

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations
5. SGD is better than ALS for all applications

Answers:

1. **False** - Works with implicit feedback too (clicks, views)

Final Pop Quiz: Comprehensive Understanding

Mastery Test

True or False? Explain your reasoning:

1. Matrix factorization can only work with explicit ratings
2. ALS always converges to the global optimum
3. A rank-1 factorization means all users have identical preferences
4. Adding regularization always improves recommendations
5. SGD is better than ALS for all applications

Answers:

1. **False** - Works with implicit feedback too (clicks, views)

Questions?

Thank you for your attention!

Next: Deep learning approaches to recommendation systems